

Interaction Models Lower Prediction Error in 121 of 160 MGnify Studies: a Microbiome Benchmark Audit

Udita Phookan

Executive summary

The positive result in this archived within-study audit is a predictive one: an interaction-augmented model lowers held-out composition error against a calibrated population prior in 121 of 160 MGnify study tables with per-study paired Wilcoxon tests and BH adjustment across all 160 study tables; the median relative gain across studies is 14.5%. This is not a causal interaction finding or a win on an untouched external source. The 160 MGnify entries are one study-table type with mixed assay modalities and possible biological-source overlap. This summary gives the paper's result hierarchy; the tests and exceptions remain below.

Two human cNODE tables do not distinguish the published interaction model's error from a simple presence-only null interval on their published-versus-recomputed comparison. Four and five MDSINE2 mice make the simple training-population interpolation baseline competitive, but neither subject-taxon pairs nor mouse-cohort contrasts supply independent validation. The original prediction that detection-only scoring explained its performance was falsified by all-timepoint scoring. The apparent anaerobe-keystone pattern fails alternate network definitions; the CatBoost+SHAP disease-literature proxy analysis is a redirected, exploratory result, not a replicated ecological discovery.

Abstract

A calibrated population prior and an interaction-augmented model were compared within 160 MGnify study tables. The interaction model reduced held-out composition error in 121 of 160 tables with per-study paired Wilcoxon tests and BH adjustment across all 160 study tables; median relative error reduction across tables was 14.5%. This is a positive predictive result on previously analyzed, heterogeneous tables, not a causal interaction estimate or an untouched external-source win. It does not identify independent cohorts across the MGnify accession list.

The six cNODE ecosystem tables give a second test of simple baselines. Published cNODE medians for two human-associated tables fall within the presence-null bootstrap

interval; the local cNODE implementation does not exactly reproduce the published medians, so these are not same-code head-to-heads. On released MDSINE2 data, a simple training-population trajectory interpolator reaches median log10 RMSE 0.692 versus full MDSINE2 1.093 in five UC-donor mice; in four healthy-donor mice it ties MDSINE2 without modules (0.919 versus 0.913) and loses to RA-MDSINE2 without modules (0.883). Subject-taxon rows are clustered within mice. All-timepoint scoring did not erase the baseline's performance, falsifying the proposed detection-only-metric explanation.

Exploratory keystone analyses did not yield a method-robust ecological discovery. An anaerobe association in ridge-inferred networks failed in graphical-lasso and betweenness networks. A downstream CatBoost+SHAP score fit a BugSigDB disease-literature proxy better than the ridge score, but followed failed primary tests and lacks untouched validation. This work offers a predictive audit and transparent failures, not a personalized or perturbation-validated digital twin. The expanded platform and independent discovery remain open goals.

1. Introduction

Research candidates for digital twins of the microbiome promise in-silico trials: remove a species, add an antibiotic, change a diet, and read out the community before touching an animal or a patient. The community therefore invests heavily in interaction-aware models: generalised Lotka-Volterra (gLV) systems, Bayesian gLV with interaction modules (MDSINE, MDSINE2), and neural ODEs (cNODE). Their value rests on a premise that is rarely tested directly: that the interaction terms carry predictive information beyond what a population average already provides.

Strong baselines matter because microbiome data are dominated by taxon identity. A taxon that is abundant in one host is usually abundant in another. A predictor that simply remembers typical abundances, conditioned on which taxa are present, may already explain most of the variance. If so, the headline accuracy of an interaction model says little about whether its interactions are right, and downstream uses such as keystone-species ranking or perturbation design rest on weaker ground than the accuracy suggests.

We test that premise on the two public benchmarks that serve as reference benchmarks for two twin tasks. Our contributions are: (i) an exact leave-one-out re-run of the cNODE benchmark with a bootstrap null interval; (ii) a verified reproduction of the MDSINE2 cross-validation metric from the authors' notebook and Source Data, matching the published MDSINE2 median; (iii) a presence-conditional population forecaster that tops the UC leaderboard and ties the healthy one; and (iv) an explicit account of why the metric allows this.

Three claim levels are kept separate throughout. Prediction asks whether a model forecasts held-out composition or trajectories. Mechanism asks whether its fitted terms recover causal ecology. Discovery asks whether it names a biological principle that survives replication. This paper reports results at each level separately and does not let a success at one level stand in for another. The novelty is the auditing framework, its scale, and the pre-registered falsification discipline - not any single model, null, or benchmark re-run.

The paper's hierarchy: the primary contribution is the predictive audit framework (does a twin beat ecological priors, and when). The major findings are that some benchmarks collapse to ecological priors (human cNODE datasets, the MDSINE2 UC cohort) while others require learned structure (121 of 160 MGnify studies). The secondary exploration is the biological-discovery programme, which mostly fails replication and is reported as such.

2. Problem statements and notation

2.1 Assemblage-to-composition (cNODE task)

Let N be the species pool. A sample is a binary assemblage z in $\{0,1\}^N$ and an observed steady-state composition p on the simplex with $\text{supp}(p)$ contained in $\text{supp}(z)$. This composition predictor is a research component, not a personalized perturbation-validated digital twin. It is a map $\phi: z \rightarrow \hat{p}$. Error is Bray-Curtis dissimilarity:

$$BC(p, q) = \sum_i |p_i - q_i| / \sum_i (p_i + q_i) \quad (1)$$

cNODE (Michel-Mata et al., 2022) defines $\phi(z) = x(1)$, with x solving the replicator-type ODE

$$dx/dt = x \odot (f(x) - 1 x^T f(x)), \quad f(x) = Wx, \quad x(0) = z / |z| \quad (2)$$

so that $\sum_i x_i$ is conserved and absent taxa stay at zero. Our presence-only null is

$$\hat{p}_i(z) = z_i m_i / \sum_j z_j m_j, \quad m_i = (1/|T|) \sum_{\{s \in T\}} p_{\{s,i\}} \quad (3)$$

where T is the training set. It has no interaction parameters; its only parameters are the N training means. Our gLV replicator baseline integrates

$$dx/dt = x \odot (r + Ax - \phi(x) 1), \quad \phi(x) = x^T (r + Ax) \quad (4)$$

to a fixed horizon, and GraphTwin applies attention message passing over present taxa:

$$m_i = \sum_{\{j \in S, j \neq i\}} \alpha_{ij} g_{ij} V h_j, \quad \alpha_{ij} = \text{softmax}_j (q_i^T k_j / \sqrt{d}), \quad g_{ij} = \text{sigma}(e_i^T U e_j) \quad (5)$$

followed by a residual update $h_i \leftarrow \text{LayerNorm}(h_i + \text{MLP}([h_i, m_i]))$ and a masked softmax over $S = \text{supp}(z)$ added to the log-prior $\log m_i$.

2.2 Held-out-subject forecasting (MDSINE2 task)

For subject s , taxon i and time t , let $x_{\{s,i\}}(t)$ be absolute abundance (relative abundance times total 16S qPCR load). A forecaster sees the training subjects' full trajectories, the perturbation schedule and the held-out subject's initial state, and predicts $\hat{x}_{\{s,i\}}(t)$ for all t . The MDSINE2 paper's evaluation (`paper_figures/fig3_cross_validation.ipynb`) computes, per (subject, taxon),

$$RMSE_{\{s,i\}} = \sqrt{(1/|D_{\{s,i\}}|) \sum_{t \in D_{\{s,i\}}} (\log_{10}(x_{\{s,i\}}(t) + \epsilon) - \log_{10}(\hat{x}_{\{s,i\}}(t) + \epsilon))^2}, \quad D_{\{s,i\}} = \{t : x_{\{s,i\}}(t) > 1e-5\} \quad (6)$$

with $\epsilon = 1e3$ in the notebook's units, and reports the distribution of $RMSE_{\{s,i\}}$ over all pairs. Note the set $D_{\{s,i\}}$: only timepoints with detected truth are scored. Our simple training-population trajectory interpolator is

$$\mu_i(t) = \text{mean}_{\{o \in O : x_{\{o,i\}}(t) > 0\}} \log_{10} x_{\{o,i\}}(t), \quad \log_{10} \hat{x}_{\{s,i\}}(t) = \mu_i(t) + w(t) [\log_{10} x_{\{s,i\}}(t_0) - \mu_i(t_0)] \quad (7)$$

$$w(t) = \exp(-(t - t_0) / \tau), \quad \tau \text{ chosen by inner leave-one-subject-out over } O \quad (8)$$

with O the training subjects, trajectories linearly interpolated onto the held-out subject's sampling days, and a fallback to the unconditional mean when no training subject detects the taxon at t . Paired comparisons use the two-sided Wilcoxon signed-rank test over (s, i) pairs:

$$W = \sum_k \text{sgn}(d_k) R_k, \quad d_k = RMSE^{\text{ours}}_k - RMSE^{\text{other}}_k \quad (9)$$

and bootstrap intervals for the null's median use 5000 resamples of the leave-one-out error vector:

$$CI_{95} = [Q_{0.025}, Q_{0.975}] \text{ of } \{ \text{median}(e^*_b) \}_{b=1..5000}, \quad e^*_b \sim \text{resample}(e) \quad (10)$$

2.3 Why the detection-conditional metric favours population priors

Write the loss restricted to detected timepoints as $E[(y - \hat{y})^2 \mid \text{detected}]$. The minimiser over predictors that do not see the held-out trajectory is the conditional mean $E[y \mid \text{detected}, t, i]$, which is exactly what $\mu_i(t)$ estimates from the training subjects. A dynamical model trained on all timepoints, or simulating extinctions, may predict low abundance where the truth happens to be detected, and pays heavily for it; its errors where the truth is undetected are never scored. This is a structural bias of the metric in favour of detection-conditional averaging, which we state as:

$$\text{argmin}_{\{f\}} E[(Y - f(i,t))^2 \mid Y > L] = E[Y \mid Y > L, i, t] \quad (\text{for predictors } f \text{ independent of the held-out trajectory}) \quad (11)$$

3. Data

All data are public, unmodified, and fetched by accession. Table 1 is the dataset manifest. We count 169 distinct primary accession-level datasets: the nine below plus 160 MGnify studies (Appendix A). The datasets ledger also records secondary trait and literature reference tables; these are not independent cohorts.

Table 1. Dataset manifest (distinct, accession-level). n = samples or subjects.

#	dataset	source / accession	n	use
1	Ocean	github.com/yixueyang/cNODE (Michel-Mata 2022)	269	cNODE benchmark
2	Drosophila gut	same	24	cNODE benchmark
3	Soil in vitro	same	48	cNODE benchmark
4	Soil in vivo	same	678	cNODE benchmark
5	Human oral	same	143	cNODE benchmark (de-dup: 143 vs 150 reported)
6	Human gut	same	106	cNODE benchmark
7	MDSINE2 healthy cohort (counts, qPCR, perturbations)	github.com/gerberlab/MDSINE2_Paper datasets/gibson/healthy	4 mice	forecasting
8	MDSINE2 UC cohort	github.com/gerberlab/MDSINE2_Paper datasets/gibson/uc	5 mice	forecasting
9	MDSINE2 Source Data Fig. 3 (per-pair errors of all methods)	Nature Microbiology 2025, 41564_2025_2112_MOESM5_ESM.xlsx	556 / 601 pairs	head-to-head

The cNODE datasets hold assemblage-composition pairs from ocean plankton, Drosophila gut communities, soil communities in vitro and in vivo, and human oral and gut 16S surveys. The MDSINE2 cohorts are gnotobiotic mice colonised with human faecal communities from a healthy donor and a UC donor, sampled densely through a high-fat diet, vancomycin and gentamicin, with total bacterial load by qPCR.

4. Methods

4.1 cNODE audit

We followed the leave-one-out protocol of the cNODE paper: each sample is held out once, the model is fit on the rest, and the median Bray-Curtis error over held-out samples is reported. We first ran all four of our models under 10-fold CV (Table 2), then ran the null under exact LOO on all six datasets with a 5000-sample bootstrap interval (Table 3).

4.2 MDSINE2 reproduction

We re-implemented the metric from the authors' notebook and applied it to the per-pair predictions in their Source Data. Our recomputed MDSINE2 median on the healthy cohort is 1.061, matching the published box plot, which verifies the pipeline. Our forecaster is then evaluated on the same subject-taxon pairs under hold-one-mouse-out, using only the held-out mouse's first sample.

Table 2. Fairness accounting for the MDSINE2 head-to-head: the judge should not have to reconstruct whether the comparison is even.

component	MDSINE2 pipeline (authors)	population forecaster (ours)
evaluation pairs	per subject-taxon, authors' Source Data	identical pairs, same Source Data
metric	authors' notebook, detected timepoints, log10 RMSE	identical re-implementation (recomputed MDSINE2 median 1.061 matches published)
split	hold-one-mouse-out	hold-one-mouse-out, same folds
taxa selection / filtering	authors'	unchanged
zero / detection handling	authors' detection rule in the notebook	identical code path
training information	MCMC-fitted dynamics on training mice	training mice's trajectories only; the held-out mouse contributes its first sample as initial state
interaction parameters	fitted gLV + interaction modules	none

4.3 Software and hermetic tests

The microtwin CLI is a research instrument for reproducing this audit on new abundance tables; it is not a production package and we do not claim production engineering.

The code is in `src/microtwin` (`data.py`, `models.py`, `evaluate.py`, `dynamics.py`, `popforecast.py`). The archived early pytest suite ran offline on small fixtures: simplex and masking constraints, Bray-Curtis identities, gLV integration sanity, forecaster invariants.

5. Results

5.1 Scale-up: interaction audit across 160 MGnify studies

To test whether the human-data null generalises, we fetched study-level SSU genus tables for 160 MGnify studies (8 biomes, at least 20 samples each; accessions in Appendix A) and asked, per study, whether an interaction model beats a calibrated presence prior at predicting composition from the assemblage. The prior is the two-way log-linear fit

$$\log p_{si} = \mu_i + c_s \text{ (present entries only)}, \quad p\text{-hat}_i(z) = \frac{z_i e^{\mu_i}}{\sum_j z_j e^{\mu_j}} \quad (12)$$

in which the sample offset c_s absorbs the closure constant, so μ is not biased by which other taxa are present. The interaction model - an association model fitted on

observational abundances, not a validated causal interaction network - adds a steady-state generalised Lotka-Volterra shift fitted by ridge regression:

$$p\text{-hat}_i(z) \propto z_i \exp(\mu_i + b_i + \sum_j W_{ij} z_j), \quad (b_i, W_i) = \operatorname{argmin}_{\sum_{s=1}^S (\log p_{si} - \mu_i - c_s - b_i - W_i z_s)^2 + \lambda \|W_i\|^2} \quad (13)$$

with lambda chosen by inner 3-fold CV. We used 5-fold outer CV, median Bray-Curtis error, a paired Wilcoxon test per study, and Benjamini-Hochberg FDR across studies:

$$q_{-}(k) = \min_{\{j \geq k\}} (m p_{-}(j) / j) \quad (14)$$

Interactions were better at FDR < 0.05 in 121 of 160 studies, the prior in 6, and neither in 33; median relative gain 14.5%.

Table 3. Interaction audit by biome (results/mgnify_audit_summary.json).

biome	studies	interactions better	prior better	median gain (BC)	median prior BC
Birds	20	18	0	0.107	0.375
Digestive system	41	21	3	0.004	0.176
Fish	15	12	1	0.113	0.420
Insecta	7	5	0	0.052	0.345
Mammals	20	14	2	0.056	0.386
Oral	20	17	0	0.040	0.326
Plants	20	18	0	0.081	0.330
Skin	17	16	0	0.075	0.374




Figure 1. Relative gain of the interaction model per study. Red: interactions better (FDR < 0.05); dark: prior better; grey: no difference.

A candidate exception, the human digestive system (21 of 41), did not survive scrutiny. A narrow study-title pattern indicating a metagenome assembly flags 33/41 digestive-system records and 0/119 records outside that biome. This flag does not certify assay type for the 127 unflagged studies, and absence of a phrase does not prove an amplicon assay. The eight unflagged digestive tables have median relative gain 0.210 and eight original-BH interaction wins; non-digestive tables have median gain 0.192 (one undefined relative-gain denominator). This is a viewed, noncausal association confounded by source, assay, pipeline and biome, not an amplicon-only external validation. Earlier audit prose stated 2/118 non-gut assemblies and 0.184 median elsewhere; the denominator was internally inconsistent and neither value reproduced under the now-pinned title rule. We withdraw that comparison rather than treating an unflagged title as confirmed amplicon. The cNODE human-table null (Section 5.4) remains task-specific and does not imply that human microbiomes lack useful predictive interaction features.

5.1.1 Original-project identity sensitivity and limits

As a post-verdict, exploratory source-family audit we asked whether a study-table accession can stand in for an independent original project. Before computing this sensitivity, we fixed a lexical selector: parse a single PRJNA/PRJEB/PRJDB token from each MGNify study name, assign tokenless or ambiguous names distinct unresolved labels, then retain the lowest MGYS identifier per known project. The old 160-table BH correction is held fixed; it is not recalculated after selection. This is a source-identity sensitivity on viewed data, not a new model fit or an untouched validation cohort.

Table 4. Exploratory project-identity sensitivity. Wins are original 160-study FDR calls; 131/160 rows lack a unique original-project token, so 'one per' does not mean independent families.

identity group	study tables	original FDR interaction wins	median relative gain
original audit	160	121	0.145
one per parsable original project	159	121	0.146

Only 29 of 160 study names contain one extractable original project token. Two records, MGYS00006825 and MGYS00006862, both name PRJNA715245 and have discordant FDR win status. Live MGNify sample lists contain 197 and 84 different SRS IDs with no exact accession intersection: the pair shares a stated original project but is not an exact-sample duplicate. Keeping one project row by lexical accession removes one distinct-sample table, leaving 159 provisional families and 121 original-FDR interaction wins; the median relative gain changes from 14.52% to 14.62%. The unchanged win count does not establish 121 independent families. The 131 unresolved study names may hide other mirrors, source aliases, or shared subjects; SRS set difference also does not prove participant independence. Parent accessions and file-column biological units need a full crosswalk before a source-family leaderboard. See

results/PREREG_20260927_source_family_sensitivity.md
results/mgnify_family_sensitivity.json.

and

5.1.2 Descriptive assay-title stratification

To show how much the positive 121/160 predictive audit relies on study titles suggestive of metagenome assemblies, we froze a narrow assembly-phrase rule and joined the original manifest and BH result tables one to one. The original 160-table BH q-values remain unchanged, including the one record with zero prior and interaction medians; that record contributes a no-difference call but has undefined relative gain and is omitted only from the finite relative-gain median. This is a post-hoc descriptive stratification of already viewed data, not a fresh significance analysis or proof of a platform effect.

Table 5. Viewed MGnify title-flag sensitivity. 'Unflagged' means unknown assay, not verified amplicon; calls retain correction across the entire 160-study audit.

study-title category	tables	BH interaction wins	BH prior wins	neither	median relative gain	undefined relative gain
unflagged (unknown)	127	108	3	16	0.201	1
assembly indicated	33	13	3	17	0.015	0

Among 33 assembly-indicated study titles, 13 have original-BH interaction wins and median relative gain 1.51%; among 127 unflagged titles, 108 have wins and the finite relative-gain median is 20.05%. All 33 flagged titles in this narrow rule happen to be in the digestive-system biome, so this comparison cannot separate assay from biome. The old two-other-assemblies assertion was not reproduced by this rule: it flags zero of 119 non-digestive titles. A broad substring search finds four non-digestive titles, but two use 'assemblage' or 'community assembly' rather than describing a metagenomic assembly. Neither title rule verifies all unflagged assay types. In particular, eight unflagged digestive records with eight BH wins do not establish eight independent amplicon cohorts or an untouched transfer result. A defensible comparison requires source method records and comparable source families, then a frozen re-analysis. This title flag cannot settle modality. Protocol and data: results/PREREG_20260927_mgnify_assay_descriptive.md and results/mgnify_assay_descriptive.json.

5.1.2a Zero-median table: an information-content diagnostic

A separate diagnostic checks the fish-associated MGYS00006086 row rather than hiding its awkward zero-over-zero relative gain. The original audit included this row among 160 BH-adjusted study tables and counted it as neither model winning, with zero median held-out Bray-Curtis error for both. The live MGnify SSU download recorded by the manifest was opened through its published URL and processed with the original genus collapse, positive-sample filter, 5% taxon-prevalence filter and sample cap. The repository does not contain the archived derived TSV, so this is a reconstruction from a

current source snapshot rather than a byte-by-byte replay of the original raw artifact. Its 142-by-11 matrix reproduces the manifest dimensions, and its 135-by-3 analyzed matrix reproduces the old audit dimensions. The source snapshot and its hash, as well as the locked check before the download, are recorded with the diagnostic.

After prevalence filtering, 3 of 11 genera remain, and 135 of 142 samples retain mass on those genera. There are 6 unique presence assemblages and 9 unique relative-composition vectors at twelve-decimal rounding. Most notably, 111 of 135 analyzed samples contain only one retained genus. For any such sample the true retained-vocabulary relative-composition vector is a unit vector. Both tested predictors condition on the observed support and normalize over present genera, so they must give the same unit vector. Thus a majority of zero errors is built into the task after filtering, not evidence that either method learned useful ecological structure. The other samples need not have zero error, and six presence patterns mean the table is not literally constant. The finding explains the zero medians without claiming that all individual held-out predictions match.

The diagnostic changes interpretation, not the analysis denominator. The row remains in the original 160-table BH family and in the no-difference category; its relative gain remains undefined rather than being set to zero. No model was refit and no p-value or q-value was recalculated. A future benchmark should set its information-content eligibility rule before viewing held-out outcomes, with minimum retained richness or nontrivial-composition coverage measured on training data only. That rule would define a different estimand and require its own complete screening ledger, not a post-hoc deletion of this fish table. Even a correctly filtered within-study result would still not establish cross-source transfer or causality. The exact MGnify endpoint and source SHA-256 are pinned in `results/audit_degenerate_table.json`; the locked check is `results/PREREG_20260927_audit_degenerate_table.md`.

5.1.2b Exact source-column linkage for the fish diagnostic

The low-information fish table also permits a narrow accession-level provenance check. We froze the match rule before reading the current MGnify analysis and sample lists: exact identity between every source SSU column and an analysis record's assembly-relationship identifier, followed by that analysis's sample relationship. No string similarity or metadata-only guess is used. The source download has 150 raw columns, of which 142 carry nonzero mass after genus collapse and are the columns recorded in the old manifest; 135 remain after the audit's 5% prevalence screen removes massless rows. Thus 150, 142 and 135 refer to distinct, explicit processing stages, not competing counts of independent subjects.

The current MGnify API lists 157 analysis records and 157 sample records, each endpoint fully covered in one page. Exactly 150 of 150 raw SSU columns, 142 of 142 nonzero columns and 135 of 135 analyzed columns match unique analysis assembly IDs. All 135 analyzed columns then link through those analysis records to distinct current

sample IDs. These matches materially strengthen the column-to-record identity for this one viewed source, rather than relying on a study-level title or aggregate count. Every matched current analysis has experiment type 'assembly'; every linked sample has explicit host scientific name *Salmo salar*, while all normalized species fields are missing. The explicit salmon name is a direct metadata observation, not a repair of the missing normalized field.

The result does not retroactively turn the source into an independent fish cohort or assay-certified amplicon dataset. The public 5.0 SSU file's assembly IDs map to today's assembly-typed analyses, but the run-level instrument and library methods that produced each abundance column are not settled by an experiment-type label. Nor does distinct sample accessions demonstrate distinct individual salmon, cages or projects, and license review remains open. The 135 columns still provide only a within-table predictive audit with a severe one-taxon prevalence artifact. No error was recomputed and no archived BH result was moved. The exact URLs, response hashes and match counts are recorded in `results/fish_column_identity.json` under the locked `results/PREREG_20260927_fish_column_identity.md` protocol.

5.1.2c Exact run-column mapping exposes a human-host conflict

The other accession-level check concerns MGYS00006755, a record filed under the human digestive biome whose sample-level metadata already looked contradictory. We fixed the link rule before opening its source matrix: exact equality of the old SSU column identifier to a current analysis's run relationship, then follow that analysis's sample relationship to the current sample record. This check is possible because run accessions, not merely a study title, appear as the SSU table columns. The source snapshot reproduced its manifest's 83 samples and 353 genera after genus collapse, then the old audit's 83 samples and 150 retained genera after the standard prevalence and top-150 filters. Those dimensions match the archived record; an old derived-file checksum is unavailable, so the source snapshot is not a byte-for-byte historical replay.

All 83 of 83 analyzed run columns exactly match unique current analysis run relationships, and all of those analysis records link to distinct current sample accessions. Their current experiment-type labels are all 'amplicon'. This is stronger than the prior study-level aggregate, because the discordant host metadata now belongs to the same sample records linked by exact IDs to the 83 audit columns. In every one of these records, the explicit host scientific name is a nonhuman animal, while the normalized species label is *Homo sapiens*. The conflict remains unresolved; we neither infer that the normalized human label wins nor silently choose the explicit host name as final truth. The public source's study abstract describes non-ruminant herbivores, consistent with the explicit nonhuman names, but it cannot repair the source metadata's contradictory normalized field.

This materially weakens any human-gut claim based on the archived digestive label, even though the assay layer currently calls these matched records amplicon. Exact run-to-sample linkage does not verify independent animals, the full historical wet-lab method, sample rights, or an external-test separation; these results were already viewed in the 160-table audit. Their original BH-adjusted win call is left intact and not reinterpreted as a human-gut amplicon validation. A later eligibility ledger must quarantine host-discordant records until the data provider or primary methods resolves the contradiction. Exact source URLs, hashes, match counts and aggregate host-name distribution are in results/conflicted_host_column_identity.json; the pre-read rule is results/PREREG_20260927_conflicted_host_column_identity.md.

5.1.2d Oral archive column mapping and pipeline versions

MGYS00002238 is archived under Human:Digestive system but its source title describes oral changes in children with autism spectrum disorder. We fixed an exact-ID comparison of the historical SSU file's columns to current MGnify run relationships, followed by linked sample descriptions. The public SSU source snapshot reproduces 111 manifest columns and 312 genera; after the old prevalence/top-150 preprocessing, 111 sample columns and 111 retained taxa match the archived audit dimensions. All 111 retained column IDs exactly match current run relationships, but the analysis list has 222 records, not 222 samples. Every run has two analysis versions: 5.0 and 4.1, both amplicon-labeled and linked to the same sample ID. Because the SSU file path explicitly names pipeline 5.0, the crosswalk selects only its one 5.0 analysis per run rather than counting 4.1 as an independent observation.

The selected 111 analyses link to 111 unique sample accessions; all linked explicit host names and normalized species fields are *Homo sapiens*. The sample descriptions identify 59 saliva and 52 dental-plaque specimens, split across patient and healthy-control descriptions. None of the matched sample descriptions says stool or gut. This is positive source evidence that the 111 archived abundance columns belong to oral material despite the digestive-system archive label, not a claim about all data ever collected by the study. The duplicate pipeline versions do not double the participant denominator; nor do 111 distinct sample accessions prove 111 distinct people. Any within-study win remains an oral-table observation, not one human-gut replication.

The revised metadata-integrity matrix marks only the old column map as known for this record. Its explicit human-host compatibility is supported on the linked columns, but independent subject/family grouping, historical wet-lab method and data rights remain unverified; the human-gut amplicon admission gate is still blocked. There was no new prediction, p-value or BH adjustment. Per-endpoint URLs, hashes, 5.0-versus-4.1 version counts and aggregate descriptions are in results/oral_column_identity.json; the locked rule and version amendment are in results/PREREG_20260927_oral_column_identity.md.

5.1.2e Viewed human-gut candidate: exact IDs, uncertain biological units

MGYS00005154 is an old human-digestive archive table with a human-gut study title, current amplicon-labeled analyses, and complete human-host names in the prior source screen. To test whether those aggregate labels belong to the old audit data rather than to other runs in the same study, we froze an exact run-ID-to-analysis-and-sample join before opening its source table. The public pipeline-5.0 SSU snapshot reproduces 529 nonzero raw sample columns and 779 genera from the manifest. The old audit's top-150-taxon and 400-sample caps yield exactly 400 columns and 150 genera, matching its archived dimensions. All 529 raw and all 400 retained column run IDs are found in the current fully paged MGnify analysis list. This is source-column identity evidence on already-viewed data, not a new model result.

The current endpoint contains 965 analyses and 529 sample records, but analysis count does not equal biological count. Among the 400 retained run IDs, 320 have two analysis records and 80 have one. We selected exactly one pipeline-5.0 analysis per retained run because the source file explicitly names that pipeline version. Different versions of the same run agree on their linked sample ID in this check. The selected 400 analysis records are amplicon-labeled and link to 400 distinct sample accessions with explicit Homo sapiens host names and normalized Homo sapiens labels. This advances exact archive-column mapping in the metadata matrix, while leaving the independent-family field unknown.

Collection-site evidence required a post-inspection correction: all 400 linked sample descriptions repeat the study-level phrase 'Human gut microbiome differentiation viewed across cultures, ages and families', and the current sample metadata provide no literal 'body site' or 'collection site' key among these records. However, a later audit of structured sample metadata found that all 400 exact-linked records bear 'environment (material): ENVO:feces'; the other 129 current sample records carry the same material value. The prior narrow key scan omitted this field. It supports source-annotated fecal material, whereas the repeated study description alone did not suffice. Neither the submitted annotation nor the description can independently prove each collected specimen's anatomical site or wet-lab assay. Nor do 400 distinct sample accessions prove 400 unrelated people or an independent source family. Rights, participant grouping, project overlap and untouched split remain unverified. The gate therefore blocks final human-gut amplicon admission on family independence even though six other amended metadata fields pass for this candidate; passing that remaining family check would still require manual validation. No score or BH result was changed. Exact source URLs, response hashes, pipeline-version counts and aggregate sample descriptions: results/gut_candidate_column_identity.json; locked method: results/PREREG_20260927_gut_candidate_column_identity.md. Post-inspection material amendment and live sample-response hashes:

results/PREREG_20260927_feces_material_amendment.md
results/gut_material_identity.json.

and

5.1.2e.1 Family and publication identifiers remain unresolved

A later source-family check read the MGnify study and all current 529 sample records without examining abundance outcomes. The study record identifies PRJEB3079 / ERP001500 and its abstract describes 531 people, including mono- and dizygotic twins. All 529 exposed sample-metadata records carry location, environment, host and depth/elevation keys, but no explicit person, family, household or donor-ID key. The 529 current sample records and the 531 people in the abstract are distinct reported counts without a verified row-level crosswalk. Neither the 400 unique exact-linked sample accessions nor a single BioProject prove 400 unrelated people or independence from a source mirror. The biological-family axis therefore remains unknown rather than false or true, and source admission remains blocked.

The abstract text matches the 2012 Nature article 'Human gut microbiome viewed across age and geography' (DOI 10.1038/nature11053), which discusses fecal samples, families and twins. Yet MGnify's separate study-publications endpoint links this accession to a different 2021 article (DOI 10.1016/j.csbj.2021.07.009). We record this publication-link conflict rather than infer whether the publication endpoint citation was misassigned; the direct ENA assay bridge below identifies the deposited 16S run family. Primary person and family mapping, rights and mirror checks remain needed. This finding is metadata-only, not a reassessment of the original BH scores or an untouched external benchmark. Provider response URLs, hashes and aggregate key counts: results/gut_family_evidence.json; the dated diagnostic: results/PREREG_20260927_gut_family_evidence.md. Primary article: <https://www.nature.com/articles/nature11053>.

5.1.2e.2 Exact ENA assay bridge: what the identifier establishes

The earlier family diagnostic could not connect the MGnify archive to the original article's sequencing deposit from its abstract alone. A subsequent direct provider check crossed that narrower bridge without opening a new predictive outcome. ENA's ERP001500 study XML lists PRJEB3079 as its secondary project accession and qiime_study_850 as its submitted alias. The 2012 Gordon Lab source portal lists qiime:850 for its Illumina V4 16S deposit, separate from qiime:621 for fecal shotgun data. The identifier, rather than resemblance of titles, connects the archive's ENA study to that 16S deposit. This does not make every source annotation correct or supply an article-to-person linkage. ENA XML: <https://www.ebi.ac.uk/ena/browser/api/xml/ERP001500> ; author portal: <https://gordonlab.wustl.edu/supplemental-data/supplemental-data-portal/yatsunenکو-et-al-2012/> .

We further compared actual run identifiers, not just study labels. ENA's current ERP001500 read-run report exposes 530 unique run accessions and 529 unique sample accessions. All 530 reported runs carry AMPLICON library strategy, PCR selection and ILLUMINA instrument platform. The original MGnify SSU table has 529 raw run-ID columns and the unchanged archived filtering retains 400. Every one of the 529 raw column IDs and all 400 retained IDs exactly appears in the ENA run set; one ENA run is not a source-table column. Thus the old columns, not merely an unrelated accession in the same BioProject, have source-supported Illumina amplicon provenance. A spot-check of one ENA experiment shows 16S V4, but V4 uniformity across all 400 retained runs has not been independently checked. Full aggregate report, URLs and hashes are frozen in results/PREREG_20260928_primary_assay_crosswalk.md. The historical assay is now better grounded than the earlier manuscript text says, while the biological-family field stays unknown.

An assay bridge and a biological-independence bridge answer different questions. The first says which read runs contributed to an abundance table and what ENA calls their sequencing strategy. The second asks whether rows represent unrelated people, a twin pair, a repeated sample, or a mirror of a published cohort. The ENA report's 529 distinct sample accessions cannot resolve those relationships. The study abstract describes 531 people and twins, but neither number should be forced onto the 400 retained run columns as a participant count. The MGnify linked-publications endpoint's 2021 citation also remains inconsistent with the matching 2012 study abstract; the exact ENA alias resolves the deposited study identity without silently resolving that bibliographic conflict. This table is already viewed, so even a later family map could not turn the diagnostic into an untouched final test. No Wilcoxon statistic or BH correction in the 160-study archive was recomputed here.

5.1.2f Exact old-column mapping reveals a site-code contradiction

MGYS00006794 has a title describing longitudinal stool and oral communities in patients with multiple sclerosis, but the archived biome places it under Human:Digestive system. We fixed an exact source-column site check before reading its original pipeline-5.0 SSU file and current MGnify analyses and samples. The source table has 184 nonzero run columns and 439 collapsed genera; after the original prevalence and top-150 filters, 184 columns and 150 genera reproduce its archived audit dimensions. All 184 old run IDs exact-match one current pipeline-5.0 analysis each, and these point to 184 distinct current sample accessions. This establishes that the site codes examined belong to the actual old table columns rather than to unrelated samples in the study.

Each linked sample's structured environment-material field is UBERON:0002097, which the EMBL-EBI Ontology Lookup Service labels 'skin of body'. Each linked environment-feature field is ENVO:2100003, labeled 'skin environment'. The source's free-text descriptions distinguish multiple-sclerosis treatment phases and controls, not the

specimen site. Its normalized species is *Homo sapiens*, but no usable explicit host scientific name was supplied in these sample records. Thus there is a direct structured-site-versus-study-text/biome metadata conflict on all 184 archived columns. We do not adjudicate specimen tissue by choosing the structured codes or the title: either the site codes or the higher-level classification may be mistaken, and the original wet-lab record would have to settle which source actually produced the table.

This finding required a dated amendment of the admission prototype. The prior six-field gate tested host identity but conflated a human host with a human GUT sample. A separate seventh `human\_gut\_site\_compatible` field now blocks unknown or conflicting collection site. MGYS00006794's field is false from the structured-site mismatch; MGYS00002238 is false from exact linked oral sample descriptions. Of the other six, MGYS00005154 is source-metadata compatible after the separately dated feces-material amendment; five remain unknown pending sample-specific site evidence. These are metadata blockers, not a retrospective assay-stratified performance analysis, and all eight viewed records remain blocked. None of the 160 original BH-adjusted calls or old within-study scores changed. The amendment is explicitly dated after this source check, not passed off as a rule frozen before the original audit. Source response and ontology hashes: results/mixed_site_column_identity.json; protocol and site-axis amendment: results/PREREG_20260927_mixed_site_column_identity.md and results/PREREG_20260927_site_gate_amendment.md.

5.1.3 Live source context for the eight unflagged digestive titles

The first study-level screen asked whether eight title-unflagged records could be described uniformly as human-gut amplicon studies, using live MGnify study endpoints without re-scoring taxon outcomes. Later exact-column checks for four records are reported above. Each endpoint returned a public study name, project IDs and abstract; the source responses and abstract strings are hashed, while the repository retains only non-private metadata and short method phrases. The MGnify study response does not provide a decisive per-run library-strategy field, so a phrase in an abstract is source context, not an assay certificate for every archived abundance column. The manifest's digestive-system biome also requires verification against collection source and organism. This screen leaves the original eight win calls untouched.

Table 6. Live public MGnify study metadata for eight prior title-unflagged digestive archive records. Phrases are not per-run assay certification; MGYS00006755 describes non-ruminant herbivores despite a human digestive archive biome label.

MGnify ID	BioProject	source abstract phrase	interpretation
MGYS00000633	PRJNA63661	no decisive phrase	unknown
MGYS00002238	PRJNA362687	no decisive phrase	unknown
MGYS00005154	PRJEB3079	no decisive phrase	unknown
MGYS00005230	PRJEB5729	no decisive phrase	unknown
MGYS00006006	PRJEB49206	ultra-deep metagenomic	study description only

MGnify ID	BioProject	source abstract phrase	interpretation
		sequencing	
MGYS00006755	PRJNA362214	V4 16S rRNA amplicons	study description only
MGYS00006794	PRJEB44538	16S rDNA sequencing	study description only
MGYS00006795	PRJNA492714	whole genome sequencing	study description only

The sources directly contradict the casual phrase 'eight amplicon human-gut studies'. MGYS00006006 describes ultra-deep metagenomic sequencing; MGYS00006795 mentions whole-genome sequencing; MGYS00006755 describes V4 16S rRNA amplicons from non-ruminant herbivores even though the archived biome reads Human:Digestive system; MGYS00006794 describes oral and stool 16S rDNA sequencing. MGYS00002238 has an oral-microbiome title and a short abstract, while several others give no decisive assay phrase in the study response. A mixed study abstract could describe multiple sample types, and neither the digestive-system label nor the absence of an assembly title can replace a sample-to-run library-strategy crosswalk. The within-study eight-of-eight BH result is still an archival arithmetic fact, but the stronger interpretation as eight independent, human gut, amplicon cohorts is withdrawn. An assay-stratified scientific result requires source-verified collection and instrument fields, harmonized taxonomy, and subject/source-family grouping before any new model test. Source URLs and response hashes are in results/gut_unflagged_source_screen.json; no external performance was computed in this screen.

5.1.4 Analysis-type and host metadata crosswalk: seven complete, one unresolved

The earlier source screen moved from abstracts to MGnify analysis records and sample-host metadata without mapping historical SSU abundance columns to analyses. Later exact joins in Sections 5.1.2c-f supply those maps for four viewed records. Seven of the eight title-unflagged digestive study IDs had complete, declared pagination for both analysis and sample endpoints. The eighth, MGYS00000633, declares 8,091 analyses and 1,563 samples across many pages; this bounded screen did not complete it, so its assay and host classification are explicitly unresolved. Each completed source record carries page URLs and response hashes, while only aggregate type and host counts enter this manuscript. The classification below is current MGnify analysis metadata, not a verified historical per-column label and not an external validation result.

Table 7. Complete MGnify analysis/sample metadata for seven of eight previously viewed digestive archive records. The eighth is unresolved due to incomplete pagination. Host conflict compares normalized Homo sapiens to an explicit nonhuman host scientific name, not a missing value.

MGnify ID	analyses	analysis types	samples	host-field conflict
MGYS00002238	222	amplicon: 222	111	0
MGYS00005154	965	amplicon: 965	529	0
MGYS00005230	1644	amplicon: 1644	915	0
MGYS00006006	457	assembly: 457	419	0
MGYS00006755	83	amplicon: 83	83	83
MGYS00006794	184	amplicon: 184	184	0

MGnify ID	analyses	analysis types	samples	host-field conflict
MGYS00006795	31	amplicon: 31	36	0
MGYS00000633	unresolved	unresolved	unresolved	unresolved

The title filter missed an assembly-derived study: MGYS00006006 has 457 of 457 current analyses typed 'assembly', despite not carrying the narrow assembly phrase in its title. In contrast, MGYS00006795's abstract mentioned whole-genome sequencing while all 31 current analyses are typed 'amplicon'; study prose and the output being audited are different evidence layers. The six completed sources besides the assembly record have all examined analyses labeled amplicon. In MGYS00006755, all 83 analyses are labeled amplicon, but all 83 normalized sample species fields say *Homo sapiens* while the source's explicit host scientific names give nonhuman animals. We cannot choose the more favorable field: the human host label is contested for that study. No outcome or original BH q-value was recomputed from these source metadata. Later exact run-to-analysis-to-sample joins for four viewed records are given in Sections 5.1.2c-f; the remaining old SSU files' columns still require joins and all require biological-family grouping before any assay-specific win rate. The seven-of-eight read reduces uncertainty in named records but does not certify eight independent human amplicon cohorts or resolve the archived gut effect. See results/PREREG_20260927_gut_analysis_metadata_crosswalk.md and results/gut_analysis_metadata_crosswalk_partial.json; MGYS00000633 remains open.

5.1.5 An explicit metadata-integrity matrix

A recurring error in a large research platform is to turn partial metadata into a single green eligibility check. The preceding screens make this risk concrete: an analysis-type count alone does not link a historical abundance column to an analysis; a normalized human species code need not agree with the source's host scientific name; and complete pagination does not establish an independent source family. The original matrix emitted a six-field status vector for each of the eight old title-unflagged digestive records; a seventh site field was added after the mixed-site conflict was found. It separates analysis pagination, uniform analysis type, sample pagination, explicit host compatibility, archive-column mapping, biological-family independence and the newly amended gut-site compatibility. `Unknown` is an intentional output rather than a false value that someone can accidentally average away. The uncompleted eighth source is unknown on every dimension. Four later exact run-to-sample joins have verified old-column maps: MGYS00006755 fails host compatibility; MGYS00002238 is oral saliva/plaque; MGYS00006794 has conflicting structured skin-site codes; and MGYS00005154 has an exact-linked feces material annotation but no independent-family certificate. The other four retain unknown archive-column mapping. Every source remains unknown on biological-family independence.

Table 8. Source metadata-integrity matrix, not an eligibility leaderboard. Study IDs shorten MGYS0000 to M (e.g., M6755 = MGYS00006755). M2238, M5154, M6755 and M6794 have exact old SSU run-column maps. M2238 is oral, M6794 has conflicting skin-site codes, M6755 has a host conflict; all eight lack an independent-family certificate.

study	analysis type	sample pages	explicit human host	gut site	old column map	source family
M0633	unknown	unknown	unknown	unknown	unknown	unknown
M2238	amplicon	yes	yes	conflict	yes	unknown
M5154	amplicon	yes	yes	yes	yes	unknown
M5230	amplicon	yes	unknown	unknown	unknown	unknown
M6006	assembly	yes	unknown	unknown	unknown	unknown
M6755	amplicon	yes	conflict	unknown	yes	unknown
M6794	amplicon	yes	unknown	conflict	yes	unknown
M6795	amplicon	yes	unknown	unknown	unknown	unknown

The distinction between false and unknown matters. MGYS00006755 is a host-field conflict on the exact 83 audited run columns: every linked sample reports a nonhuman host scientific name while the normalized species says *Homo sapiens*. Its known column mapping does not repair the biological-host disagreement. MGYS00005230, MGYS00006794 and MGYS00006795 have missing explicit host names for some or all samples and are therefore unknown on the explicit-host criterion, even when normalized species labels say human. MGYS00006006 is assembly-typed at the analysis layer and also has seven samples without usable explicit host names. MGYS00002238 and MGYS00005154 have complete human names and later exact run-column maps: the former links to oral saliva/plaque; the latter has a repeated gut-oriented description and a later identified structured feces material annotation on all 400 exact-linked samples, though no literal body-site field. MGYS00006794 has a fourth exact map, but its structured codes point to skin in conflict with its title and archive biome; explicit host names are missing. None of these four exact-mapped sources has independently verified subjects, families or rights. These statuses do not revise the 121/160 BH test or identify an assay-specific effect. They define why a future leaderboard admission procedure must check all needed evidence separately before treating any old accession row as a biological dataset. Machine-readable statuses and synthetic fail-closed tests live in `results/metadata_integrity_matrix.json` and `tests/test_metadata_integrity_matrix.py`.

5.1.5a Fail-closed source-admission prototype

To make the integrity matrix operational rather than a decorative warning, a local admission function now evaluates seven amended fields for each already-viewed digestive candidate. A human-gut amplicon metadata review requires complete analysis and sample pagination, uniform current amplicon analysis labels, exact old SSU column-to-analysis mapping, an explicit human host field without conflict or missing names, a separate source-annotated gut-compatible collection material, and verified biological-family independence. The seventh site axis was added after the skin-code discrepancy;

the original six-field gate lacked it. A blank, false or unknown value blocks the review and names every failed field. A normalized Homo sapiens code cannot substitute for explicit host or site; a title or accession count cannot prove an independent dataset. Even a feces material annotation is source metadata, not an independent tissue assay. These fields are necessary screening checks, not sufficient conditions for an external benchmark.

Applied to the 8 viewed title-unflagged records, the gate blocks 8; 0 reach metadata-ready manual review, and zero become independently validated external datasets. MGYS00006755 illustrates why the gate has several independent axes: its exact historical run-column map is now known and its current analyses are amplicon-labeled, yet its host fields conflict in every linked sample and its biological-family independence remains unknown. MGYS00005154 now passes six amended metadata fields, including structured feces material on its exact-linked 400 samples, but remains blocked on unknown biological-family independence. MGYS00006006 has assembly-typed analyses even though the narrow title rule missed it; MGYS00000633 remains unpaged and unknown. These are explicit engineering blockers for this human-gut amplicon question, not declarations that a nonhuman or assembly source has no use for another biological question.

Even a synthetic record passing all seven amended fields returns only 'metadata ready for manual review'. Historical library methods, source-specific rights, participant lineage, a frozen holdout and a same-task comparator must still be checked before any dataset could be admitted to an external leaderboard. Unit tests assert that each field independently blocks on an unverified value, that duplicate study identities fail, and that all eight real records remain blocked. No performance score is read by this gate, and the 160-study BH calls remain as archived. Protocol and exact reason-coded output: `results/PREREG_20260927_source_admission_gate.md` and `results/source_admission_screen.json`; implementation: `src/microtwin/source_admission.py`.

5.1.6 A conservative source-lineage guard

The original-project sensitivity motivates an explicit partition rule for later cross-source evaluation. A study accession, its sequencing-project accession, and a derived assembly project may be different handles for related biological material. If a researcher randomises those handles independently into training and test sets, the nominal number of held-out studies can increase while the number of underlying source projects does not. Source-family registration therefore precedes split generation: each verified alias maps to one canonical family, and every registered alias in that family must be assigned to the same partition. An accession missing from the registry is not presumed independent; it remains ineligible for the known-family separation certificate until lineage is checked. The check is deliberately asymmetric: a failed split proves that known source-family

separation was violated, whereas a passing split proves only that no collision in the registered aliases was found.

For the one currently confirmed MGnify lineage pair, two assembly records carry separate MGYS identifiers, separate ERP secondary accessions, and separate derived PRJEB BioProject identifiers, but both study titles identify PRJNA715245 as their original data project. The source-lineage guard registers those seven identifiers under PRJNA715245 and rejects a train/test split spanning MGYS00006825 and the ERP identifier of MGYS00006862. A same-partition assignment passes this narrow mechanical check; an unknown accession fails. This engineering test does not retrospectively split the old within-study audit, change its 160-table BH procedure, or imply that the pair has duplicate sample accessions: the live MGnify lists have no exact SRS intersection. Conversely, zero exact SRS overlap cannot exclude donor overlap, reuse of underlying reads under changed accession labels, or laboratory and analysis-pipeline correlations. The appropriate unit for a future external benchmark must be justified by its biological question and provenance, not selected to preserve a favorable win count.

A scalable provenance ledger would keep, for every candidate, the original experiment identifier, parent sequencing project, assembly parent, sample accessions, participant or cage identifier when permitted, collection site and date, assay modality, raw-file checksums, license, and train/development/final-test status. Matching would proceed from exact IDs and file hashes to documented project relationships, and only then to a clearly labeled manual review of near-duplicates; a name similarity alone cannot be promoted to identity. Modality-specific harmonisation should occur only within training folds, with its fitted taxonomy map and normalization frozen before opening a final-test source. For repeated human samples, subject-level grouping prevents a visit from appearing in both train and test; for longitudinal mouse data, cage and donor effects may require an even wider grouping. These are prospective safeguards. The existing 131 token-unresolved MGnify rows show why a 500- or 1000-row manifest does not yet equal a 500- or 1000-independent-source benchmark. Until lineage is resolved, the positive within-study predictive association remains the proper headline, not a cross-source performance claim.

5.2 MDSINE2 benchmark: UC beat, healthy tie

Table 9. UC cohort, official MDSINE2 metric, hold-one-mouse-out.

method	median	mean	pairs	ours better	mean diff (ours-other)	Wilcoxon p (2-sided)
PresenceConditionalPopulation (ours)	0.692	0.842	601	-	-	-
RA-MDSINE2 (No	0.784	0.883	601	384	-0.041	5.6e-06

method	median	mean	pairs	ours better	mean diff (ours-other)	Wilcoxon p (2-sided)
Modules)						
MDSINE2 (No Modules)	0.805	0.889	601	376	-0.047	8.0e-07
MDSINE2_s eed10	0.965	1.253	601	431	-0.411	1.0e-29
RA-MDSINE2	1.009	1.140	601	427	-0.298	2.8e-29
MDSINE2	1.093	1.417	601	471	-0.575	4.3e-44
gLV (elastic net)	1.582	1.819	601	512	-0.977	4.5e-72
gLV (ridge)	1.970	2.589	601	553	-1.747	4.9e-85
gLV-RA (ridge)	4.536	4.696	601	599	-3.854	4.2e-100
LRA	4.536	4.696	601	599	-3.854	8.7e-100
gLV-RA (elastic net)	4.536	4.696	601	599	-3.854	4.2e-100
cLV	4.536	4.696	601	599	-3.854	4.2e-100

Table 10. Healthy cohort, same protocol.

method	median	mean	pairs	ours better	mean diff (ours-other)	Wilcoxon p (2-sided)
RA-MDSINE2 (No Modules)	0.883	0.990	556	285	0.095	4.5e-02
MDSINE2 (No Modules)	0.913	1.024	556	299	0.061	5.3e-01
PresenceCon ditionalPopu lation (ours)	0.919	1.085	556	-	-	-
RA-MDSINE2	0.957	1.183	556	311	-0.098	3.6e-02
MDSINE2	1.061	1.195	556	335	-0.110	1.1e-04
gLV (elastic net)	1.446	1.592	556	425	-0.507	1.2e-35
gLV (ridge)	1.924	2.098	556	472	-1.013	1.4e-65
gLV-RA (ridge)	4.643	4.770	556	555	-3.685	9.1e-93
gLV-RA (elastic net)	4.643	4.770	556	555	-3.685	9.1e-93
cLV	4.643	4.770	556	555	-3.685	9.1e-93
LRA	4.643	4.770	556	555	-3.685	1.3e-92




Figure 2. Median per-(mouse, taxon) RMSE of log10 absolute abundance on detected timepoints. Red: our presence-conditional population forecaster.

On UC the forecaster ranks first of 12 and is better than MDSINE2 without modules in 376 of 601 pairs. On healthy it ranks third of 11. Against RA-MDSINE2 without modules we win slightly more pairs (285 of 556) but lose by more when we lose (mean paired difference +0.095 in their favour); the signed-rank test favours them (two-sided $p = 0.045$). Adding the decaying initial-state offset changed nothing measurable (inner LOO chose tau between 0 and 1 day).

5.2.1 Biological-unit sensitivity: mice, not taxon pairs

The box-plot and Wilcoxon table use 556 or 601 taxon-mouse errors, but the independently held-out biological units are four healthy-donor mice and five UC-donor mice. Taxon errors within one mouse share the same animal, treatment course and measurement conditions, so the pair-level p -values cannot serve as mouse-level replication. To expose the unit shift, we took each mouse's median paired taxon RMSE difference between the population interpolator and MDSINE2 without modules, then enumerated every n^n resample of n mouse-level differences, using the mean across mice and 2.5%/97.5% percentiles. The same mouse can recur in a resample; these are descriptive intervals for the already viewed cohorts, not fresh animals or multiplicity-adjusted confirmation.

Table 11. Archived biological-unit sensitivity against MDSINE2 without modules. Negative gap favors the simple population baseline. Exhaustive mouse resampling has 4^4 or 5^5 draws; these are descriptive viewed-cohort intervals, not independent-study inference.

cohort	score	mice	negative mouse gaps	mean mouse-median gap	mouse-resample 95% interval
healthy	detected-only	4	3	-0.036	[-0.0978, +0.0396]
healthy	all timepoints	4	4	-0.238	[-0.2514, -0.2284]
uc	detected-only	5	5	-0.106	[-0.1351, -0.0782]
uc	all timepoints	5	5	-0.173	[-0.1840, -0.1576]

On the official detected-only score the average mouse-median gap is -0.0356 for healthy and -0.1061 for UC; the healthy mouse interval crosses zero, while the UC interval in this viewed cohort does not. The 4/4 and 5/5 all-timepoint negative signs contradict the earlier prediction that adding undetected times would erase the baseline's advantage. They do not show that detected-only scoring is harmless in general, or that either cohort is stable while the other is perturbed: the source design includes diet and antibiotics in both mouse cohorts. The all-timepoint metric uses the same released measurements, not an independent experiment. The baseline and score choices were developed on these cohorts, so even an interval below zero cannot establish a new-source beat.

A defensible future power design would simulate correlated taxon trajectories at the mouse level, predefine effect size and perturbation schedule, refit both methods in every simulated training fold, and estimate power over independent mouse cohorts. Four and five mice cannot determine a universal minimum of ten; any $n \geq 10$ rule is only a planning hypothesis until such design-specific curves or an external cohort are available. Detection and conditional-abundance endpoints should be defined separately before inspecting future outcomes. No such prospective power curve or dual endpoint is claimed here.

5.2.2 Intervention-window diagnostic on previously viewed mice

Both source cohorts have the same intervention timetable: high-fat diet from days 21.5 through 28.5, vancomycin from 35.5 through 42.5, and gentamicin from 50.5 through 57.5. Therefore the earlier hypothesis that the UC-versus-healthy contrast separates perturbed and stable cohorts is false at the design level. We instead predeclared seven windows around those periods and compared the same two predictors within each cohort. Figure 3 time indices were mapped mechanically to the public post-preprocessing sample schedule, not aligned against held-out truth. The training-population baseline was fit on the other mice for each held-out animal; published MDSINE2 without modules predictions supplied the comparator. Full-period medians and pair counts reproduce the

archived head-to-head outputs for both scores before slicing windows. This diagnostic was designed after seeing aggregate full-period results, so it cannot become an untouched confirmation.

The healthy cohort has 4 held-out mice and 141 taxon labels. A window/taxon contributes if it has at least one scored timepoint; the detected-only score further requires true abundance above $1e-5$. Each mouse's median taxon-level error gap (population baseline minus MDSINE2-NM) is calculated before averaging across mice. Negative numbers favor the baseline. Pair and timepoint counts below are within-mouse repeated measurements, not independent biological units.

Table 12. Predeclared intervention-window diagnostic on viewed healthy mice. Scores reuse published Figure 3 predictions and truth; windows follow source intervention days. No treatment effect is identified.

window	score	mice	pairs	timepoints	baseline-favorable mice	mean mouse gap
pre	all	4	564	10716	4/4	-0.311
diet	all	4	564	4794	4/4	-0.195
between diet vanco	all	4	564	5076	4/4	-0.104
vancomycin	all	4	564	4935	4/4	-0.309
between vanco gent	all	4	564	5076	4/4	-0.122
gentamicin	all	4	564	5076	4/4	-0.163
after gent	all	4	564	6204	4/4	-0.059
pre	detected	4	514	5858	4/4	-0.254
diet	detected	4	524	3412	4/4	-0.075
between diet vanco	detected	4	486	3329	3/4	-0.036
vancomycin	detected	4	446	2143	2/4	-0.002
between vanco gent	detected	4	340	1874	4/4	-0.128
gentamicin	detected	4	390	2508	3/4	-0.118
after gent	detected	4	399	3218	0/4	0.025

The uc cohort has 5 held-out mice and 121 taxon labels. A window/taxon contributes if it has at least one scored timepoint; the detected-only score further requires true abundance above $1e-5$. Each mouse's median taxon-level error gap (population baseline minus MDSINE2-NM) is calculated before averaging across mice. Negative numbers favor the baseline. Pair and timepoint counts below are within-mouse repeated measurements, not independent biological units.

Table 13. Predeclared intervention-window diagnostic on viewed uc mice. Scores reuse published Figure 3 predictions and truth; windows follow source intervention days. No treatment effect is identified.

window	score	mice	pairs	timepoints	baseline-favorable mice	mean mouse gap
pre	all	5	605	11495	5/5	-0.220
diet	all	5	605	5203	5/5	-0.128
between diet vanco	all	5	605	5445	4/5	-0.061
vancomycin	all	5	605	5445	5/5	-0.205
between vanco gent	all	5	605	5445	5/5	-0.092
gentamicin	all	5	605	5324	5/5	-0.105
after gent	all	5	605	6655	5/5	-0.093
pre	detected	5	547	7346	5/5	-0.174
diet	detected	5	571	3713	5/5	-0.089
between diet vanco	detected	5	524	3846	3/5	-0.018
vancomycin	detected	5	501	2592	3/5	-0.023
between vanco gent	detected	5	388	2186	5/5	-0.263
gentamicin	detected	5	426	2992	5/5	-0.154
after gent	detected	5	436	3888	4/5	-0.037

The all-timepoint error gap favors the simple baseline in all seven windows for all four healthy mice. In UC it favors the baseline in five of five mice in six windows, and four of five between diet and vancomycin. This is a pattern of predictive error, not a causal antibiotic response: the before/after windows move along one trajectory, the models may differ in extrapolation drift, and the data were viewed during method development. The detected-only endpoint is not uniformly baseline-favorable: after gentamicin in the healthy cohort its mean mouse-level gap is +0.025 with zero of four mice favoring the baseline. A legitimate test of whether interactions help because of a perturbation would require randomized intervention holdout, matching a no-change arm and freezing taxa, metrics and intervention contrasts before an unseen cohort is opened.

5.2.3 Why a detection-plus-abundance score needs a compatible model output

A two-part microbiome forecast score would ideally measure both whether a taxon is detected and, conditional on detection, how far its abundance forecast misses. The first component needs calibrated probabilities or a threshold chosen using training-only data in a common absolute-abundance unit. The second can use the source paper's detected-truth log10 RMSE. To show why the distinction matters without claiming a new benchmark, we predeclared a diagnostic on the same viewed Figure 3 data: apply the truth definition (absolute abundance >1e-5) to both models' raw absolute predictions and count timepoint-level true/false detections, then compute conditional RMSE without changing the original prediction arrays. This is not a score optimized for either method.

Table 14. Exploratory decomposition on viewed mouse cohorts.

Sensitivity/specificity apply truth's $>1e-5$ cutoff to continuous forecasts; they are NOT a fair calibrated detection leaderboard. Conditional RMSE medians reproduce the archived official paired metric.

cohort	method	mice	truth-negative times	sensitivity	specificity	positive forecast rate	conditional RMSE median
healthy	population baseline	4	19535	0.951	0.631	0.679	0.919
healthy	MDSINE2 no modules	4	19535	1.000	0.000	1.000	0.913
uc	population baseline	5	18449	0.971	0.581	0.745	0.692
uc	MDSINE2 no modules	5	18449	1.000	0.000	1.000	0.805

At this cutoff the released MDSINE2-without-modules continuous abundance forecasts are positive at almost all 41,877 healthy and 45,012 UC taxon-timepoints. Only eight healthy and five UC truth-negative timepoints are forecast below the cutoff, giving specificity near zero; sensitivity is 100%. The training-population baseline scores more true negatives but misses some true positives. It would be misleading to advertise the resulting balanced-accuracy gap as a new win: MDSINE2-NM did not expose a calibrated detection classifier, the truth cutoff is not a prefit model-output threshold, taxon-timepoints are repeated within four or five mice, and both source cohorts were already inspected. Conditional-abundance pair medians agree exactly with the archived official result: healthy 0.919 baseline versus 0.913 comparator, UC 0.692 versus 0.805. A proper future two-part benchmark would freeze a train-only detection calibrator and evaluate detection and abundance separately at source-family and mouse levels on a truly untouched cohort; no such evaluation exists here.

5.2.3 Mouse-level two-part sensitivity

Pooling tens of thousands of taxon-timepoints is useful for checking a forecast-output threshold, but it hides the actual sample size of four healthy-donor and five UC-donor mice. We therefore applied the same fixed threshold separately to each held-out mouse, with no new threshold fit. The observed detection statuses, prediction construction, and detected-only abundance errors are unchanged; per-mouse confusion matrices sum exactly to the pooled counts above and per-mouse conditional RMSE matches the earlier calculation. The table reports equal-weight mouse medians and median within-mouse baseline-minus-MDSINE2-no-modules differences, not confidence intervals over pseudo-replicated timepoints.

Table 15. Exploratory viewed-cohort per-mouse decomposition. Difference is baseline minus MDSINE2 without modules; lower conditional RMSE is better. Detection cutoff applied directly to continuous forecasts is not a calibrated detector benchmark.

cohort	endpoint	mice	baseline median	comparator median	paired median difference	baseline higher/lower mice
healthy	sensitivity	4	0.951	1.000	-0.049	0/4
healthy	specificity	4	0.639	0.000	0.639	4/0
healthy	balanced acc.	4	0.795	0.500	0.295	4/0
healthy	positive rate	4	0.679	1.000	-0.321	0/4
healthy	conditional RMSE	4	0.923	0.917	-0.027	1/3
uc	sensitivity	5	0.973	1.000	-0.027	0/5
uc	specificity	5	0.578	0.000	0.578	5/0
uc	balanced acc.	5	0.775	0.500	0.275	5/0
uc	positive rate	5	0.746	1.000	-0.253	0/5
uc	conditional RMSE	5	0.712	0.813	-0.130	0/5

Every healthy mouse has baseline specificity greater than 0.595 and the continuous MDSINE2-no-modules output has zero or near-zero specificity at the raw truth cutoff. The pattern repeats in every UC mouse, with baseline specificity 0.572-0.590 and comparator specificity 0-0.0011. This consistency across mice prevents one unusually sampled subject from explaining the pooled output-threshold incompatibility, but it does not create a fair detection leaderboard: the continuous comparator was not trained or calibrated to issue binary calls at this cutoff. Predicting almost all timepoints as positive also gives it perfect sensitivity by construction. The abundance story is separate. The baseline has lower conditional RMSE in three of four healthy mice yet a slightly worse pooled taxon-pair median, illustrating how aggregation and the biological unit alter the description without changing any forecast. It has lower conditional RMSE in all five UC mice. Neither within-cohort pattern establishes source transfer, clinical detection accuracy, or the response to diet or antibiotics. A prospective head-to-head would define a train-only detection score and threshold for each method, preserve mouse or donor grouping, and test the frozen pipelines on a disjoint source family.

5.3 What predicts where interactions help?

Across the 160 studies we regressed the audit gain (prior error minus interaction error) on standardised study covariates: Shannon diversity and Bray-Curtis dispersion (scikit-bio), co-occurrence network density and modularity (networkx; edges where |Spearman rho| > 0.3 among genera with prevalence >= 20%), log sample and taxa counts, and assembly and human-gut flags. Inference uses heteroskedasticity-robust (HC3) standard errors:

$$gain_s = b_0 + \sum_k b_k z_{sk} + e_s, \quad Var(b) = \frac{(Z'Z)^{-1} Z' diag(e_s^2 / (1 - h_{ss})^2)}{Z (Z'Z)^{-1}} \quad (15)$$

Table 16. Predictors of interaction gain (OLS $R^2 = 0.47$; random forest 5-fold CV $R^2 = 0.37$; results/predictors_audit.json).

covariate	std. beta	p (HC3)	RF permutation importance	marginal Spearman rho
shannon	-0.003	0.63	0.039	-0.350
beta_bc	0.014	0.0051	0.268	0.510
net_density	0.026	0.00076	0.600	0.610
net_modularity	-0.001	0.89	0.022	-0.420
n	0.012	0.034	0.089	0.240
taxa	0.017	0.0061	0.127	0.540
assembly	-0.002	0.7	0.022	-0.550
human_gut	-0.007	0.24	0.001	-0.460

Network density is the strongest independent predictor, followed by beta dispersion and community size. Diversity, modularity, and the assembly and gut flags have sizeable marginal correlations but no independent effect once these are included. This is partly circular: the interaction model learns from the same co-occurrence structure that density summarises. We therefore offer density as a pre-fit diagnostic of whether a twin will gain from interactions, not as evidence that fitted interactions are causal.

5.3.1 Local model capacity and budget inventory

The critique rightly asks whether the local model ranking reflects ecological information or a capacity and optimization mismatch. To make that question testable, we counted trainable scalar weights in the exact local implementations at each of the six already viewed cNODE dataset dimensions and paired the counts with their archived ten-fold sample-error vectors. The closed-form presence mean counts one fitted mean per taxon, while cNODE has a full n -by- n matrix and the gLV implementation has that matrix plus n intrinsic rates. GraphTwin's default two-layer, 32-dimensional architecture has roughly fourteen to sixteen thousand trainable scalars. The table is an inventory of available comparisons, not a matched-parameter or matched-FLOP experiment. It preserves the archived same-fold vectors rather than rerunning or choosing a width after seeing the test errors.

Table 17. Local implementation and archived same-fold error inventory, on six viewed cNODE ecosystem tables. Counts omit non-trainable buffers; zero epochs means a closed-form prior. Error medians are local ten-fold values, not the published cNODE LOO result.

dataset	taxa	model	fitted scalars	epochs	10-fold BC median	better/worse samples vs prior
Drosophila	5	prior	5	0	0.172	0/0
Drosophila	5	cNODE	25	200	0.104	16/7
Drosophila	5	gLV	30	200	0.052	17/6
Drosophila	5	GraphTwin	13761	150	0.093	15/8
Soil vitro	8	prior	8	0	0.205	0/0
Soil vitro	8	cNODE	64	200	0.157	22/21
Soil vitro	8	gLV	72	200	0.117	21/22

dataset	taxa	model	fitted scalars	epochs	10-fold BC median	better/worse samples vs prior
Soil vitro	8	GraphTwin	13857	150	0.140	21/22
Oral	73	prior	73	0	0.204	0/0
Oral	73	cNODE	5329	200	0.245	21/122
Oral	73	gLV	5402	200	0.216	69/74
Oral	73	GraphTwin	15937	150	0.202	75/68
Gut	58	prior	58	0	0.256	0/0
Gut	58	cNODE	3364	200	0.275	31/75
Gut	58	gLV	3422	200	0.244	58/48
Gut	58	GraphTwin	15457	150	0.287	49/57
Ocean	73	prior	73	0	0.093	0/0
Ocean	73	cNODE	5329	200	0.132	51/217
Ocean	73	gLV	5402	200	0.070	180/88
Ocean	73	GraphTwin	15937	150	0.071	158/110
Soil vivo	36	prior	36	0	0.124	0/0
Soil vivo	36	cNODE	1296	200	0.123	314/364
Soil vivo	36	gLV	1332	200	0.106	466/212
Soil vivo	36	GraphTwin	14753	150	0.119	407/271

Increasing scalar count does not order the six local scores. In human gut, the 58-parameter population prior has median BC 0.256, compared with cNODE's 3,364 parameters at 0.275, gLV's 3,422 at 0.244 and GraphTwin's 15,457 at 0.287. On *Drosophila* gut, the 30-parameter gLV has median 0.052 versus 13,761-parameter GraphTwin at 0.093, though the previously reported gLV ten-fold apparent beat over published cNODE was retracted after a protocol-matched leave-one-out check. On Ocean, gLV and GraphTwin both have lower local ten-fold medians than the 73-parameter prior, 0.070 and 0.071 versus 0.093. These observations show no simple monotone model-size ranking; they cannot tell whether larger networks lack signal, are over- or underfit, or optimize a mismatched objective. cNODE and gLV use 200 epochs at learning rate 0.01, while GraphTwin uses 150 at 0.005, and different forward passes make even equal epochs unequal compute. Optimizer convergence, seed variance, FLOPs and wall-time by fold were not pinned in the old result JSONs. A credible capacity pivot must preselect widths and depths, parameter-matched variants within architecture, identical frozen folds, several training seeds and train-only early stopping, and then report model failure and compute along with error on an independent source. This inventory answers what was fit locally, not why one class wins.

5.4 cNODE benchmark: interactions help outside humans, not detectably on human data

Table 18. 10-fold CV median Bray-Curtis on four cNODE datasets. Published cNODE uses LOO, so this table is indicative only.

dataset	n	taxa	null	cNODE (ours)	gLV (ours)	GraphTwin (ours)	cNODE published
Drosophila Gut	24	5	0.172	0.104	0.052	0.093	0.066
Soil Vitro	48	8	0.205	0.157	0.117	0.140	0.079

dataset	n	taxa	null	cNODE (ours)	gLV (ours)	GraphTwinn (ours)	cNODE published
Human Oral	143	73	0.204	0.245	0.216	0.202	0.211
Human Gut	106	58	0.256	0.275	0.244	0.287	0.242

Table 19. Exact leave-one-out null vs published cNODE.

dataset	n	null LOO median	95% CI	published cNODE	cNODE inside null CI?
Ocean	269	0.094	[0.089, 0.101]	0.060	no - cNODE better
Drosophila gut	24	0.170	[0.100, 0.193]	0.066	no - cNODE better
Soil in vitro	48	0.198	[0.085, 0.268]	0.079	no - cNODE better
Soil in vivo	678	0.124	[0.119, 0.127]	0.107	no - cNODE better
Human oral	143	0.204	[0.188, 0.222]	0.211	YES
Human gut	106	0.259	[0.236, 0.271]	0.242	YES




Figure 3. Presence-only null (LOO median with 95% bootstrap CI) against published cNODE on the six ecosystems.

Interpretation. On non-human systems cNODE's interactions buy a large, clear gain. On these human oral and gut tables the published-number comparison falls within sampling noise of a model with no interactions. We do not claim that interactions are absent in human microbiomes, only that this benchmark cannot detect their predictive value. Our GraphTwin, a larger model, does no better (human gut 0.287 under 10-fold CV), which does not distinguish limited signal from model-fitting or capacity limitations.

5.4.1 Internal split-calibrated error radii, not a reliability index

A local research predictor can state what its uncertainty check actually measured without turning it into a clinical or cross-source promise. We fitted the same PresenceMean composition prior separately on 60% of each of six previously viewed cNODE ecosystem tables, calibrated on floor 20%, and checked the remaining samples after assemblage deduplication under one seed. For each labeled calibration sample we computed Bray-Curtis error between the fixed predictor and observed composition. At target miscoverage $\alpha=0.1$, the radius is the $\text{ceil}((n_{\text{cal}}+1)(1-\alpha))$ -th sorted calibration error; if that rank exceeds n_{cal} , it returns the vacuous Bray-Curtis maximum 1. This is an error-radius check for a composition predictor, not a set of intervals for individual taxa.

Table 20. One 60/20/20 within-table split for the PresenceMean baseline, seed zero. Empirical test coverage is descriptive; nominal 90% marginal coverage requires exchangeability and does not hold deterministically in each finite test set.

ecosystem	train	cal	test	radius	covered	empirical coverage	vacuous
Ocean	161	53	55	0.223	52/55	94.5%	no
Drosophila Gut	14	4	6	1.000	6/6	100.0%	yes
Soil Vitro	28	9	11	0.441	9/11	81.8%	no
Soil Vivo	406	135	137	0.226	122/137	89.1%	no
Human Oral	85	28	30	0.407	29/30	96.7%	no
Human Gut	63	21	22	0.430	19/22	86.4%	no

The Ocean test covered 52/55 observations at radius 0.223; Human Oral covered 29/30 at radius 0.407. Soil in vitro covered 9/11, Soil in vivo 122/137 and Human Gut 19/22, all below 90% as realized proportions. Drosophila Gut covered 6/6 only because four calibration samples forced the vacuous radius 1.0; it provides no useful discrimination. Small finite test proportions are not direct proof of a failure of marginal conformal theory, and a radius calibrated on one table need not transfer to a new biome, assay, participant or perturbation. The source tables were already viewed and their split is within-dataset, not a held-out biological source. No Twin Reliability Index is calibrated or reported. Source SHA-256 values, splits, quantile ranks and exact result fractions are in `results/internal_conformal_check.json`; implementation: `scripts/internal_conformal_check.py` and `src/microtwin/conformal.py`.

5.4.2 End-to-end predictor check after input-only abstention

The preceding radius calculation checked the frozen model directly. A separate already-run end-to-end check fed six previously viewed tables into the released local calibrated-prediction path with the same seeded candidate split. The command refuses a calibration or query sample with a present taxon absent from the training vocabulary. The check therefore removed such rows using only their input taxon presence, before measuring errors, then reported coverage on the admissible test rows. This conditional denominator

differs from the all-candidate split in Section 5.4.1: it cannot be presented as unconditional 90% coverage or as successful prediction of abstained rows.

Table 21. Actual local calibrated prediction path on six already-viewed cNODE tables. Rows with unseen present taxa were excluded by input-only abstention before scoring; split seed zero. This is an admissible-row sensitivity, not an external validation result.

ecosystem	cal kept/candidate	test kept/candidate	radius	covered/scored	scored coverage	vacuous
Ocean	52/53	55/55	0.193	51/55	92.7%	no
Drosophila Gut	4/4	6/6	1.000	6/6	100.0%	yes
Soil Vitro	9/9	11/11	0.441	9/11	81.8%	no
Soil Vivo	132/135	135/137	0.225	120/135	88.9%	no
Human Oral	26/28	28/30	0.407	27/28	96.4%	no
Human Gut	21/21	22/22	0.430	19/22	86.4%	no

Ocean loses one of 53 calibration candidates and covers 51/55 scored tests at radius 0.193. Soil in vivo loses three calibration and two test candidates, then covers 120/135 at radius 0.225. Human Oral loses two calibration and two test rows, covering 27/28 at radius 0.407. Soil in vitro still covers 9/11, Human Gut 19/22, and Drosophila Gut 6/6 with the same vacuous radius 1.0. For every scored table, the command's predictions match the archived PresenceMean implementation to better than $2e-15$ absolute maximum deviation; this is a numerical implementation check, not independent ecological replication. No experiment here verifies sample exchangeability across participants or source families, and no Twin Reliability Index or clinical guarantee follows. Candidate and scored counts, source hashes and per-table numerical deviations are in `results/local_calibrated_six_dataset_check.json`; runnable check: `scripts/local_calibrated_six_dataset_check.py`.

Measured extension: first-period stool-metabolite forecasting

A personalized research platform needs a task with measurable targets, known sample units and an honest boundary around what its model predicts. The next measured extension therefore uses targeted stool short-chain fatty acid (SCFA) concentrations rather than borrowing an unverified taxonomic table or treating predicted gene functions as observed biology. The public 2026 yogurt and rolled-oats crossover study supplies an explicitly CC BY 4.0 metabolite deposit, Zenodo record 15363886. Its eight analytes are 2-methylbutyric acid, acetic acid, butyric acid, hexanoic acid, isobutyric acid, isovaleric acid, propionic acid and valeric acid, measured in micrograms per gram of stool. This is a different measurement contract from relative microbial abundance: concentrations are not normalized to sum to one, and predicting them does not certify taxonomic composition, causal ecological interactions or clinical benefit.

The original trial randomized 119 adults; 110 completed all four phases. Each participant consumed 250 grams of yogurt daily in one four-week intervention and the same yogurt plus 50 grams of rolled oats daily in another. Treatment order was randomized, with four-week washout phases. The deposit contains 462 samples, of which 440 are main samples and 22 are quality-control duplicates. All main samples form unique submitted subject-time coordinates across four phases. The analysis deliberately uses only the first baseline and first post-intervention sample. It does not learn from the second intervention, pool repeated periods as independent people or assume that washout eliminated carryover. Group A receives yogurt first and group B receives yogurt plus oats first. The model sees that intervention label, but the task remains a forecast in observed treated participants, not an estimate of what the same person would have experienced without treatment.

Metadata inspection preceded any concentration scoring. This order found an important difference between nominal phase labels and actual collection dates. First-period intervals range from 24 to 138 days in the deposited metadata, while later washout intervals include two nonpositive values and some date endpoints do not parse. These entries were not silently repaired. The frozen first-period task retains only baseline-to-follow-up intervals of 21 through 35 days, inclusive. This bounds collection separation without pretending that dietary start, adherence or exact exposure duration were measured by the sample label. The result cannot support a continuous-time kinetic model or a clinical scheduling recommendation. Public scripts elsewhere in the author repository include illustrative participant clinical records; those records are neither model features nor public outputs of this analysis.

The concentration workbook has 3,696 unique sample-analyte records, exactly eight for each of its 462 sample labels, and those labels match the metabolite metadata. Units and material are uniform. The local reproduction script verifies the original workbook MD5 and metadata SHA-256 before processing. Primary assay eligibility is strict: all baseline and target cells must be finite nonnegative numbers, within their stated lower and upper quantitation limits, and free of nonempty analysis comments. The laboratory report says that values outside calibration range may be extrapolated and flagged BLOQ or ALOQ. The analysis does not convert those flags to zero, remove an inconvenient metabolite, or impute a favorable response. Complete-case filtering leaves 84 subjects, 42 in each treatment-order group. Twenty-five subjects fail the assay rule and one fails the collection interval. This restriction is material: the estimates apply to quantifiable complete cases, not all randomized participants or all low-concentration specimens.

A frozen five-fold subject split holds out entire submitted people. Within each arm, sorted subject labels are shuffled with a fixed random seed and assigned cyclically to folds. Log1p baseline concentrations are standardized using training-fold statistics. A ridge model with fixed alpha 10 predicts change in log concentration from eight baseline features, the assigned first intervention, and eight baseline-by-intervention products.

Predicted target log concentrations are clipped only below zero; there is no compositional projection. No hyperparameter search, effect-based feature selection or alternate model variant follows the score. The comparators are persistence, training-only same-arm mean log change, and the mean change of the three nearest same-arm baseline donors. Neighbors are selected in training-standardized baseline space. This is a demanding practical comparison because a personalized predictor must add value over carrying the last measurement forward or borrowing an observed donor response.

The primary error averages absolute log concentration errors across analytes after scaling by each target-log training-fold standard deviation, then averages over held-out subjects. Scaling prevents a large natural range in one analyte from dominating every other target; it does not create a universal reliability score. Secondary results retain unscaled log errors and per-analyte errors. The useful-win rule was frozen before outcome inspection: at least ten percent lower primary error than each comparator, with a negative upper bound on each paired subject-bootstrap interval. The bootstrap uses 2,000 subject resamples and is descriptive within this study. It does not establish independent-source replication, correct for all previous development choices or supply calibrated patient-level prediction intervals.

Table 22. First-period SCFA held-subject development forecasts, 84 complete-case subjects; smaller error is better.

Forecast	Standardized log error	Unscaled log error
ridge	0.703	0.447
persistence	0.799	0.495
arm_mean	0.807	0.502
nearest_three	0.776	0.497

The fixed ridge forecast has primary mean error 0.702517, compared with 0.798983 for persistence, 0.806896 for same-arm mean change and 0.775684 for nearest-three donors. Relative improvements are 12.07, 12.94 and 9.43 percent. The corresponding paired model-minus-reference bootstrap intervals are [-0.169426, -0.022469], [-0.172099, -0.034693] and [-0.123910, -0.020608]. Thus the model improves the aggregate metric against all three fixed baselines in this development evaluation, but the third improvement falls below the required ten percent. The recorded verdict is `useful_win=False`. This distinction matters: a favorable direction of error change does not satisfy a preregistered usefulness rule. The threshold is not lowered to rescue the result, and a close miss is not promoted to a win against a leading external tool.

Aggregate accuracy also hides target-specific weaknesses. Hexanoic acid has mean absolute log error 0.690508 under ridge, worse than persistence at 0.601259. A user should not infer that every metabolite benefits from the aggregate improvement. All eligible subjects are scored once, predictions are finite, and replacing held-out targets leaves their predictions unchanged. The mandatory replacement check refits the five folds for verification, giving ten small fits in total: five scoring fits and five verification fits.

This exceeds the brief five-fit execution estimate but stays inside the frozen single-thread, 120-second evaluation limit; measured evaluation time is 0.287 seconds. No extra scientific variant or paid compute was used. These checks establish local execution properties, not true biological identity or generalization to a new laboratory.

Source admission is part of the measurement, not a dataset count

Several apparently suitable antibiotic cohorts did not supply the full evidence needed for a comparable task. CEREMI has multiple publications and accessions from the same trial, not independent replication. Hagan has a well-defined antibiotic course and a primary deposit, but this bounded inventory inspection did not retrieve a taxonomic table with an exact phase, subject and collection-time join. Raymond has clear cefprozil dosing and untreated controls, yet its inspected supplements contain quality-control and statistical tables rather than an identified complete abundance matrix. Its final sample label denotes ninety days after treatment end, so treating that label as ninety days after first dose would alter the time axis. Zaura provides an important further warning: the primary study selected its shotgun pairs for the largest observed microbiome changes. Such a selected subset cannot be an unbiased multimodal validation sample merely because two assays exist.

Palleja illustrates why processed-data availability alone is insufficient. Its author archive is retrievable and its primary sample table has 57 specimens from twelve donors. An educational mirror has 57 abundance-table headers, but eight primary sample IDs carry an opt suffix absent from the mirror. The analysis does not strip the suffix and assert provenance. The earlier clinical report also describes stool collected before named visits and an additional antibiotic course in one participant. The bounded author archive contains more than one gigabyte of uncompressed tables and no README, license or sample-map member in its inspected inventory. The approved extraction stop was honored. This is not proof that the source is unusable forever; it is a precisely described admission gap. A later MISO diet deposit similarly uses an other-open label without explicit reuse terms in the inspected archive. Its seven-day homogeneous diet and measured outputs cannot be substituted for four-week yogurt/oat SCFA validation without a separate task contract.

The licensed SCFA route therefore adds a small, reproducible quantitative test and exposes a modest predictive signal, while leaving the larger scientific promises open. Its concentrations are now viewed development data, not an untouched holdout. The taxonomic GitHub tables stay metadata-only because the article license and metabolite-deposit license do not automatically establish their own reuse scope. Neither the source hunt nor the SCFA forecast closes independent top-tool benchmarking, causal perturbation validation, replicated keystone discovery, foundation-scale training or external calibration. The appropriate next scientific question is whether a locked

transferable task works in an independently recruited, measurement-compatible and explicitly reusable cohort. If those requirements fail, the platform should abstain rather than accumulate nominal datasets or relabel a same-study result as discovery.

Operational contract for a local research run

The current platform is a set of local research programs, not a hosted patient service. Its practical contribution is to make a limited prediction reproducible while keeping questions that it cannot answer separate. A researcher first decides whether the target is composition conditional on an observed assemblage, a future measured metabolite, or a response to an intervention. These targets need different input and validation contracts. A species-removal screen changes the assemblage supplied to a composition predictor, but does not thereby identify the composition that would occur after removing that organism in a person. A model can calculate a number for the changed input even when the biological contrast is unidentified. The workflow must stop that numerical operation from becoming an unsupported treatment recommendation.

For the implemented composition baseline, each labeled training row represents a submitted sample and each column a taxon. Counts are normalized within each row before the population mean is computed, so a deeply sequenced sample does not receive more weight merely because it has more reads. The query supplies binary presence indicators, not its unknown outcome abundances. The prediction masks the population mean to the taxa present in the query and renormalizes their weights. A small smoothing term is applied only after the unseen-present-taxon check. Consequently, smoothing cannot manufacture evidence for an organism absent from all training samples. This rule is intentionally simple. Its behavior is inspectable, but its simplicity is not proof that it is the best predictor or that a population average captures a particular person.

The table contract is strict at the local file boundary. CSV or TSV files are limited to fifty million bytes each. The first column must be `sample_id`, remaining column names must be unique and nonempty, and every numeric entry must be finite. Training counts must be integers; submitted relative abundances must sum to one within the implemented tolerance. Query taxa must match both training names and column order exactly, query identifiers must not overlap training identifiers, and every query must contain at least one present taxon. A present taxon with zero empirical training abundance leads to abstention. These restrictions reduce silent numerical errors and accidental leakage. They do not detect every renamed duplicate, identify a laboratory method, establish a biological taxon equivalence, or recover an omitted organism.

The strict prediction route adds a measurement contract before writing an output. The contract supplies source family, specimen material, assay, taxonomy version, processing pipeline and unit for training and query. Declared measurement fields must match, the source family must match the requested source label, and cross-source prediction abstains

because an independently validated bridge is not implemented. Matching declarations establish consistency between submitted labels, not their truth. The researcher must retain the source documents behind the declarations. A copied assay name or a shared genus label cannot turn a V4 amplicon table into a shotgun table, make saliva interchangeable with feces, or establish that two processing pipelines have equivalent detection limits.

Subject maps address a second boundary that sample identifiers alone cannot protect. The checked route can require exact sample-to-submitted-subject maps for training and query, with one mapping for every sample. Overlap between submitted subjects then fails even if their sample names differ. When the query map is absent, the report explicitly leaves query-subject separation unverified. This distinction matters in longitudinal studies, where many valid sample rows can come from one person. Renaming that person across files still defeats an exact-label check. Source recruitment records, family relationships and specimen provenance therefore remain part of the scientific admission process, rather than claims supplied by a successful software return code.

Calibration, integrity and the meaning of an abstention

A calibrated local run adds a separate labeled calibration table after the predictor has been defined from training data. The reported Bray-Curtis radius uses the finite-sample rank ceiling of $(n_{\text{cal}}+1)(1-\alpha)$. If that rank exceeds the available calibration count, the radius is the maximum distance, one. That output is valid as an explicit lack of useful precision, not a reason to reduce the nominal coverage target after seeing the result. The radius concerns marginal error under exchangeability of calibration and future queries. It is not a conditional interval for a named individual, a probability of clinical harm, or a guarantee for a new intervention or laboratory. The observed same-dataset coverage diagnostics earlier in this paper must be read with their precise candidate, excluded and scored denominators.

The reliability ledger preserves five separate evidence components: prediction accuracy, predictive interaction dependence, uncertainty calibration, transferability and robustness. Even when a caller supplies all five numbers, the implementation does not average them into a score. It records them as reported but unvalidated and leaves the composite index unavailable. This is a deliberate separation between an engineering interface and a scientific result. A high number for one component cannot replace missing transfer evidence, and arbitrary weights would hide a decision that has never been calibrated on independent sources. The correct output for the present platform is the evidence vector and its missing fields, not a patient-facing percentage of trust.

A checked composition run can write an integrity receipt alongside its new prediction CSV. The receipt binds hashes and byte counts of the training file, query, measurement contract, prediction and any included subject maps. Input hashes are rechecked during the run to detect changes between inspection and prediction. Existing output and receipt

paths are refused. If receipt creation fails after prediction creation, the implementation removes that new prediction output rather than leave an apparently complete paired deliverable. This is local failure handling, not a transaction guarantee across machines, a cryptographic signature, or authentication of the researcher. Replacing both an artifact and its receipt can defeat a plain hash comparison. The receipt also does not embed the input tables, although its run report may contain local paths and source labels and should be treated as private working material.

The workflow therefore separates refusal reasons from scientific negatives. A malformed query is an input failure; a missing source bridge is an admission failure; a vacuous calibration radius is an uncertainty limitation; and a frozen model that misses its usefulness threshold is a measured negative. None should be erased by trying nearby settings until a favorable status appears. The researcher records the original reason, checks whether the task itself was correctly defined, and changes the protocol only prospectively. When an intervention target remains unidentified, collecting suitable longitudinal or perturbation evidence is a different action from rerunning a composition baseline. This separation is the operational form of the paper's scientific gates: executable arithmetic is useful, but its authority ends at the measurement and validation contract it actually satisfies.

Implementation anchors are `src/microtwin/local_predict.py`, `src/microtwin/cli.py`, `src/microtwin/local_calibrated_predict.py`, `src/microtwin/reliability.py` and `src/microtwin/run_bundle.py`. These anchors describe the current local research workflow, not a deployed access-controlled service. Assay declarations, caller authorization flags and submitted pseudonyms still require source-grounded verification outside the numerical routine. This operational section adds no benchmark outcome and does not close external prediction, calibrated composite reliability, biological intervention validation or replicated discovery.

Measurement hierarchy and the unit of a forecast

A microbiome dataset has several nested units that must remain distinct: an analytical run, a biological specimen, a collection occasion, a participant, and a recruited source family. A matrix column is not automatically the unit of inference. Two sequencing runs can measure one specimen; two aliquots can share one collection; several stools can belong to one person; and two deposited studies can reuse one recruitment. The platform treats the reconstruction of this hierarchy as part of defining the task, before fitting or splitting. A method intended to generalize to a new person must hold that person out, whereas a method intended to transfer to a new recruitment or laboratory needs a source-level boundary. More rows improve measurement of the units already sampled but do not automatically add independent biological replication.

Compatible raw sequencing runs of one specimen can be summed before composition normalization. The implemented run-collapse routine checks the supplied run-to-sample relationships, requires all run columns to be explained, refuses a run assigned to multiple samples, and excludes identified assembly columns. It also rejects an identifier that is simultaneously a run and an assembly. These checks prevent a pooled assembly from being counted again as an independent measurement beside the reads that generated it. Compatibility of the assays, pipelines and raw count interpretation remains a source requirement. The routine does not verify that requirement and must not be used to sum unrelated abundance units merely because their columns share names.

A numerical example shows why the order of aggregation matters. Suppose two compatible runs from one hypothetical specimen have counts (90,10) and (1,9) for taxa A and B. Pooling raw counts gives (91,19), with relative abundance of A equal to $91/110$, approximately 0.8273. Averaging the two separately normalized runs instead gives A a weight of $(0.9+0.1)/2$, or 0.5. The second calculation weights the ten-read run equally with the hundred-read run. Neither arithmetic rule can be selected by which later gives a better model score: the source measurement contract determines whether compatible raw reads are being pooled or two genuinely separate measurements are being averaged. These are hypothetical counts, not a new empirical result.

Technical concentration replicates require a different decision. Repeated injections from one extract measure analytical variation, not repeated biological exposure. Summing their concentrations would usually change the unit and cannot borrow the raw-read pooling rule. A prospectively specified summary, duplicate-quality criterion and quantitation treatment are needed before learning from such measurements. Distinct stools from the same day are also not automatically technical replicates: they can be separate specimens with real within-day variation. The source must identify which level was repeated. The SCFA development task in this paper excludes labeled quality-control duplicates rather than treating them as additional people. The independent high-fiber lead illustrates the separate need to distinguish analytical samples from its participant denominator; it has not supplied an admitted quantitative task here.

Repeated collection times remain valuable after this distinction. They can estimate a within-person trajectory, expose transient changes, and support a baseline-to-follow-up endpoint. They should not be shuffled across training and test as if they were unrelated people when the target is a new participant. A participant with densely collected stools also should not dominate a population endpoint merely through a larger row count unless that row-weighted target was deliberately chosen. One defensible design scores a fixed participant-level summary and then weights participants equally. Another targets a randomly chosen collection occasion and weights occasions. These answer different questions. A paper must state which population its average describes instead of allowing file layout to choose the estimand.

Exposure windows, source boundaries and paired comparisons

Time labels need the same discipline. A visit label, a collection date, a day since the first dose and a day since the last dose are different coordinates. The local course-design screen takes explicit numerical start and end times, requires a strictly earlier baseline and strictly later follow-up, and selects the nearest eligible samples around the course. Nonnegative timing margins can be added prospectively. A specimen exactly at the start or end does not qualify under those strict inequalities. The routine refuses duplicate submitted source-subject-time coordinates so that technical-replicate resolution is not hidden inside a pairing rule. It never reads outcomes, estimates an antibiotic response, or establishes that the supplied treatment dates are correct.

The screen also checks whether another submitted course overlaps the selected baseline-to-follow-up window. If another course falls between those observations, the default rule blocks attribution to a single isolated exposure. This is a design check rather than an adjustment model: recording an overlap does not estimate its effect or remove confounding. Unknown courses, adherence, concurrent medication, illness and dietary changes can still matter even when the submitted window passes. A nominal recovery phase cannot replace a verified treatment timeline. The Raymond distinction between ninety days after treatment end and ninety days after first dose, and the Palleja collection-versus-visit and extra-course findings, are concrete source-admission reasons to preserve the original temporal meaning rather than quietly rename the axis.

Source grouping addresses dependence above the participant. The lineage guard registers documented accession aliases under a shared source and refuses a train, validation or test split that places one known family in multiple partitions. An unknown accession fails rather than becoming independent by default. A separate exact-mirror screen excludes old accessions but labels all remaining records only as candidates with near-duplicate status unchecked. This asymmetry is intentional: finding a known overlap establishes a problem, while failing to find one does not establish independence. A source-family registry needs positive recruitment, sample and processing evidence. Distinct repository URLs, projects or study titles are identifiers to investigate, not independent experiments to count.

Evaluation aggregation then has to follow the intended boundary. The implemented study-aware comparison computes each submitted study's median candidate error and median comparator error, subtracts them, and averages those study differences with equal study weight. Its bootstrap resamples studies together, not individual rows from within a study. A large source cannot dominate merely because it has many specimens. Nevertheless, supplied study identifiers may still conceal repeated recruitment, related participants or laboratory dependence; the routine marks external-win certification false regardless of its interval. With fewer than two submitted studies it reports insufficient

studies and does not manufacture a study-bootstrap interval. These are properties of a statistic, not permission to call a viewed development study external validation.

Paired losses also do not make every possible contrast equivalent. The difference between two study medians is generally different from the median of samplewise paired differences. Likewise, a mean participant loss can differ from a median taxon loss pooled over participants. These choices are not interchangeable translations of one result. The paired prediction files, sample coverage, outcome mask, weighting rule and comparator selection must all be frozen for the same task. The mouse-level sensitivities earlier in this paper retain the original forecasts and explicitly change the biological aggregation level; they do not rewrite the archived taxon-pair benchmark. A future independent comparison should specify the endpoint and resampling unit before observing whether either aggregation favors the new model.

Implementation anchors are src/microtwin/sample_collapse.py, src/microtwin/course_design.py, src/microtwin/subject_partition.py, src/microtwin/source_family.py, src/microtwin/mirror_guard.py and src/microtwin/leaderboard_stats.py. This method synthesis introduces no fresh fit, cohort outcome or empirical benchmark. It identifies the contracts required to turn specimen rows into a scientifically interpretable forecast and explains why the archived 160-table audit, the viewed mouse results and the first-period SCFA test retain separate scopes. The remaining external-test and perturbation-validation gates require new qualifying evidence, not a larger count of analytical rows.

5.5 Stacking campaign: TwinStack and ConstStack

The base models of Section 5.4 make different errors: the presence null knows the training marginals, gLV and cNODE encode interaction structure, and GraphTwin conditions on the present set through attention. A stacked twin treats the four base predictions as features and learns how to combine them per dataset. Both stackers below were pre-registered (PREREG_graphtwin2.md, PREREG_graphtwin2b.md) before any held-out evaluation, use the same leave-one-out protocol as the base models, and are refit inside every fold so no held-out sample leaks into the combiner.

5.5.1 TwinStack: learned gating

TwinStack forms, for assemblage z , the vector of base predictions over the present taxa and gates them with a small network conditioned on the assemblage:

$$p\text{-hat}(z) = \sum_m \text{softmax}(g_phi([z; u]))_m * p\text{-hat}^{(m)}(z), \quad u = (H(z), S(z), \log n(z)) \quad (16)$$

where $p\text{-hat}^{(m)}$ are the four base predictions restricted to $\text{supp}(z)$ and renormalised, and u carries assemblage-level summaries (shannon entropy, richness, log pool size). The

gate g_{ϕ} is a two-layer perceptron trained by Adam on the Bray-Curtis loss through the mixture, with inner cross-validation on the training folds only.

5.5.2 ConstStack: constant weights with inner-OOF selection

ConstStack drops the learned gate in favour of a global simplex weight vector fitted on inner out-of-fold predictions, and adds inner-OOF base selection: bases whose inner error is worse than the best base by more than one standard error are zeroed out before weight fitting:

$$p\text{-}\hat{z}(z) = \sum_{m \in M^*} w_m p\text{-}\hat{z}^{(m)}(z), \quad w = \operatorname{argmin}_{w \text{ in simplex}} \sum_{s \text{ in OOF}} BC(\sum_m w_m p\text{-}\hat{z}^{(m)}(z_s), p_s) \quad (17)$$

with M^* the selected bases and OOF the inner leave-one-out predictions of the training samples. Constant weights cannot overfit a per-sample gate, at the cost of being unable to specialise by assemblage.

5.5.3 Decision rule

Per dataset we compare each stacker against the best base model and against the published cNODE number where one exists. Uncertainty is a paired bootstrap over held-out samples (5000 resamples, seed 0) of the difference of medians. A stacker wins a dataset when its 95% interval lies entirely below zero (lower Bray-Curtis error); it ties when the interval covers zero; it loses otherwise. Losses are reported, not hidden. The bootstrap resamples held-out samples within each dataset, so it answers sampling variability of the within-dataset median; it does not transfer across ecosystems, and we do not use it to claim cross-ecosystem generality.

Table 23. Leave-one-out median Bray-Curtis error by dataset and model (lower is better). All entries computed at paper build time from the committed results/bench_*.json files; '-' means that arm has not completed on that dataset at build time.

data set	n	taxa	presence null	cNODE (ours)	gLV (ours)	GraphTwin	Twinstack	ConstStack	cNODE2	LightGBM	catb	ridgeclr	xgb	cNODE published
Ocean	269	73	0.093	0.132	0.070	0.071	0.070	0.070	0.075	0.074	0.778	0.069	0.072	0.060
Drosophila_Gut	24	5	0.172	0.104	0.052	0.093	0.087	0.072	0.084	0.293	0.258	0.173	0.077	0.066
Human_Gut	106	58	0.256	0.275	0.244	0.287	0.244	0.248	0.254	0.278	0.424	0.250	0.278	0.242
Human	143	73	0.204	0.245	0.216	0.202	0.202	0.202	0.222	0.206	0.379	0.200	0.205	0.211

data set	n	taxa	presence null	cNODE (ours)	gLV (ours)	GraphTwin	TwinStack	ConstStack	cNODE2	LightGBM	catb	ridgeclr	xgb	cNODE published
_Oral														
Soil_Vitro	48	8	0.205	0.157	0.117	0.140	0.177	0.143	0.114	0.127	0.226	0.175	0.097	0.079
Soil_Vivo	678	36	0.124	0.123	0.106	0.119	0.106	0.106	0.110	0.110	0.333	0.112	0.111	0.107

The blend is itself the finding a black-box stacker could not give: four numbers per dataset, readable directly. Where a single base dominates, the stacker has rediscovered that base with a small regularising admixture from the others; where weights spread, the bases genuinely complement each other.

Table 24. ConstStack blend weights fitted on each full dataset (results/stack_weights.json; inner-OOF selection, 400-step simplex fit, seed 0). '0 (dropped)' means inner-OOF selection excluded that base.

dataset	presence null	cNODE (ours)	gLV (ours)	GraphTwin
Ocean	0.005	0 (dropped)	0.945	0.050
Drosophila_Gut	0.085	0.003	0.814	0.098
Human_Gut	0.007	0 (dropped)	0.782	0.211
Human_Oral	0.195	0 (dropped)	0.005	0.800
Soil_Vitro	0.410	0.293	0.005	0.291
Soil_Vivo	0.002	0 (dropped)	0.996	0.002

Table 25. Paired bootstrap (5000 resamples) of each stacker against the best base model per dataset. Negative differences favour the stacker. WIN/LOSS require the whole interval on one side of zero.

dataset	stacker	best base	stacker median	base median	diff of medians	95% CI	verdict
Ocean	TwinStack	gLV (ours)	0.070	0.070	+0.0000	[-0.0000, +0.0001]	tie
Ocean	ConstStack	gLV (ours)	0.070	0.070	+0.0000	[-0.0006, +0.0003]	tie
Drosophila_Gut	TwinStack	gLV (ours)	0.087	0.052	+0.0348	[-0.0060, +0.0622]	tie
Drosophila_Gut	ConstStack	gLV (ours)	0.072	0.052	+0.0198	[-0.0165, +0.0483]	tie
Human_Gut	TwinStack	gLV (ours)	0.244	0.244	-0.0002	[-0.0295, +0.0175]	tie
Human_Gut	ConstStack	gLV (ours)	0.248	0.244	+0.0037	[-0.0169, +0.0163]	tie
Human_Oral	TwinStack	GraphTwin	0.202	0.202	+0.0001	[-0.0022, +0.0083]	tie
Human_Oral	ConstStack	GraphTwin	0.202	0.202	+0.0000	[-0.0022, +0.0031]	tie
Soil_Vitro	TwinStack	gLV (ours)	0.177	0.117	+0.0600	[+0.0075, +0.1203]	LOSS
Soil_Vitro	ConstStack	gLV (ours)	0.143	0.117	+0.0261	[-0.0374, +0.0780]	tie

dataset	stacker	best base	stacker median	base median	diff of medians	95% CI	verdict
Soil_Vivo	TwinStack	gLV (ours)	0.106	0.106	+0.0000	[-0.0000, +0.0000]	tie
Soil_Vivo	ConstStack	gLV (ours)	0.106	0.106	+0.0000	[-0.0004, +0.0002]	tie

One point the tables make explicit: on Human_Oral both stackers land at 0.2024, below the published cNODE median 0.211 - a point win over the published number. It sits inside the presence-null bootstrap interval (0.204 [0.188, 0.222]), so we claim a published-point win, not a new error record. Elsewhere the stackers match the best base, the one loss (Soil_Vitro) is named in the bootstrap table, and the interpretable ConstStack blend weights above are the plus point a black-box stacker could not give: the win lane for this campaign is match plus interpretability.

5.6 Keystone consensus across studies (exploratory)

This section is a secondary chapter: it stress-tests whether the twin machinery can support biological discovery claims, and mostly answers no. The audit framework of Sections 5.1-5.5 is the primary contribution.

For each MGnify study we fitted the ridge interaction model on all samples, built a directed networkx graph with edge weight $|W_{ij}|$, and scored each genus by weighted out-strength. A genus is 'top' in a study if it is in that study's top 10%. Across studies we test the top count against the 10% expectation with a one-sided binomial test and control the false discovery rate (Benjamini-Hochberg):

$$s_j = \sum_i |W_{ij}|, \quad p_j = P(\text{Binom}(n_j, 0.1) \geq k_j), \quad q_{-}(r) = \min_{\{r' \geq r\}} (m_{p_{-}(r') / r'}) \quad (18)$$

Table 26. Top genera by keystone consensus (248 genera modelled in ≥ 20 studies; results/keystone_consensus.csv).

genus	studies	top	fraction top	p	q (BH)
Cardiobacterium	30	13	0.433	2.3e-06	0.00057
Eikenella	28	11	0.393	4.2e-05	0.0052
Comamonas	51	13	0.255	1.2e-03	0.099
Johnsonella	35	10	0.286	1.7e-03	0.099
Simonsiella	20	7	0.350	2.4e-03	0.099
Desulfobulbus	20	7	0.350	2.4e-03	0.099
Sutterella	56	13	0.232	3.0e-03	0.11

Two genera pass $FDR < 0.05$: Cardiobacterium and Eikenella. Both are oral HACEK genera, so the signal most likely reflects the oral studies, where interaction models fit best. It is not a cross-biome keystone law. Out-strength from a ridge model is a statistical dependence score, not a causal keystone effect, and it scales with each genus's variance and prevalence. No experimental validation was done.

Phylogeny. Resolving each genus to its phylum with NCBI Taxonomy (240 of 248 genera; five eukaryote homonyms such as Bacillus were caught and re-queried within

Bacteria and Archaea), we tested phylum enrichment among the 25 lowest-p genera with one-sided Fisher exact tests and BH correction. Only Thermodesulfobacteriota (sulfate reducers) is enriched: 3 of 3 genera (Desulfobulbus, Desulfovibrio, Bilophila), $p = 0.001$, $q = 0.019$. With three genera, this is a hypothesis (hydrogen and sulfur cross-feeding hubs), not a result.

Literature cross-check. Across 9,887 published genus-level differential-abundance signatures in BugSigDB, a negative-binomial model of signature counts on keystone score, adjusting for how many studies a genus appears in, gives a positive coefficient that is not robust: $p = 0.025$ with dispersion fixed at 1 and $p = 0.062$ with dispersion estimated ($\alpha = 1.16$). Keystone-consensus genera are not clearly over-represented in the disease literature.

$$\log E[n_j] = c_0 + c_1 (-\log_{10} p_j) + c_2 \log(\text{studies}_j), \quad \text{Var}(n_j) = \mu_j + \alpha \mu_j^2 \quad (19)$$

GTDB replication. Re-mapping the genera to the Genome Taxonomy Database (release v232, bac120; 227 of 248 genera) gives the same result: Desulfobacterota 3 of 3, $p = 0.0012$, $q = 0.020$. NCBI and GTDB disagree on the phylum of only 7 genera, mostly renamings. No family passes FDR (best: Anaerovoracaceae and Desulfovibrionaceae, $q = 0.71$). The result does not depend on the taxonomy, but it still rests on three genera.

5.6.1 Physiological traits of keystones: the anaerobe-keystone candidate

We joined each genus to the Madin et al. (2020) trait database (14,893 species; condensed_species_NCBI.csv), aggregating species to genus as the share of strict anaerobes, Gram-negative, motile and sporulating species, and the median genome size, GC content and doubling time. We test each trait against keystone frequency (Spearman, BH across traits) and top-25 membership (Mann-Whitney).

Table 27. Keystone frequency vs genus traits (results/keystone_traits_tests.csv).

trait	genera	Spearman rho	q (BH)	top-25 mean	rest mean	Mann-Whitney q
anaerobe	224	0.303	2.8e-05	0.703	0.411	0.08
gram_neg	225	-0.018	0.91	0.654	0.505	0.22
motile	200	-0.104	0.2	0.443	0.337	0.65
spore	214	-0.137	0.081	0.0455	0.144	0.22
genome_size	222	-0.246	0.00077	3.04e+06	3.64e+06	0.22
gc_content	192	-0.153	0.08	49	50.6	0.65
doubling_h	57	0.008	0.96	10.6	4.36	1

Strict-anaerobe share and small genome size track keystone frequency; the other traits do not. The anaerobe effect keeps its sign and significance in a rank regression with genome size and log study count ($p = 5.8e-04$, HC3) and under permutation within study-count quintiles ($p = 2.0e-04$). Because biome could confound this (anaerobes dominate gut studies), we ran two falsification tests inside studies. First, within each of 155 studies we compared keystones with the other genera: keystones have the higher anaerobe share in

62% of studies (mean difference +0.071, sign-flip $p = 2.0\text{e-}04$). Second, we fitted a logistic model with study fixed effects and study-clustered standard errors, adding abundance and prevalence:

$$\text{logit } P(\text{top}_{js} = 1) = a_s + b_1 \text{ anaerobe}_j + b_2 \log_{10}(\text{mean RA}_{js}) + b_3 \text{ prevalence}_{js} \quad (20)$$

On 10,599 genus-study rows from 156 studies, $b_1 = 0.40$ ($p = 1.4\text{e-}04$), almost the same as without abundance terms (0.41). Anaerobe share is nearly uncorrelated with abundance ($r = 0.09$). So the association is not a biome mix or abundance artifact.

Genomic replication (KEGG). To avoid relying on one phenotype database, we scored marker genes across 11,951 KEGG genomes (203 of 248 genera mapped): pyruvate:ferredoxin oxidoreductase (PFOR; K00169 or K03737), [FeFe]-hydrogenase, complete dissimilatory sulfate reduction (*dsrA*, *dsrB* and *aprA*), aa3 cytochrome c oxidase (*coxA*, K02274), catalase and cytochrome bd. The genomic anaerobe index is

$$\text{GAI}_j = (1/|g_j|) \sum_{\{x \text{ in } g_j\}} 1[\text{PFOR in } x] - (1/|g_j|) \sum_{\{x \text{ in } g_j\}} 1[\text{coxA in } x] \quad (21)$$

It agrees with the Madin anaerobe share (Spearman 0.84) and tracks keystone frequency ($\rho = 0.36$, $p = 9\text{e-}08$). In the within-study fixed-effect logit with abundance and prevalence (13,294 rows, 159 studies), GAI has coefficient 0.32 ($p = 7\text{e-}08$). Most of this is the absence of aerobic respiration (*coxA* -0.52, $p = 6\text{e-}06$; PFOR alone $p = 0.08$), plus complete sulfate reduction (+0.70, $p = 3\text{e-}05$).

Third source (BV-BRC). Oxygen-requirement annotations of 11,323 BV-BRC genomes cover 102 keystone genera (≥ 3 annotated genomes). Their anaerobe share agrees with Madin ($\rho = 0.83$) and KEGG GAI ($\rho = 0.81$), tracks keystone frequency ($\rho = 0.34$, $p = 6\text{e-}04$), and keeps its effect in the fixed-effect logit (coefficient 0.44, $p = 6\text{e-}04$; 6,804 rows, 151 studies). The three trait sources overlap partly in their curated inputs, so they are not fully independent.

Table 28. Within-study anaerobe difference by biome
(results/keystone_traits_biome.json).

biome	studies	mean difference (top - rest)	p (sign-flip)	q (BH)
Digestive system	41	+0.065	0.038	0.10
Oral	19	+0.014	0.720	0.72
Skin	17	+0.110	0.034	0.10
Mammals	20	+0.056	0.174	0.28
Plants	18	+0.041	0.605	0.69
Birds	19	+0.178	0.020	0.10
Fish	14	+0.081	0.134	0.27
Insecta	7	-0.026	0.602	0.69

It is positive in 7 of 8 biomes but no single biome passes FDR (smallest $q = 0.10$). We name it the anaerobe-keystone hypothesis: genera made of strict anaerobes rank as network hubs more often, independent of biome, abundance and prevalence. It is falsified if (a) it fails in an independent cohort with a single biome and adequate power, or (b) perturbation experiments show inferred out-strength does not predict community response. Keystone status here comes from our own inferred networks; the gglasso and igraph tests below fail. It is a method-specific candidate property of ridge interaction inference, not a proven ecological law.

5.6.2 Stress tests and change of direction

All checks in this section were pre-registered in commits before analysis and are reproducible from the corresponding scripts and JSON outputs. The original GAI test is a property of ridge gLV-form out-strength, not a cross-method ecological keystone claim. Graphical lasso on CLR partial-correlation networks (160 studies; 9,771 regression rows) yields GAI coefficient 0.079, $p=0.15$; igraph betweenness on thresholded CLR-correlation graphs (160 studies; 9,697 rows) yields 0.063, $p=0.34$. Agreement with ridge keystone labels is low (Cohen kappa 0.045 and 0.022 respectively). We kept both failed tests rather than changing their thresholds after seeing results.

$$\text{logit } P(\text{top_js}^{\text{method}} = 1) = \alpha_s + \beta_G \text{GAI}_j + \beta_O \text{oral}_j + \beta_A \log_{10}(\text{RA}_{\text{js}} + 10^{-6}) + \beta_P \text{prevalence}_{\text{js}} \quad (22)$$

Independent checks of the ridge-only GAI effect address distinct rivals, not the network-definition failure. Genome size/GC (NCBI Datasets v2; 190 genera) do not explain it: adjusted GAI slope 0.028, HC3 $p=1.3e-6$, versus genome-size $p=0.33$. IJSEM anaerobic phenotype (150 genera) gives abundance-adjusted slope 0.059, $p=6.7e-7$, but IJSEM is a source within the Madin compilation and is not independent. Phylogenetic GLS on the GTDB tree gives GAI slope 0.037 at ML lambda ($p=4.3e-7$) and 0.032 under Brownian covariance ($p=0.0064$); an Open Tree synthetic topology gives 0.028, Brownian $p=0.017$. Synthetic-tree branch lengths are artificial. ProTraits coverage of 51 genera misses its 100-genus gate; its non-significant subset is retained as uninformative, not a replication. For every test in this paragraph, the ridge endpoint remains the same.

Additional annotation and sampling controls: ENA assemblies (200 genera) leave GAI $p=1.8e-9$; Ensembl Genomes (203) $p=2.2e-7$; RNACentral rRNA sequence counts (203) $p=8.2e-6$. The assembly or rRNA count terms themselves were null. UniProt EC calls (198) yield a related index with abundance-adjusted $p=0.0013$; InterPro signatures (198) yield $p=0.00024$, but their protein sets overlap. A GBIF occurrence/country adjustment reduces the GAI slope to 0.019 ($p=0.024$), indicating shared variance with a noisy generalism proxy. Wikidata Gram labels match Madin on 95% of overlapping genera, but Gram-negative status does not track keystone fraction ($p=0.90$). None of these controls removes the network-method caveat.

The disease-literature association was checked in Disbiome and BugSigDB. Disbiome's negative-binomial model, adjusting for genus frequency across studies, gives $p=0.024$; controlling the number of Europe PMC articles gives $p=0.0017$. Crossref works count gives Disbiome $p=0.0080$ and BugSigDB $p=0.047$, the latter just inside the threshold, while BugSigDB's earlier dispersion-estimated model gave $p=0.062$. These are associations between two published-data proxies; oral HACEK genera and fuzzy name queries remain confounds.

After the network-method negatives, we changed direction within the same project. HOMD v4.2 classifies 46 of 203 GAI genera as oral. The ridge genus-level GAI slope remains 0.024 ($p=3.7e-6$) after adjustment for oral status; HOMD oral coefficient is 0.042 ($p=0.0084$). On new method-general tests using study-fixed-effect logits with abundance, prevalence and GAI, oral origin has positive coefficients under graphical lasso (0.211; $p=0.0265$) and igrph (0.322; $p=0.00105$). GAI is null in those methods. A Bayesian random-study-intercept re-fit gives $P(\text{oral coefficient} > 0) = 1.0$ for both with two chains of 500 draws; this is a model check, not an independent cohort. Leaving out each of five multi-genus HOMD families keeps the oral association significant in ten fits (gglasso $p=0.0197$ -0.0433; igrph $p=0.0003$ -0.0028).

A separate, pre-registered PubTator3 oral/dental literature fraction did not replicate the binary HOMD effect: despite 203 measured genera, graphical-lasso $p=0.57$ and igrph $p=0.23$. The index correlates with HOMD membership ($\rho=0.54$) but is a different, literature-biased construct. Bio.Phylo gives no oral clustering on the Open Tree topology ($p=0.53$). Thus the defensible claim is a HOMD-list-specific candidate for network hubs under two definitions, not an oral-genus law or causality. The ridge oral coefficient in the per-study model is nonsignificant ($p=0.17$) and its fit did not converge; we do not claim a three-method result.

5.6.3 Predictive and software checks

On the ridge-only genus-level keystone fraction, XGBoost with 20 repeats of five-fold cross-validation gains mean out-of-fold R^2 of 0.27 from GAI (20/20 positive deltas) but has negative absolute R^2 (-0.11 with GAI, -0.38 without). That model is not usable for prediction. Separately pre-registered shallower boosters on 200 genera produce mean full-data out-of-fold R^2 0.133 for LightGBM (reduced -0.071) and 0.119 for CatBoost (reduced -0.045), both with 20/20 positive deltas. The modest improvement is within the same dataset and outcome. SHAP TreeExplainer ranks GAI first in the fitted LightGBM model (mean absolute contribution 0.0265 versus ENA 0.00772 and abundance 0.00576), an in-sample explanation rather than proof of causality. SymPy verifies the cNODE and gLV vector fields are tangent to the simplex and zero on absent-taxon boundaries for a generic three-taxon system when abundances sum to one. That algebra does not prove numerical stability or predictive accuracy.

Table 29. Key post-audit tests. Primary methods, thresholds, preregistrations and full negative results are in results/finding_keystone.md and the JSON files.

analysis	scope	pre-registered verdict	key caveat
gglasso / igraph GAI	160 MGnify studies	FAIL / FAIL	ridge-specific effect
HOMD oral on gglasso / igraph	46 of 203 GAI genera oral	PASS / PASS	PubTator index failed
HOMD family deletions	5 families; 10 fits	PASS	same curated list
LightGBM / CatBoost	200 genera; repeated CV	PASS / PASS	modest, ridge-only R2
Bio.Phylo oral topology	160 tips; 38 oral	FAIL	synthetic topology

5.7 Pre-registered discovery arms

Two discovery tests were pre-registered (PREREG_twindiscovery.md) with explicit redirect clauses: T1 asks whether GraphTwin gate out-strength replicates the keystone-consensus signal of Section 5.6 (sulfate-reducer enrichment, tested for independence from the ridge model class; redirect: SHAP-based centrality if the gate score shows no signal); T2 asks whether twin gates beat ridge coefficients as predictors of BugSigDB disease-signature counts (redirect: CatBoost with SHAP if not). Both arms are fit per MGnify study under the same caps as the audit. Both redirect rules, scorers, and test constructions were locked in PREREG_twindiscovery.md (committed 2026-09-25 15:55 IST) before any T1/T2 run; the completion-order stopping-rule amendment was committed 2026-09-26 07:12 IST before any T1 output was read, the redirect path was fixed before those results existed, but it remains downstream of failed primary tests and must be treated as exploratory rather than independent replication.

T1 (model-class independence of the keystone enrichment). T1: 37 studies, 56 genera. Top enrichment: Actinomycetota $q=0.906$. Desulfobacterota $q=$ None. NULL - redirect to SHAP centrality per pre-reg

Table 30. Top 15 genera by GraphTwin gate out-strength consensus (results/twindiscovery_t1_consensus.csv).

genus	studies	top	fraction top	p	q (BH)
Acinetobacter	32	22	0.688	2.3627632274605656e-15	1.3231474073779166e-13
Pseudomonas	34	20	0.588	3.4367075799906436e-12	7.966283892670575e-11
Staphylococcus	31	19	0.613	4.267652085359237e-12	7.966283892670575e-11
Abiotrophia	20	15	0.750	9.481303720000011e-12	8.849216805333343e-11
Peptostreptococcus	20	15	0.750	9.481303720000011e-12	8.849216805333343e-11
Solobacterium	20	15	0.750	9.481303720000011e-12	8.849216805333343e-11
Bifidobacterium	26	17	0.654	1.2810380798769867e-11	1.0248304639015893e-10
Lautropia	21	15	0.714	3.0079956880000025e-11	1.6844775852800012e-10
Treponema	21	15	0.714	3.0079956880000025e-11	1.6844775852800012e-10

genus	studies	top	fraction top	p	q (BH)
				0025e-11	0012e-10
Peptococcus	21	15	0.714	3.007995688000 0025e-11	1.684477585280 0012e-10
Fastidiosipila	24	16	0.667	3.338913051015 403e-11	1.699810280516 9325e-10
Lactobacillus	34	19	0.559	4.164943296236 0486e-11	1.943640204910 156e-10
Blautia	28	17	0.607	7.224555155610 398e-11	3.112116067032 171e-10
Atopobium	26	16	0.615	1.980191282012 2747e-10	7.392714119512 492e-10
Haemophilus	26	16	0.615	1.980191282012 2747e-10	7.392714119512 492e-10

Table 31. Phylum enrichment among the 25 lowest-p genera by twin-gate score (one-sided Fisher, BH).

phylum	in top 25	genera	p	q (BH)
Actinomycetota	4	7	0.377453351926159 8	0.905724408776573 5
Campylobacterota	1	1	0.446428571428571 45	0.905724408776573 5
Spirochaetota	1	1	0.446428571428571 45	0.905724408776573 5
Pseudomonadota	4	8	0.517556805015184 9	0.905724408776573 5
Fusobacteriota	1	2	0.698051948051948	0.909292705075169
Bacillota	11	27	0.798163055615896 7	0.909292705075169
Bacteroidota	2	7	0.909292705075169	0.909292705075169

T2 (twin gates vs ridge as disease-literature predictor). T2: 56 genera. Twin-gate coef - 2.28 (p=0.168), AIC 767.1 vs ridge 766.1. TWIN DOES NOT WIN - redirect to CatBoost+SHAP per pre-reg

Table 32. Negative-binomial models of BugSigDB signature counts (results/twindiscovery_t2.json).

model	score coef	p	AIC	ML alpha	ML coef	ML p	ML AIC
ridge keystone score	-0.351	0.241	766.1	0.62	-0.351	0.152	762.0
twin gate frequency	-2.28	0.168	767.1	0.63	-2.28	0.0737	763.7

Redirects (locked in the same pre-registration before any T1 run). R1, the scorer redirect to SHAP-attribution centrality from the validated LightGBM keystone model: Desulfobacterota 1/3 in top decile, q = 0.462; no phylum passes FDR (best q = 0.462). The sulfate-reducer enrichment is therefore specific to the ridge out-strength definition; it does not survive either model-class change. R2, the predictor redirect to CatBoost + SHAP per-genus attribution ranking in the same negative-binomial comparison: coefficient +37.07 (p = 2.6e-35), AIC 2343.9 vs ridge 2443.1 on 200 genera - the redirect WINS. Gradient-boosted interaction attributions fit BugSigDB disease-signature counts

better than the linear ridge score in this viewed comparison; this exploratory proxy association is not an independently replicated biological discovery.

Table 33. T2 redirect: negative-binomial models of BugSigDB signature counts (results/keystone_redirects.json).

model	score coef	p	AIC
ridge keystone score	-0.351	0.241	2443.1
CatBoost+SHAP attribution	+37.07	2.6e-35	2343.9

Appendix B. Full MGnify interaction audit (per study)

Table 34. All 160 audited studies (results/mgnify_audit.csv). Gain is (prior - interaction) / prior; positive favours interactions.

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
MGYS000 06862	root:Host-associated: Human:Digestive system	84	55	0.173	0.170	0.004	0.1174934141540573
MGYS000 05791	root:Host-associated: Human:Digestive system	243	47	0.178	0.172	0.006	4.8384265757998056e-05
MGYS000 06825	root:Host-associated: Human:Digestive system	197	57	0.174	0.168	0.006	0.0051888744323153
MGYS000 06824	root:Host-associated: Human:Digestive system	40	66	0.175	0.180	-0.004	0.97349265041521
MGYS000 06822	root:Host-associated: Human:Digestive system	55	58	0.192	0.189	0.003	0.0874094952304293
MGYS000 06814	root:Host-associated: Human:Digestive system	68	70	0.177	0.175	0.002	0.3404810744509645
MGYS000 06815	root:Host-associated: Human:Digestive system	67	46	0.167	0.167	-0.000	0.851347514124289
MGYS000 06811	root:Host-associated: Human:Digestive	98	39	0.160	0.154	0.007	0.0116937067032781

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
	system						
MGYS000 06807	root:Host-associated: Human:Digestive system	60	65	0.179	0.180	-0.001	0.192573487700825
MGYS000 06805	root:Host-associated: Human:Digestive system	86	66	0.180	0.162	0.017	0.0159873022676302
MGYS000 06806	root:Host-associated: Human:Digestive system	74	60	0.178	0.180	-0.002	0.0059561819902366
MGYS000 06760	root:Host-associated: Human:Digestive system	218	69	0.169	0.164	0.005	1.1800880841672732e-09
MGYS000 02338	root:Host-associated: Human:Digestive system	303	44	0.154	0.148	0.006	0.002728300451495
MGYS000 06797	root:Host-associated: Human:Digestive system	175	72	0.163	0.160	0.002	0.0254205801722315
MGYS000 06551	root:Host-associated: Human:Digestive system	400	63	0.167	0.162	0.005	0.0001113060975781
MGYS000 06789	root:Host-associated: Human:Digestive system	108	63	0.178	0.179	-0.001	0.0169520637921749
MGYS000 06787	root:Host-associated: Human:Digestive system	108	56	0.180	0.177	0.004	0.5418861206076111
MGYS000 06786	root:Host-associated: Human:Digestive system	134	52	0.173	0.169	0.004	0.0144364716231558
MGYS000 06759	root:Host-associated: Human:Digestive system	400	72	0.183	0.176	0.006	3.964715623438228e-19
MGYS000 06780	root:Host-associated: Human:Digestive system	24	27	0.129	0.130	-0.001	0.943392056822656

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
	estive system						
MGYS000 06779	root:Host-associated: Human:Digestive system	57	63	0.185	0.185	-0.000	0.6420785027937037
MGYS000 06769	root:Host-associated: Human:Digestive system	116	54	0.186	0.187	-0.001	0.8169950973439684
MGYS000 06765	root:Host-associated: Human:Digestive system	28	43	0.154	0.159	-0.005	0.6455853804945946
MGYS000 05333	root:Host-associated: Human:Digestive system	244	53	0.165	0.158	0.008	1.8156083569124025e-07
MGYS000 05230	root:Host-associated: Human:Digestive system	400	114	0.363	0.311	0.052	2.5108503132493804e-33
MGYS000 02675	root:Host-associated: Human:Digestive system	24	23	0.142	0.136	0.007	0.704756498336792
MGYS000 06069	root:Host-associated: Human:Digestive system	400	48	0.167	0.158	0.009	4.343497031756574e-07
MGYS000 06006	root:Host-associated: Human:Digestive system	400	106	0.181	0.176	0.005	3.3971201874549978e-22
MGYS000 02061	root:Host-associated: Human:Digestive system	52	40	0.142	0.142	-0.000	0.160154524509844
MGYS000 00633	root:Host-associated: Human:Digestive system	400	68	0.572	0.438	0.135	7.59928368626225e-21
MGYS000 02238	root:Host-associated: Human:Digestive system	111	111	0.402	0.301	0.101	2.366731168120885e-13
MGYS000 06562	root:Host-associated:	400	44	0.167	0.163	0.004	8.224596795214574e-

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
	Human:Digestive system						11
MGYS00006649	root:Host-associated: Human:Digestive system	20	44	0.176	0.176	-0.000	0.2305126190185547
MGYS00006641	root:Host-associated: Human:Digestive system	87	52	0.159	0.156	0.002	0.175661084603356
MGYS00006642	root:Host-associated: Human:Digestive system	57	19	0.149	0.150	-0.001	0.0689068793681605
MGYS00006364	root:Host-associated: Human:Digestive system	160	33	0.190	0.191	-0.001	0.0006405387930733
MGYS00006640	root:Host-associated: Human:Digestive system	60	53	0.170	0.176	-0.006	0.6725238936938076
MGYS00005253	root:Host-associated: Human:Digestive system:Oral	400	83	0.315	0.286	0.030	5.630794366587372e-13
MGYS00006320	root:Host-associated: Human:Digestive system:Oral	392	68	0.347	0.284	0.063	2.6018524371287146e-18
MGYS00006263	root:Host-associated: Human:Digestive system:Oral	400	101	0.292	0.259	0.032	1.6144709003228016e-16
MGYS00006017	root:Host-associated: Human:Digestive system:Oral	53	144	0.314	0.291	0.023	0.0063127550930951
MGYS00005592	root:Host-associated: Human:Digestive system:Oral	160	104	0.323	0.257	0.065	2.954937077626768e-12
MGYS000	root:Host-	71	112	0.330	0.246	0.083	3.12962737

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
05569	associated: Human:Digestive system:Oral						22221307e-09
MGYS000 05528	root:Host-associated: Human:Digestive system:Oral	400	68	0.312	0.267	0.045	2.94477499 14759767e-24
MGYS000 05506	root:Host-associated: Human:Digestive system:Oral	400	78	0.259	0.237	0.022	1.25438678 96369403e-26
MGYS000 05223	root:Host-associated: Human:Digestive system:Oral	400	99	0.288	0.218	0.070	6.02598443 8246577e-36
MGYS000 05143	root:Host-associated: Human:Digestive system:Oral	400	142	0.372	0.326	0.046	7.39947150 4901985e-36
MGYS000 02272	root:Host-associated: Human:Digestive system:Oral	40	96	0.525	0.491	0.035	0.01019742 31565691
MGYS000 02269	root:Host-associated: Human:Digestive system:Oral	147	105	0.410	0.342	0.069	1.76030560 13155362e-16
MGYS000 02229	root:Host-associated: Human:Digestive system:Oral	100	85	0.448	0.358	0.090	5.56345569 5355892e-12
MGYS000 05232	root:Host-associated: Human:Skin	400	96	0.365	0.325	0.040	9.31108298 430499e-22
MGYS000 06246	root:Host-associated: Human:Skin	50	93	0.354	0.256	0.099	0.00651664 99849652
MGYS000 05384	root:Host-associated: Human:Skin	400	99	0.422	0.260	0.162	8.32184474 8869964e-43

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
MGYS00006795	root:Host-associated: Human:Digestive system	31	150	0.582	0.505	0.078	0.0004954077303409576
MGYS00006794	root:Host-associated: Human:Digestive system	184	150	0.399	0.314	0.085	3.4225595177609226e-22
MGYS00006755	root:Host-associated: Human:Digestive system	83	150	0.562	0.375	0.187	1.3575561479460155e-13
MGYS00005154	root:Host-associated: Human:Digestive system	400	150	0.492	0.389	0.102	9.623147369597524e-46
MGYS00006276	root:Host-associated: Human:Digestive system:Oral	91	150	0.298	0.283	0.015	0.0001867751839014596
MGYS00005189	root:Host-associated: Human:Digestive system:Oral	400	150	0.656	0.407	0.249	6.400607441863894e-60
MGYS00005197	root:Host-associated: Human:Digestive system:Oral	400	150	0.197	0.170	0.027	7.298198067237661e-42
MGYS00002159	root:Host-associated: Human:Digestive system:Oral	21	150	0.573	0.494	0.080	0.08879852294921875
MGYS00003719	root:Host-associated: Human:Digestive system:Oral	21	150	0.306	0.287	0.018	0.05469322204589844
MGYS00002395	root:Host-associated: Human:Digestive system:Oral	68	150	0.351	0.346	0.005	0.025728924772950645
MGYS00002394	root:Host-associated: Human:Dig	20	150	0.417	0.400	0.016	0.16495704650878906

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
	estive system:Oral						
MGYS00006648	root:Host-associated:Human:Skin	67	150	0.470	0.344	0.126	1.0192740768516179e-07
MGYS00005156	root:Host-associated:Human:Skin	400	150	0.399	0.339	0.060	3.7584514383668004e-33
MGYS00005172	root:Host-associated:Human:Skin	400	150	0.401	0.302	0.099	1.9412619707417544e-34
MGYS00005533	root:Host-associated:Human:Skin	400	150	0.317	0.316	0.001	8.614275319605027e-22
MGYS00005212	root:Host-associated:Human:Skin	400	150	0.374	0.292	0.083	4.0222879063051704e-33
MGYS00005238	root:Host-associated:Human:Skin	400	106	0.396	0.358	0.039	2.7355102946682556e-13
MGYS00005101	root:Host-associated:Human:Skin	71	91	0.367	0.292	0.075	0.001176166597423719
MGYS00005102	root:Host-associated:Human:Skin	179	150	0.381	0.264	0.118	7.298919552265234e-19
MGYS00005135	root:Host-associated:Human:Skin	109	150	0.470	0.420	0.050	1.4252348445599325e-07
MGYS00005098	root:Host-associated:Human:Skin	101	150	0.127	0.102	0.025	1.5388621128648606e-10
MGYS00005037	root:Host-associated:Human:Skin	106	150	0.311	0.233	0.078	6.006215258961562e-07
MGYS00005090	root:Host-associated:Human:Skin	120	7	0.082	0.077	0.005	0.9498798977422346
MGYS00002184	root:Host-associated:Human:Skin	400	150	0.471	0.261	0.210	2.878775655656423e-63
MGYS00002072	root:Host-associated:Human:Skin	258	150	0.233	0.190	0.043	2.0058912688460236e-19

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
	n						
MGYS000 06855	root:Host-associated: Mammals	116	150	0.293	0.232	0.061	1.6553837550067384e-12
MGYS000 06799	root:Host-associated: Mammals	27	150	0.406	0.417	-0.011	0.03189583122730255
MGYS000 02673	root:Host-associated: Mammals	93	150	0.650	0.533	0.116	1.1556875270009098e-13
MGYS000 06747	root:Host-associated: Mammals	101	150	0.568	0.386	0.182	4.170399034793852e-18
MGYS000 01223	root:Host-associated: Mammals	30	98	0.194	0.173	0.021	0.08406547829508781
MGYS000 01224	root:Host-associated: Mammals	66	89	0.389	0.242	0.147	1.6687130603014278e-09
MGYS000 05267	root:Host-associated: Mammals	400	71	0.343	0.254	0.088	3.2675706524688328e-52
MGYS000 00426	root:Host-associated: Mammals	44	66	0.272	0.245	0.027	0.4011977486671867
MGYS000 05228	root:Host-associated: Mammals	400	63	0.257	0.206	0.052	2.070366009312925e-33
MGYS000 06537	root:Host-associated: Mammals	85	150	0.495	0.407	0.088	0.00037059731217404824
MGYS000 06536	root:Host-associated: Mammals	65	150	0.384	0.384	-0.000	8.467639424268034e-05
MGYS000 06538	root:Host-associated: Mammals	40	150	0.515	0.512	0.003	0.08415465162215696
MGYS000 05628	root:Host-associated: Mammals	400	150	0.338	0.276	0.062	7.073539193762629e-38
MGYS000 06450	root:Host-associated: Mammals	144	81	0.444	0.310	0.134	1.8256817238051636e-13
MGYS000 06456	root:Host-associated: Mammals	24	76	0.331	0.313	0.017	0.975736119417176
MGYS000 06446	root:Host-associated: Mammals	400	62	0.317	0.247	0.070	4.779278261362484e-25
MGYS000 06433	root:Host-associated: Mammals	25	150	0.557	0.537	0.020	0.018744707107543945
MGYS000 06430	root:Host-associated: Mammals	32	150	0.397	0.299	0.099	0.014118533720624406
MGYS000 06427	root:Host-associated: Mammals	85	109	0.389	0.341	0.048	0.00245958589664498

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
MGYS00006424	root:Host-associated: Mammals	142	150	0.280	0.252	0.027	2.5100543778701536e-14
MGYS00006652	root:Host-associated: Plants	400	150	0.250	0.175	0.076	2.8719573482573203e-56
MGYS00001412	root:Host-associated: Plants	48	98	0.727	0.639	0.088	0.004499547313045582
MGYS00001368	root:Host-associated: Plants	400	150	0.276	0.190	0.086	5.315213480136056e-65
MGYS00003172	root:Host-associated: Plants	45	150	0.247	0.101	0.145	5.684341886080802e-14
MGYS00006578	root:Host-associated: Plants	27	150	0.162	0.170	-0.009	0.9153143465518951
MGYS00000957	root:Host-associated: Plants	70	61	0.381	0.351	0.030	0.0010592552437449615
MGYS00001704	root:Host-associated: Plants	383	150	0.252	0.164	0.088	2.432607986872169e-55
MGYS00000867	root:Host-associated: Plants	32	150	0.001	0.001	0.000	0.34500763198357975
MGYS00006031	root:Host-associated: Plants	89	150	0.477	0.402	0.075	1.1054300201400448e-07
MGYS00006021	root:Host-associated: Plants	135	18	0.233	0.216	0.017	0.0035382477408565484
MGYS00006018	root:Host-associated: Plants	181	150	0.452	0.287	0.166	7.976256015721753e-13
MGYS00001349	root:Host-associated: Plants	108	150	0.402	0.211	0.190	9.800115982359504e-18
MGYS00005260	root:Host-associated: Plants	400	150	0.396	0.248	0.148	1.2745313376300463e-61
MGYS00005227	root:Host-associated: Plants	400	150	0.223	0.201	0.022	2.4755145503466964e-21
MGYS00000983	root:Host-associated: Plants	400	150	0.178	0.141	0.036	7.250624660443202e-39
MGYS00001347	root:Host-associated: Plants	400	150	0.471	0.218	0.252	6.620371482312783e-66
MGYS00005740	root:Host-associated: Plants	34	150	0.400	0.339	0.062	3.2188836485147476e-07
MGYS00005726	root:Host-associated: Plants	33	150	0.382	0.337	0.046	5.099480040371418e-05
MGYS000	root:Host-	400	150	0.311	0.171	0.140	2.36530085

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
05169	associated: Plants						5.1996965e-59
MGYS000 05151	root:Host-associated: Plants	400	150	0.348	0.111	0.238	6.538056051035535e-60
MGYS000 06750	root:Host-associated: Birds	340	89	0.142	0.135	0.007	1.1556877632363245e-12
MGYS000 06533	root:Host-associated: Birds	94	150	0.404	0.322	0.081	2.1203049163202955e-09
MGYS000 06532	root:Host-associated: Birds	40	150	0.478	0.252	0.226	5.723268259316683e-08
MGYS000 05766	root:Host-associated: Birds	138	150	0.474	0.245	0.229	7.806550735727007e-23
MGYS000 06062	root:Host-associated: Birds	220	92	0.186	0.166	0.020	2.1652952739860382e-13
MGYS000 06042	root:Host-associated: Birds	355	96	0.284	0.124	0.160	1.1028252358036625e-43
MGYS000 05621	root:Host-associated: Birds	60	150	0.265	0.173	0.092	1.0181154208756968e-06
MGYS000 05611	root:Host-associated: Birds	117	115	0.334	0.178	0.156	3.79430763036704e-14
MGYS000 05610	root:Host-associated: Birds	251	87	0.376	0.317	0.059	1.1443357306387731e-14
MGYS000 05609	root:Host-associated: Birds	314	68	0.470	0.395	0.075	1.3785359922199967e-23
MGYS000 05600	root:Host-associated: Birds	286	150	0.270	0.175	0.095	2.8587965104295026e-40
MGYS000 05593	root:Host-associated: Birds	150	129	0.589	0.361	0.228	1.9022991582307127e-20
MGYS000 05576	root:Host-associated: Birds	34	150	0.374	0.220	0.153	3.573950380086899e-08
MGYS000 05187	root:Host-associated: Birds	400	4	0.147	0.131	0.016	0.20038180242578474
MGYS000 05316	root:Host-associated: Birds	400	145	0.542	0.257	0.285	1.506095419477192e-57
MGYS000 05317	root:Host-associated: Birds	400	150	0.491	0.349	0.141	1.3522978810793892e-53
MGYS000 05309	root:Host-associated: Birds	72	150	0.312	0.194	0.118	1.6709026814241643e-12
MGYS000 04995	root:Host-associated:	43	150	0.256	0.223	0.033	0.10220764011114625

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
	Birds						
MGYS000 03706	root:Host-associated: Birds	400	150	0.516	0.422	0.094	2.0461254236329354e-34
MGYS000 03681	root:Host-associated: Birds	122	150	0.674	0.444	0.231	4.108703914852822e-12
MGYS000 06764	root:Host-associated: Fish	21	150	0.262	0.181	0.081	1.811981201171875e-05
MGYS000 06763	root:Host-associated: Fish	39	150	0.420	0.364	0.055	2.3796659661456943e-05
MGYS000 06762	root:Host-associated: Fish	137	150	0.428	0.290	0.138	1.851401098428866e-14
MGYS000 06761	root:Host-associated: Fish	123	150	0.295	0.231	0.064	8.880400175577094e-10
MGYS000 06758	root:Host-associated: Fish	135	143	0.522	0.312	0.209	2.601298498785815e-10
MGYS000 06041	root:Host-associated: Fish	400	30	0.278	0.164	0.113	6.995323507416229e-32
MGYS000 06117	root:Host-associated: Fish	90	72	0.028	0.030	-0.002	1.0043017278491953e-06
MGYS000 06086	root:Host-associated: Fish	135	3	0.000	0.000	0.000	1.0
MGYS000 05503	root:Host-associated: Fish	400	150	0.594	0.467	0.127	1.9972646703767122e-30
MGYS000 03166	root:Host-associated: Fish	70	150	0.417	0.352	0.065	8.973570568522901e-07
MGYS000 04672	root:Host-associated: Fish	33	150	0.422	0.372	0.050	0.05752792255952954
MGYS000 02905	root:Host-associated: Fish	66	134	0.191	0.069	0.122	2.107866932199474e-09
MGYS000 02608	root:Host-associated: Fish	101	150	0.589	0.429	0.161	4.037951900494873e-09
MGYS000 02413	root:Host-associated: Fish	173	150	0.588	0.356	0.232	1.508508341940648e-16
MGYS000 02065	root:Host-associated: Fish	214	150	0.499	0.363	0.136	6.023085537966328e-20
MGYS000 06798	root:Host-associated: Insecta	24	150	0.567	0.174	0.393	8.344650268554688e-06
MGYS000 06019	root:Host-associated: Insecta	20	150	0.345	0.215	0.130	0.000209808349609375

study	biome	n	taxa	prior median	interaction median	gain	Wilcoxon p
MGYS00005862	root:Host-associated:Insecta	21	150	0.321	0.346	-0.026	0.05950736999511719
MGYS00005215	root:Host-associated:Insecta	400	92	0.296	0.244	0.052	6.345809298780021e-27
MGYS00005163	root:Host-associated:Insecta	386	150	0.548	0.467	0.081	2.789722004048674e-31
MGYS00003800	root:Host-associated:Insecta	154	96	0.184	0.184	0.000	0.6706023900543987
MGYS00003705	root:Host-associated:Insecta	384	70	0.408	0.370	0.038	2.7686978173767286e-09

Appendix C. Keystone consensus, all genera

Table 35. All 248 genera modelled in ≥ 20 studies (results/keystone_consensus.csv).

genus	studies	top	fraction top	p	q (BH)
Cardiobacterium	30	13	0.433	2.302322724161098e-06	0.0005709760355919523
Eikenella	28	11	0.393	4.232321631540403e-05	0.0052480788231101
Comamonas	51	13	0.255	0.0012261499891986897	0.09862502890393288
Johnsonella	35	10	0.286	0.0017422103851516847	0.09862502890393288
Simonsiella	20	7	0.350	0.002386089408966118	0.09862502890393288
Desulfobulbus	20	7	0.350	0.002386089408966118	0.09862502890393288
Sutterella	56	13	0.232	0.0030302254033037486	0.10735655714561852
Agathobacter	77	16	0.208	0.003583040147751078	0.11107424458028342
Collinsella	67	14	0.209	0.005889699276757321	0.16110995747664922
Desulfovibrio	68	14	0.206	0.00674761803544495	0.16110995747664922
Anaerosporbacter	30	8	0.267	0.00778361923792132	0.16110995747664922
Solobacterium	49	11	0.224	0.008004971888016863	0.16110995747664922
Bilophila	56	12	0.214	0.008445280029017902	0.16110995747664922
Barnesiella	47	10	0.213	0.016260359901383525	0.2880406611102225
Bergeyella	49	10	0.204	0.02150126888113899	0.3504201263576335
Anaerovorax	23	6	0.261	0.022607750087589258	0.3504201263576335
Odoribacter	57	11	0.193	0.02423840456349196	0.3535955489262357
Cetobacterium	30	7	0.233	0.025826788661	0.355835754897

genus	studies	top	fraction top	p	q (BH)
				909898	42525
Mogibacterium	45	9	0.200	0.031994021841 86363	0.412814207952 3551
Eubacterium	60	11	0.183	0.034209123078 73303	0.412814207952 3551
Subdoligranulum	83	14	0.169	0.035467001433 08487	0.412814207952 3551
Slackia	39	8	0.205	0.036620615221 57989	0.412814207952 3551
Campylobacter	69	12	0.174	0.040017895148 75476	0.431497304212 66
Enhydrobacter	20	5	0.250	0.043174495284 46339	0.446136451272 78834
Fastidiosipila	86	14	0.163	0.046060001870 58948	0.453852358283 953
Stomatobaculum	34	7	0.206	0.048144326629 12472	0.453852358283 953
Phascolarctobacterium	56	10	0.179	0.049411345458 33359	0.453852358283 953
Fournierella	28	6	0.214	0.055006948868 291425	0.470404252390 906
Propionibacterium	28	6	0.214	0.055006948868 291425	0.470404252390 906
Moryella	43	8	0.186	0.060669351591 22358	0.483186617866 3752
Aggregatibacter	36	7	0.194	0.062828142375 83728	0.483186617866 3752
Fretibacterium	29	6	0.207	0.063717444452 2859	0.483186617866 3752
Leptotrichia	51	9	0.176	0.064294993506 41282	0.483186617866 3752
Dolosigranulum	37	7	0.189	0.071085058271 75729	0.518502777982 2297
Anaeroplasma	30	6	0.200	0.073190108399 87991	0.518604196662 0062
Papillibacter	24	5	0.208	0.085074885878 67089	0.577705898153 1825
Cutibacterium	54	9	0.167	0.086189992869 62803	0.577705898153 1825
Anaerotruncus	39	7	0.179	0.089433889739 4184	0.583673806720 4149
Abiotrophia	40	7	0.175	0.099516423965 55122	0.626546230147 6376
Lautropia	56	9	0.161	0.102970794170 76241	0.626546230147 6376
Moraxella	33	6	0.182	0.106108635750 80959	0.626546230147 6376
Butyricimonas	33	6	0.182	0.106108635750 80959	0.626546230147 6376
Butyricicoccus	66	10	0.152	0.119923548936 55092	0.672906616565 6145
Enterorhabdus	42	7	0.167	0.121447206815 06566	0.672906616565 6145
Hydrogenoanaerobacterium	27	5	0.185	0.126557448308 2461	0.672906616565 6145
Tumebacillus	27	5	0.185	0.126557448308 2461	0.672906616565 6145
Leuconostoc	35	6	0.171	0.131635774841	0.672906616565

genus	studies	top	fraction top	p	q (BH)
				83725	6145
Saccharofermentans	20	4	0.200	0.1329533234343351	0.6729066165656145
Marinomonas	20	4	0.200	0.1329533234343351	0.6729066165656145
Parasutterella	60	9	0.150	0.1416356047983447	0.691053967992992
Ammopiptanthus	28	5	0.179	0.14211190470823626	0.691053967992992
Finegoldia	36	6	0.167	0.14539730133503725	0.6934332832901777
Alistipes	78	11	0.141	0.1536206356148474	0.718828634575135
Ruminococcus	45	7	0.156	0.15850738539835585	0.7279598440517083
Paucibacter	88	12	0.136	0.1670511156120358	0.753248666759725
Carnobacterium	22	4	0.182	0.17192789706645586	0.7613949727228759
Erysipelothrix	23	4	0.174	0.19273101351862001	0.8385489710985572
Oribacterium	65	9	0.138	0.198722696795465	0.8497108414702642
Neisseria	66	9	0.136	0.21116466751003035	0.8710989390747419
Fluviicola	32	5	0.156	0.21149828869932752	0.8710989390747419
Ezakiella	24	4	0.167	0.2142622390466099	0.8710989390747419
Akkermansia	67	9	0.134	0.22390696165563617	0.8769946942869572
Selenomonas	33	5	0.152	0.23031786958299133	0.8769946942869572
Methanobrevibacter	33	5	0.152	0.23031786958299133	0.8769946942869572
Nocardia	25	4	0.160	0.23640864244682802	0.8769946942869572
Fibrobacter	25	4	0.160	0.23640864244682802	0.8769946942869572
Megasphaera	68	9	0.132	0.2369300182146215	0.8769946942869572
Lachnoanaerobaculum	34	5	0.147	0.2495917162121227	0.8891872058135635
Kingella	34	5	0.147	0.2495917162121227	0.8891872058135635
Parabacteroides	87	11	0.126	0.2509802597054413	0.8891872058135635
Catonella	35	5	0.143	0.2692510397738372	0.9404825051255159
Corynebacterium	91	11	0.121	0.29998647993678423	1.0
Tannerella	28	4	0.143	0.3054336969081328	1.0
Candidatus_Soleaferrea	37	5	0.135	0.30945356374288324	1.0
Lachnospira	55	7	0.127	0.30955723920441663	1.0
Anaerostipes	83	10	0.120	0.316254147757	1.0

genus	studies	top	fraction top	p	q (BH)
				0478	
Bradyrhizobium	65	8	0.123	0.32314240394111926	1.0
Olsenella	47	6	0.128	0.3285688971461327	1.0
Acetitomaculum	29	4	0.138	0.328952034984918	1.0
Pedobacter	38	5	0.132	0.329863552054893	1.0
Atopobium	57	7	0.123	0.34289143772806113	1.0
Gemmata	21	3	0.143	0.3515911755169639	1.0
Negativibacillus	58	7	0.121	0.35972337214077327	1.0
Treponema	40	5	0.125	0.3709823034657993	1.0
Anaerofilum	22	3	0.136	0.37995906158809667	1.0
Phaselicystis	22	3	0.136	0.37995906158809667	1.0
Marvinbryantia	50	6	0.120	0.3838769922757232	1.0
Brevundimonas	69	8	0.116	0.3854441012423575	1.0
Stenotrophomonas	60	7	0.117	0.39354859789913854	1.0
Coxiella	33	4	0.121	0.4230563358743078	1.0
Phenylobacterium	43	5	0.116	0.43252905747178994	1.0
Coprobacter	24	3	0.125	0.4357262730487911	1.0
Filifactor	24	3	0.125	0.4357262730487911	1.0
Shewanella	24	3	0.125	0.4357262730487911	1.0
Blautia	102	11	0.108	0.4432030664009525	1.0
Holdemania	34	4	0.118	0.4461849518292653	1.0
Enterobacter	35	4	0.114	0.46901513402996614	1.0
Aeromonas	36	4	0.111	0.4914885946337808	1.0
Caulobacter	36	4	0.111	0.4914885946337808	1.0
Pseudopropionibacterium	27	3	0.111	0.5154188583079999	1.0
Parvimonas	48	5	0.104	0.5314392758472106	1.0
Mycoplasma	48	5	0.104	0.5314392758472106	1.0
Brachybacterium	38	4	0.105	0.5351640227227555	1.0
Sphingobacterium	39	4	0.103	0.5562806425628786	1.0
Acinetobacter	101	10	0.099	0.561751462265	1.0

genus	studies	top	fraction top	p	q (BH)
				9117	
Enterococcus	91	9	0.099	0.565055679835745	1.0
Arthrobacter	50	5	0.100	0.5688015931709377	1.0
Dorea	71	7	0.099	0.5736490594127721	1.0
Intestinimonas	40	4	0.100	0.5768693469069974	1.0
Massilia	61	6	0.098	0.5794548905240392	1.0
Facklamia	30	3	0.100	0.5886487604404943	1.0
Williamsia	20	2	0.100	0.6082530018748321	1.0
Cloacibacterium	20	2	0.100	0.6082530018748321	1.0
Methylobacterium	84	8	0.095	0.6111403728323437	1.0
Peredibacter	31	3	0.097	0.6114143824771845	1.0
Terrisporobacter	31	3	0.097	0.6114143824771845	1.0
Granulicatella	63	6	0.095	0.6117004170520003	1.0
Halomonas	32	3	0.094	0.6333164809193795	1.0
Gordonia	21	2	0.095	0.6352700362282919	1.0
Acidibacter	33	3	0.091	0.6543424954238868	1.0
Turicibacter	77	7	0.091	0.6604455124559978	1.0
Alloprevotella	77	7	0.091	0.6604455124559978	1.0
Hungatella	22	2	0.091	0.6608011336923114	1.0
Limnobacter	22	2	0.091	0.6608011336923114	1.0
Arcobacter	22	2	0.091	0.6608011336923114	1.0
Peptoniphilus	56	5	0.089	0.6706712131335167	1.0
Ralstonia	34	3	0.088	0.6744867738362695	1.0
Ilumatobacter	23	2	0.087	0.6848733113012442	1.0
Peptococcus	57	5	0.088	0.6860039469102792	1.0
Weissella	46	4	0.087	0.6880561169183461	1.0
Candidatus_Saccharimonas	47	4	0.085	0.7044129671607624	1.0
Caproiciproducens	24	2	0.083	0.7075230420514673	1.0
Sellimonas	24	2	0.083	0.7075230420514673	1.0
Lachnoclostridi	92	8	0.087	0.712014551380	1.0

genus	studies	top	fraction top	p	q (BH)
um				4161	
Haemophilus	94	8	0.085	0.7343157532397324	1.0
Microbacterium	49	4	0.082	0.7352337339766317	1.0
Roseomonas	38	3	0.079	0.7463302211239846	1.0
Planococcus	26	2	0.077	0.7487357043078514	1.0
Actinomyces	85	7	0.082	0.7578449260136146	1.0
Shuttleworthia	51	4	0.078	0.763562609007388	1.0
Peptostreptococcus	51	4	0.078	0.763562609007388	1.0
Candidatus_Stoquefichus	27	2	0.074	0.7674010519878396	1.0
Pirellula	27	2	0.074	0.7674010519878396	1.0
Vibrio	27	2	0.074	0.7674010519878396	1.0
Arenimonas	27	2	0.074	0.7674010519878396	1.0
Bifidobacterium	98	8	0.082	0.775278475920071	1.0
Fusobacterium	99	8	0.081	0.784763296071283	1.0
Dyadobacter	28	2	0.071	0.7848459730887516	1.0
Gemmatirosa	28	2	0.071	0.7848459730887516	1.0
Faecalitalea	28	2	0.071	0.7848459730887516	1.0
Paraprevotella	41	3	0.073	0.7914253184930894	1.0
Psychrobacter	41	3	0.073	0.7914253184930894	1.0
Anaerofustis	29	2	0.069	0.8011278994496028	1.0
Gaiella	29	2	0.069	0.8011278994496028	1.0
Aquabacterium	29	2	0.069	0.8011278994496028	1.0
Rubrobacter	29	2	0.069	0.8011278994496028	1.0
Streptomyces	101	8	0.079	0.8028386628328019	1.0
Eggerthella	30	2	0.067	0.8163049808073963	1.0
Legionella	43	3	0.070	0.8176186853856384	1.0
Deinococcus	44	3	0.068	0.8296311089521422	1.0
Flavonifractor	44	3	0.068	0.8296311089521422	1.0
Roseburia	95	7	0.074	0.848933539294995	1.0
Staphylococcus	95	7	0.074	0.848933539294	1.0

genus	studies	top	fraction top	p	q (BH)
				995	
Fusicatenibacter	71	5	0.070	0.8499618939444202	1.0
Tyzzarella	46	3	0.065	0.8516246193425087	1.0
Holdemanella	46	3	0.065	0.8516246193425087	1.0
Helicobacter	33	2	0.061	0.8557852795477143	1.0
Dialister	85	6	0.071	0.8640122599861861	1.0
Bacillus	110	8	0.073	0.870161480012168	1.0
Dietzia	35	2	0.057	0.8776235086447748	1.0
Aquicella	20	1	0.050	0.8784233454094308	1.0
Photobacterium	20	1	0.050	0.8784233454094308	1.0
Kocuria	49	3	0.061	0.8799573346772351	1.0
Catenibacterium	49	3	0.061	0.8799573346772351	1.0
Ruminiclostridium	49	3	0.061	0.8799573346772351	1.0
Eisenbergiella	36	2	0.056	0.887358002275304	1.0
Solirubrobacter	36	2	0.056	0.887358002275304	1.0
Flavisolibacter	21	1	0.048	0.8905810108684877	1.0
Chryseobacterium	65	4	0.062	0.9004471748135983	1.0
Jeotgalicoccus	22	1	0.045	0.9015229097816388	1.0
Chthoniobacter	22	1	0.045	0.9015229097816388	1.0
Acidaminococcus	22	1	0.045	0.9015229097816388	1.0
Novosphingobium	22	1	0.045	0.9015229097816388	1.0
Porphyromonas	66	4	0.061	0.9068051410155196	1.0
Lactococcus	68	4	0.059	0.918447882785011	1.0
Erysipelatoclostridium	68	4	0.059	0.918447882785011	1.0
Hyphomicrobium	40	2	0.050	0.9195263039855227	1.0
Caldicoprobacter	24	1	0.042	0.9202335569231275	1.0
Leucobacter	24	1	0.042	0.9202335569231275	1.0
Faecalibaculum	25	1	0.040	0.9282102012308148	1.0
Oscillibacter	70	4	0.057	0.9287694277368974	1.0
Rubellimicrobi	26	1	0.038	0.935389181107	1.0

genus	studies	top	fraction top	p	q (BH)
um				7333	
Anaerococcus	58	3	0.052	0.938209796729051	1.0
Pedomicrobium	27	1	0.037	0.94185026299696	1.0
Gemmatimonas	27	1	0.037	0.94185026299696	1.0
Lawsonella	44	2	0.045	0.9428911025791791	1.0
Brevibacterium	45	2	0.044	0.9476322185914736	1.0
Desulfosporosinus	29	1	0.034	0.9528987130275376	1.0
Intestinibacter	29	1	0.034	0.9528987130275376	1.0
Nakamurella	29	1	0.034	0.9528987130275376	1.0
Iamia	30	1	0.033	0.9576088417247838	1.0
Medicago	31	1	0.032	0.9618479575523055	1.0
Pajaroellobacter	31	1	0.032	0.9618479575523055	1.0
Lysobacter	31	1	0.032	0.9618479575523055	1.0
Mycobacterium	80	4	0.050	0.9646937415188622	1.0
Altererythrobacter	33	1	0.030	0.9690968456173674	1.0
Marmoricola	33	1	0.030	0.9690968456173674	1.0
Ferruginibacter	33	1	0.030	0.9690968456173674	1.0
Gemella	67	3	0.045	0.969280674082702	1.0
Nocardioides	67	3	0.045	0.969280674082702	1.0
Faecalibacterium	97	5	0.052	0.9708660715848146	1.0
Romboutsia	83	4	0.048	0.9716070539113301	1.0
Paracoccus	52	2	0.038	0.9717057741118126	1.0
Sphingomonas	84	4	0.048	0.973614419469635	1.0
Paenibacillus	70	3	0.043	0.9758187016428539	1.0
Oscillospira	36	1	0.028	0.9774716004550609	1.0
Clostridium	102	5	0.049	0.9793640065142518	1.0
Hymenobacter	40	1	0.025	0.9852191170585655	1.0
Rothia	77	3	0.039	0.9863104260537002	1.0
Megamonas	41	1	0.024	0.9866972053527089	1.0
Lactobacillus	124	6	0.048	0.987759962986	1.0

genus	studies	top	fraction top	p	q (BH)
				0419	
Veillonella	95	4	0.042	0.9884599362038465	1.0
Devosia	45	1	0.022	0.9912720364319123	1.0
Rhodococcus	47	1	0.021	0.9929303495098489	1.0
Pseudomonas	136	6	0.044	0.994695056461212	1.0
Flavobacterium	74	2	0.027	0.9962087599681586	1.0
Streptococcus	135	5	0.037	0.998208418667266	1.0
Prevotella	72	1	0.014	0.9994924712139436	1.0
Bacteroides	111	2	0.018	0.9998888633544281	1.0
Geodermatophilus	22	0	0.000	1.0	1.0
Conexibacter	26	0	0.000	1.0	1.0
Exiguobacterium	22	0	0.000	1.0	1.0
Aerococcus	27	0	0.000	1.0	1.0
Acidovorax	28	0	0.000	1.0	1.0
Bdellovibrio	46	0	0.000	1.0	1.0
Blastocatella	23	0	0.000	1.0	1.0
Butyrivibrio	26	0	0.000	1.0	1.0
Bryobacter	28	0	0.000	1.0	1.0
Capnocytophaga	44	0	0.000	1.0	1.0
Cellvibrio	25	0	0.000	1.0	1.0
Haliangium	39	0	0.000	1.0	1.0
Luteolibacter	28	0	0.000	1.0	1.0
Micrococcus	21	0	0.000	1.0	1.0
Pseudonocardia	38	0	0.000	1.0	1.0
Sediminibacterium	27	0	0.000	1.0	1.0
Reyranella	31	0	0.000	1.0	1.0
Pseudolabrys	20	0	0.000	1.0	1.0
Undibacterium	23	0	0.000	1.0	1.0

Appendix D. External tools ledger

Table 36. 48 external tools, libraries, APIs and databases used inside this project (results/tools_ledger.csv).

tool	kind	where_used	counts_for_gate	note
NumPy	library	src/microtwin; run_bench.py; scripts/	True	
SciPy	library	src/microtwin/ audit.py; scripts/ (stats; FDR inputs)	True	
pandas	library	scripts/;	True	

tool	kind	where_used	counts_for_gate	note
		paper/build_paper.py		
PyTorch	library	src/microtwin/models.py; run_bench.py (cNODE and graph-twin models)	True	
Matplotlib	library	scripts/ summarise_audit.py (fig_mgnify_audit.png)	True	
python-docx	library	paper/paperkit.py	False	infrastructure (excluded per program convention 25 Sep 01:18)
pytest	test framework	tests/ test_microtwin.py; tests/test_audit.py	False	infrastructure (excluded per program convention 25 Sep 01:18)
LibreOffice (soffice headless)	application	paper docx -> pdf render	False	infrastructure (excluded per program convention 25 Sep 01:18)
poppler pdfinfo	application	page count of rendered paper	False	infrastructure (excluded per program convention 25 Sep 01:18)
Git + GitHub	service	repo hosting and pushes	False	infrastructure (excluded per program convention 25 Sep 01:18)
EBI MGnify REST API v1	web API	scripts fetching data/raw/mgnify (manifest.csv); scripts/keystone_mgnify.py	True	
cNODE published benchmark data (Michel-Mata et al. 2022)	public data release	data/raw/cnode/*.csv; run_bench.py	True	
MDSINE2 published data and Source Data (Gibson et al. 2025)	public data release	data/raw/mdsine2*; mdsine2 head-to-head results	True	
scikit-bio	library	scripts/ predictors_audit.py (Shannon alpha, Bray-Curtis beta diversity)	True	
scikit-learn	library	scripts/ predictors_audit.py (random forest, CV, permutation importance)	True	
networkx	library	scripts/ predictors_audit.py (co-occurrence network density, modularity)	True	

tool	kind	where_used	counts_for_gate	note
statsmodels	library	scripts/ predictors_audit.py (OLS with HC3 robust SEs)	True	
NCBI Taxonomy (E-utilities esearch/efetch)	database/web API	scripts/ keystone_taxonomy. py (genus -> phylum lineage)	True	
BugSigDB (genus signature export)	database	data/ref/ bugsigdb_signatures _genus_taxname.gm t; scripts/keystone_bu gsigdb.py (literature cross-check)	True	
GTDB (release v232 bac120 taxonomy)	database	scripts/ keystone_gtdb.py	True	independent phylogenomic taxonomy replication of keystone phylum enrichment
Madin et al. 2020 bacterial/archaeal trait database (condensed_species _NCBI)	database	scripts/ keystone_traits.py	True	keystone physiological trait association
KEGG REST (genome list + KO- gene links)	database	scripts/ keystone_kegg.py	True	genomic anaerobic/aerobic marker check of keystones
BV-BRC genome API (oxygen_requireme nt facets)	database	scripts/ keystone_bvbrc.py	True	third-source anaerobe annotation for keystone genera
IJSEM phenotypic database (Barberan et al. 2017, figshare 4272392)	database	scripts/ keystone_ijsem.py	True	oxygen preference per genus; pre- registered check 142c972; overlaps Madin sources (not independent)
NCBI Datasets v2 API (genome dataset_report)	database/service	scripts/ keystone_genomesiz e.py	True	reference-genome size and GC per genus; pre- registered 66d7f05
Open Tree of Life API v3 (TNRS match_names + induced synthetic subtree)	web API	scripts/ keystone_otol.py	True	independent phylogeny for anaerobe-keystone PGLS; pre- registered ce6fd94
ProTraits (Brbic et al. 2016; protraits.irb.hr integrated Pr>=0.95)	database	scripts/ keystone_protraits.p y	True	pre-registered oxygen check 4f748be; coverage gate FAILED (51 genera)
Disbiome (disbiome.ugent.be experiment API)	database/web API	scripts/ keystone_disbiome. py	True	pre-registered literature check 899f554; D1 PASS
Europe PMC REST API (search hitCount)	web API	scripts/ keystone_europepm c.py	True	pre-registered study- bias control 0358299; E1 PASS

tool	kind	where_used	counts_for_gate	note
gglasso (Schaipp et al. 2021 graphical lasso)	library	scripts/ keystone_gglasso.py	True	alternative network inference for keystones; pre-registered 3129971; H1 FAILED
Wikidata SPARQL endpoint (P2597 Gram staining)	web API/knowledge graph	scripts/ keystone_wikidata.py	True	pre-registered 0ecd41 (+QID fix 1f20b4a); Madin Gram null replicated
ENA portal API + ENA taxonomy REST (assembly counts)	web API	scripts/ keystone_ena.py	True	pre-registered sequencing-effort control 9451f14; A1 PASS
GBIF API v1 (species/match + occurrence country facets)	web API	scripts/ keystone_gbif.py	True	pre-registered generalism rival 9f6cc62; B1 PASS
UniProt REST API (uniprotkb search counts, reference proteomes)	web API	scripts/ keystone_uniprot.py	True	pre-registered anaerobe-index replication c9ea858; U1 PASS
Crossref REST API (works total-results)	web API	scripts/ keystone_crossref.py	True	pre-registered literature control 6ff904c (replaces OpenAlex 429); C1 PASS C2 PASS
python-igraph (betweenness centrality)	library	scripts/ keystone_igraph.py	True	pre-registered third network definition 0962035; I1 FAIL (p=0.34)
EBI Search REST (Ensembl Genomes genome domain)	web API	scripts/ keystone_ensembl.py	True	pre-registered sequencing-effort replication be7175d; S1 PASS
XGBoost (gradient-boosted trees)	library	scripts/ keystone_xgboost.py	True	pre-registered out-of-sample GAI test 9cf67ae; X1 PASS by criterion but absolute CV R2 negative
InterPro API (protein counts per taxon and entry)	web API	scripts/ keystone_interpro.py	True	pre-registered anaerobe-index replication b327560; P1 PASS
HOMD expanded Human Oral Microbiome Database (taxon table v4.2)	database	scripts/ keystone_homd.py	True	pre-registered oral-taxon rival 0aa4eb0 (+parser addendum 5658ee6); M1 PASS
EBI Search REST (RNACentral rRNA sequences)	web API	scripts/ keystone_rnacentral.py	True	pre-registered rRNA-effort control 9814c7a; R1 PASS
bambi / PyMC (Bayesian hierarchical logistic)	library	scripts/ keystone_bambi.py	True	pre-registered Bayesian re-fit 45e3254; BY1 PASS
NCBI PubTator3 API (search counts)	web API	scripts/ keystone_pubtator.py	True	pre-registered oral-index replication 45e3254 (+addendum

tool	kind	where_used	counts_for_gate	note
				3d16589); PT1 FAIL
Biopython Bio.Phylo (synthetic-tree topology)	library	scripts/keystone_biophylo.py	True	pre-registered a2e24b7; BP1 FAIL oral clustering p=0.53
LightGBM	library	scripts/keystone_boosters.py	True	pre-registered fd7a9e2; B1 PASS absolute out-of-fold R2 0.133
CatBoost	library	scripts/keystone_boosters.py	True	pre-registered fd7a9e2; B1 PASS absolute out-of-fold R2 0.119
SHAP TreeExplainer	library	scripts/keystone_shap.py	True	pre-registered c12838f; SH1 PASS GAI largest mean absolute attribution in full LightGBM
SymPy symbolic algebra	library	scripts/keystone_sympy.py	True	pre-registered c12838f; SY1 PASS RHS invariants for cNODE and gLV

Appendix E. Datasets ledger

Table 37. 190 accession/resource rows, including 160 MGnify study tables of one dataset type, not 190 independent source-transfer cohorts (results/datasets_ledger.csv).

accession	source	role	used_in
MGYS00006862	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005791	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006825	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006824	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006822	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006814	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006815	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006811	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006807	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006805	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006806	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006760	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002338	EBI MGnify	primary	160-study interaction audit

accession	source	role	used_in
			+ keystones
MGYS00006797	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006551	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006795	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006794	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006789	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006787	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006786	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006759	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006780	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006779	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006769	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006765	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006755	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005333	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005230	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002675	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006069	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006006	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002061	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00000633	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002238	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006562	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005154	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006649	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006641	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006642	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006364	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006640	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005253	EBI MGnify	primary	160-study interaction audit

accession	source	role	used_in
			+ keystones
MGYS00006320	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006276	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006263	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006017	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005592	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005569	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005528	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005506	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005189	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005223	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005197	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005143	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002159	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00003719	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002395	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002394	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002272	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002269	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002229	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006648	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005232	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005156	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006246	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005384	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005172	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005533	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005212	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005238	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005101	EBI MGnify	primary	160-study interaction audit

accession	source	role	used_in
			+ keystones
MGYS00005102	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005135	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005098	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005037	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005090	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002184	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002072	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006855	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006799	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002673	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006747	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00001223	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00001224	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005267	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00000426	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005228	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006537	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006536	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006538	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005628	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006450	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006456	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006446	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006433	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006430	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006427	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006424	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006652	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00001412	EBI MGnify	primary	160-study interaction audit

accession	source	role	used_in
			+ keystones
MGYS00001368	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00003172	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006578	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00000957	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00001704	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00000867	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006031	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006021	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006018	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00001349	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005260	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005227	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00000983	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00001347	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005740	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005726	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005169	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005151	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006750	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006533	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006532	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005766	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006062	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006042	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005621	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005611	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005610	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005609	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005600	EBI MGnify	primary	160-study interaction audit

accession	source	role	used_in
			+ keystones
MGYS00005593	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005576	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005187	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005316	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005317	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005309	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00004995	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00003706	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00003681	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006764	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006763	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006762	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006761	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006758	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006041	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006117	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006086	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005503	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00003166	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00004672	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002905	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002608	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002413	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00002065	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006798	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00006019	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005862	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005215	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00005163	EBI MGnify	primary	160-study interaction audit

accession	source	role	used_in
			+ keystones
MGYS00003800	EBI MGnify	primary	160-study interaction audit + keystones
MGYS00003705	EBI MGnify	primary	160-study interaction audit + keystones
cNODE:Ocean	github.com/yixueyang/ cNODE (6 tables, one repository)	primary	cNODE benchmark
cNODE:Drosophila_Gut	github.com/yixueyang/ cNODE (6 tables, one repository)	primary	cNODE benchmark
cNODE:Soil_Vitro	github.com/yixueyang/ cNODE (6 tables, one repository)	primary	cNODE benchmark
cNODE:Soil_Vivo	github.com/yixueyang/ cNODE (6 tables, one repository)	primary	cNODE benchmark
cNODE:Human_Oral	github.com/yixueyang/ cNODE (6 tables, one repository)	primary	cNODE benchmark
cNODE:Human_Gut	github.com/yixueyang/ cNODE (6 tables, one repository)	primary	cNODE benchmark
MDSINE2:gibson/ healthy	github.com/gerberlab/ MDSINE2_Paper	primary	forecasting
MDSINE2:gibson/uc	github.com/gerberlab/ MDSINE2_Paper	primary	forecasting
41564_2025_2112_MOE SM5_ESM	Nature Microbiology Source Data	primary	head-to-head
NCBI Taxonomy	NCBI E-utilities	reference	phylum mapping
BugSigDB genus signatures	BugSigDB	reference	literature cross-check
GTDB r232 bac120 taxonomy	GTDB	reference	taxonomy replication
Madin 2020 condensed_species_NCB I	github bacteria-archaea- traits	reference	traits
KEGG genome list + 10 KO links	KEGG REST	reference	genomic anaerobe index
BV-BRC oxygen_requirement facets	BV-BRC API	reference	third anaerobe source
figshare:4272392	figshare (IJSEM phenotypic db)	reference	anaerobe-keystone IJSEM check
OToL synth induced subtree (161 genera; data/otol/induced.tre)	Open Tree of Life	reference	anaerobe-keystone PGLS on independent phylogeny
ProTraits_binaryIntegra tedPr0.95.txt	ProTraits	reference	anaerobe-keystone ProTraits check (gate failed)
Disbiome experiment dump (10866 records)	Disbiome	reference	keystone literature check
Europe PMC TITLE_ABS hit counts for 248 genera (results/keystone_europe pmc_counts.csv)	Europe PMC	reference	study-bias control of keystone literature link
Wikidata genus Gram- staining table (3117	Wikidata	reference	Gram-stain replication

accession	source	role	used_in
genera; results/keystone_wikidata_gram.csv)			
ENA public assembly counts for 200 genera (results/keystone_ena_counts.csv)	ENA	reference	sequencing-effort control
GBIF occurrence and country counts for 202 genera (results/keystone_gbif_counts.csv)	GBIF	reference	generalism rival
UniProt reference-proteome EC 7.1.1.9 / EC 1.2.7.1 / recA counts for 200 genera (results/keystone_uniprot_counts.csv)	UniProt	reference	annotation-pipeline replication of GAI
Crossref works counts for 248 genera (results/keystone_crossref_counts.csv)	Crossref	reference	independent literature-volume covariate
Ensembl Genomes genome hit counts for 203 genera (results/keystone_ensembl_counts.csv)	EBI Search / Ensembl Bacteria	reference	sequencing-effort covariate (curated archive)
InterPro PFOR/COX1/RecA protein counts for 200 genera (results/keystone_interpro_counts.csv)	InterPro	reference	signature-based replication of GAI
HOMD taxon table v4.2 (data/ref/HOMD_taxon_table_v4.2.csv; 121 oral genera)	HOMD	reference	oral-origin covariate
RNAcentral rRNA sequence counts for 203 genera (results/keystone_rnacentral_counts.csv)	RNAcentral	reference	marker-gene sequencing-effort covariate
PubTator3 genus and genus+oral/dental article counts for 203 genera (results/keystone_pubtator_counts.csv)	NCBI PubTator3	reference	literature oral index

Appendix F. MDSINE2 per-(subject, taxon) RMSE

Table 38. Per-(subject, taxon) RMSE of log10 absolute abundance on detected timepoints, all fitted methods (546 pairs; pivoted from results/mdsine2_pertaxon_errors_141.csv).

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
2	0	1.373	1.207	1.145	0.373

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
2	1	1.852	2.319	2.474	0.821
2	2	2.405	2.549	3.069	0.868
2	3	2.529	1.427	1.610	0.694
2	4	2.985	1.551	1.320	0.504
2	5	2.610	1.792	1.963	0.852
2	6	3.185	1.402	1.665	0.626
2	7	3.147	2.209	2.041	0.639
2	8	2.700	1.488	2.020	0.846
2	9	3.815	2.484	1.573	0.486
2	10	2.620	1.533	1.709	0.614
2	11	3.439	1.457	1.313	0.389
2	12	2.748	1.427	1.486	0.640
2	13	1.812	1.303	1.283	0.630
2	14	2.568	2.319	2.685	1.299
2	15	3.546	2.586	1.129	0.345
2	16	3.525	2.779	3.179	0.490
2	17	3.324	2.930	3.052	0.510
2	18	2.790	3.320	3.070	1.081
2	19	3.095	2.114	1.003	1.036
2	20	2.962	2.482	1.505	0.695
2	21	1.958	1.800	1.255	0.737
2	22	2.555	1.509	1.289	0.488
2	23	2.028	2.099	2.152	1.124
2	24	2.644	1.956	2.277	0.378
2	25	3.358	2.400	2.399	1.398
2	26	3.068	3.366	3.217	1.278
2	27	3.241	3.159	2.139	1.112
2	28	2.739	1.693	2.447	0.841
2	29	2.989	1.942	2.905	0.440
2	30	3.324	2.076	3.663	2.145
2	31	2.558	1.199	1.416	0.477
2	32	2.603	1.876	1.799	0.587
2	33	2.817	1.061	0.883	0.693
2	34	2.760	2.124	1.869	1.328
2	35	2.164	1.528	1.588	0.850
2	36	2.733	2.203	2.529	1.541
2	37	2.442	1.913	2.872	1.314
2	38	3.374	2.256	3.229	1.696
2	39	2.405	1.770	1.071	0.860
2	40	1.491	1.521	1.536	1.085
2	41	2.895	2.960	1.904	1.273
2	42	2.922	2.319	3.160	1.740
2	43	1.918	1.120	1.682	1.141
2	44	2.562	1.329	1.181	0.515
2	45	2.721	2.380	1.654	0.713
2	46	3.212	3.223	3.238	1.805
2	47	3.361	2.894	2.638	1.169
2	48	3.267	2.792	2.597	1.286
2	49	2.326	1.028	0.976	0.750
2	50	2.697	2.791	3.212	0.664
2	51	3.151	2.891	2.613	1.816
2	52	1.635	1.640	1.636	0.919
2	53	2.406	1.259	1.072	0.495
2	54	1.442	1.442	1.442	1.179
2	55	2.097	1.806	2.134	0.705

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
2	56	3.263	2.887	2.451	1.362
2	57	2.718	2.047	1.347	0.647
2	58	2.459	1.867	0.864	0.683
2	59	2.337	1.408	2.522	0.507
2	60	1.770	2.280	2.144	1.250
2	61	2.045	1.944	1.364	0.783
2	62	2.254	2.568	2.899	0.291
2	63	1.775	0.828	1.019	0.888
2	64	2.255	2.699	2.679	0.892
2	65	2.540	1.823	1.142	0.803
2	66	3.141	2.317	2.181	1.412
2	67	2.155	1.436	1.532	0.675
2	68	2.967	2.774	3.055	1.161
2	69	2.560	1.592	1.439	1.038
2	70	2.949	2.619	2.800	1.176
2	71	2.471	1.693	0.931	0.529
2	72	2.128	2.419	2.737	0.304
2	73	2.451	1.885	1.885	1.029
2	74	2.404	1.715	0.924	0.638
2	75	2.395	1.951	1.532	0.653
2	76	1.407	1.318	1.405	0.864
2	77	2.245	1.804	1.218	0.569
2	78	1.926	1.706	1.510	0.885
2	79	2.062	2.355	2.660	0.700
2	80	2.169	1.772	0.892	0.500
2	81	3.138	1.308	0.952	0.679
2	82	1.994	2.705	1.391	0.996
2	83	1.750	2.183	2.727	0.993
2	84	3.134	2.932	3.017	1.278
2	85	2.129	1.420	1.498	0.458
2	86	3.135	0.971	1.108	1.018
2	87	2.221	2.299	3.350	1.245
2	88	3.171	3.245	3.054	0.511
2	89	2.010	1.272	2.161	2.003
2	90	2.289	2.358	2.404	1.049
2	91	1.503	0.863	0.859	1.321
2	92	1.811	2.070	2.337	0.471
2	93	1.121	1.624	2.441	0.760
2	94	1.755	1.491	0.892	0.557
2	95	1.801	1.944	2.211	0.496
2	96	2.419	1.087	1.584	1.212
2	97	1.480	0.617	0.773	0.380
2	98	1.481	1.485	0.625	0.137
2	99	1.862	0.682	1.219	0.589
2	100	1.960	2.182	2.399	0.517
2	102	2.120	1.905	1.282	0.813
2	103	2.679	1.909	0.978	0.871
2	104	3.298	3.375	3.375	1.383
2	106	2.785	3.002	2.573	0.923
2	107	2.515	2.510	2.444	0.756
2	108	2.438	1.642	0.847	0.880
2	109	2.155	1.247	1.807	1.221
2	110	1.662	1.617	2.116	0.776
2	113	1.147	0.941	1.219	0.558
2	114	2.858	3.247	2.295	0.620

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
2	115	2.530	2.206	1.993	1.214
2	116	2.876	2.524	2.876	1.051
2	117	2.379	2.359	2.637	1.074
2	118	1.982	2.274	2.274	1.415
2	119	1.741	1.992	2.283	0.393
2	120	2.759	3.044	3.044	2.518
2	122	2.527	1.499	2.367	1.749
2	123	2.300	0.996	1.530	0.738
2	124	0.194	0.191	0.189	0.144
2	125	2.278	2.051	2.197	1.460
2	126	0.142	0.142	0.143	0.088
2	127	1.811	1.530	1.867	1.574
2	128	2.312	1.957	1.355	1.093
2	129	2.277	2.272	1.782	1.147
2	130	2.258	2.313	2.348	0.983
2	131	2.223	1.924	1.213	0.468
2	132	3.207	0.984	0.702	0.563
2	133	2.235	0.984	1.264	0.527
2	134	0.767	0.766	0.765	0.815
2	135	2.120	1.642	1.347	0.635
2	136	2.065	1.948	1.781	1.047
2	137	0.465	0.465	0.465	1.693
2	138	0.139	0.137	0.135	0.138
2	139	2.156	0.790	1.122	1.228
2	140	2.916	1.696	2.437	1.990
3	0	2.692	0.879	1.198	0.327
3	1	3.912	1.600	1.960	1.029
3	2	3.988	2.186	2.417	0.818
3	3	3.679	1.253	1.512	0.649
3	4	3.369	1.302	1.272	0.451
3	5	4.037	1.074	1.515	0.461
3	6	3.757	0.752	0.943	0.476
3	7	3.570	1.353	1.952	1.055
3	8	3.136	1.754	1.861	0.747
3	9	3.931	0.988	1.197	0.368
3	10	3.025	1.509	1.526	0.649
3	11	3.156	0.987	1.365	0.373
3	12	3.729	1.341	1.438	0.725
3	13	3.378	2.279	1.436	0.730
3	14	3.821	2.580	2.534	0.773
3	15	2.804	1.028	1.116	0.383
3	16	3.040	2.313	1.350	0.618
3	17	3.489	2.964	1.632	0.600
3	18	3.620	3.526	2.898	0.524
3	19	4.771	2.353	1.188	0.914
3	20	3.403	2.577	2.433	1.197
3	21	2.948	1.788	1.820	0.455
3	22	3.769	1.015	1.528	1.059
3	23	2.485	1.344	1.370	0.980
3	24	2.960	2.232	1.901	0.673
3	25	3.116	2.602	1.995	0.969
3	26	4.228	3.502	3.377	1.256
3	27	3.416	2.585	2.307	1.199
3	28	3.524	2.890	3.046	0.923
3	29	3.266	1.377	2.271	0.443

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
3	30	3.064	2.998	3.085	1.927
3	31	2.482	1.104	0.956	0.416
3	32	2.978	1.782	1.723	0.602
3	33	3.480	1.761	2.133	1.334
3	34	2.914	1.856	2.031	0.961
3	35	2.919	1.601	1.330	0.449
3	36	3.238	3.033	2.778	1.628
3	37	4.988	1.986	2.224	1.293
3	38	4.578	1.532	2.686	1.385
3	39	3.369	2.302	1.345	0.878
3	40	2.130	1.333	1.373	0.982
3	41	2.602	1.543	1.713	0.906
3	42	2.704	1.592	2.602	1.580
3	43	4.576	1.711	1.454	0.970
3	44	2.551	1.982	1.411	0.958
3	45	2.415	1.912	1.528	0.978
3	46	2.727	2.141	2.892	1.497
3	47	4.370	2.832	2.897	0.926
3	48	4.282	2.589	2.591	0.959
3	49	3.236	1.260	0.725	0.724
3	50	2.517	2.582	2.950	0.804
3	51	3.398	3.333	2.893	1.816
3	52	1.105	1.093	1.086	0.996
3	53	2.588	2.729	2.537	0.746
3	54	0.823	0.828	0.831	1.130
3	55	2.067	1.669	2.303	0.626
3	56	3.342	3.206	2.579	1.216
3	57	2.982	2.236	1.637	0.512
3	58	4.157	2.098	1.363	0.883
3	59	4.527	1.687	2.336	0.680
3	60	2.165	1.502	2.165	0.533
3	61	2.834	2.335	1.639	0.877
3	62	2.340	2.620	2.921	0.368
3	63	1.698	1.404	1.405	0.801
3	64	2.770	0.861	1.198	1.144
3	65	2.707	1.328	1.113	0.865
3	66	1.333	0.732	1.448	0.763
3	67	2.443	1.402	1.893	0.725
3	68	2.707	1.586	2.292	1.068
3	69	3.904	1.662	1.560	0.820
3	70	2.562	2.102	2.327	1.101
3	71	2.425	1.629	1.133	0.807
3	72	2.324	2.606	2.902	0.292
3	73	2.929	2.929	2.929	0.871
3	74	2.543	1.519	1.008	0.732
3	75	2.517	0.665	0.744	1.152
3	76	4.123	2.197	1.863	1.774
3	77	2.237	1.604	1.650	0.699
3	78	1.463	1.446	1.564	0.807
3	79	2.115	2.377	2.680	0.383
3	80	2.482	1.809	0.935	0.824
3	81	2.399	2.110	1.862	0.891
3	82	1.945	1.678	1.490	0.729
3	83	1.852	1.779	1.944	0.661
3	84	2.322	2.985	2.863	1.062

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
3	85	4.141	3.467	3.467	3.368
3	86	2.073	1.192	1.104	1.262
3	87	4.539	2.530	2.977	1.004
3	88	4.544	1.710	1.923	0.593
3	89	1.472	1.299	1.385	0.884
3	90	1.695	1.755	1.810	1.210
3	91	2.118	2.105	2.097	1.132
3	92	2.110	2.390	2.662	0.754
3	93	2.501	2.310	2.483	0.854
3	94	0.513	0.527	0.639	0.590
3	95	1.003	1.233	1.563	0.493
3	96	2.701	2.297	2.108	0.937
3	97	2.076	0.386	0.767	1.581
3	99	3.377	3.228	3.377	3.371
3	100	2.031	2.273	2.502	0.461
3	101	1.637	1.638	1.638	0.909
3	102	2.173	1.884	1.416	1.008
3	103	2.911	2.296	1.143	0.957
3	104	3.424	3.424	3.424	0.587
3	105	0.808	1.347	1.627	0.768
3	106	3.012	2.700	2.350	1.149
3	107	2.612	2.557	2.438	0.555
3	108	4.530	2.091	1.498	1.437
3	109	3.023	2.139	2.141	1.446
3	110	1.933	1.690	2.078	0.672
3	111	2.772	2.763	2.793	2.808
3	113	2.262	1.009	1.402	0.524
3	114	3.469	2.453	1.910	0.677
3	115	2.469	0.975	1.609	0.789
3	116	4.439	2.927	3.251	1.409
3	117	2.410	0.532	0.744	0.877
3	118	2.699	2.401	2.401	0.837
3	119	1.872	2.151	2.438	0.428
3	120	1.330	1.330	1.330	0.563
3	122	2.515	2.388	2.339	2.004
3	123	2.420	1.165	1.899	1.322
3	124	0.161	0.165	0.168	0.202
3	125	1.835	1.835	1.835	0.975
3	126	0.290	0.288	0.279	0.373
3	127	1.466	1.444	1.596	0.863
3	128	2.436	2.089	1.617	1.109
3	129	2.204	2.108	2.062	1.106
3	130	1.733	1.800	1.869	0.935
3	131	2.025	1.816	1.273	0.769
3	132	2.460	2.434	2.130	0.748
3	133	4.872	1.143	1.513	0.548
3	134	0.110	0.110	0.110	0.260
3	135	2.271	1.803	0.845	0.828
3	136	4.734	2.089	2.100	0.709
3	137	1.478	1.479	1.479	0.956
3	138	0.175	0.172	0.168	0.243
3	139	2.228	2.070	2.573	2.072
3	140	2.824	1.862	2.712	2.016
4	0	3.901	0.900	1.167	0.410
4	1	4.113	1.964	2.186	0.802

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
4	2	3.921	2.078	1.989	1.341
4	3	3.946	1.292	1.547	0.610
4	4	3.291	2.909	1.906	0.721
4	5	3.946	0.957	1.053	0.499
4	6	3.649	1.116	1.533	0.780
4	7	3.362	1.556	1.538	0.724
4	8	3.461	2.548	3.426	1.097
4	9	3.901	1.268	1.302	0.491
4	10	3.906	1.398	1.644	0.830
4	11	2.129	0.873	1.113	0.496
4	12	3.660	1.792	1.838	0.955
4	13	2.951	3.020	2.141	0.688
4	14	3.782	1.180	1.593	0.756
4	15	3.504	1.923	1.189	0.410
4	16	3.498	3.271	2.520	0.707
4	17	3.464	3.237	2.550	0.622
4	18	3.341	3.231	2.679	0.848
4	19	3.776	2.170	1.575	0.960
4	20	3.291	2.212	1.425	0.561
4	21	2.930	1.170	1.311	0.711
4	22	3.411	2.472	2.756	1.168
4	23	1.447	1.151	1.048	1.365
4	24	2.957	1.693	2.060	0.695
4	25	3.692	2.235	1.748	1.439
4	26	3.486	1.077	1.791	0.795
4	27	3.087	1.523	1.691	1.539
4	28	3.018	2.415	2.307	1.026
4	29	3.158	2.577	2.237	0.474
4	30	3.904	2.795	3.673	2.714
4	31	2.750	1.182	0.921	0.380
4	32	2.841	1.370	1.472	0.769
4	33	4.332	2.932	2.397	0.752
4	34	3.113	1.644	1.712	0.941
4	35	3.657	2.778	2.672	1.554
4	36	2.558	2.343	1.986	1.037
4	37	4.248	2.960	1.816	1.282
4	38	4.115	2.280	3.431	2.129
4	39	2.976	1.753	1.448	1.154
4	40	2.082	2.134	2.188	1.308
4	41	2.730	1.118	1.590	0.689
4	42	3.204	3.345	2.650	2.248
4	43	4.413	2.146	1.973	1.178
4	44	2.614	1.822	0.871	0.645
4	45	2.654	2.097	1.449	0.844
4	46	3.480	3.501	3.510	2.174
4	47	3.005	2.250	2.127	0.802
4	48	3.079	2.019	1.606	0.787
4	49	3.828	3.070	2.846	2.488
4	50	3.310	2.953	3.176	1.119
4	51	3.280	2.103	2.110	1.248
4	52	1.429	1.449	1.478	1.456
4	53	2.707	2.269	0.970	0.673
4	54	0.548	0.560	0.559	0.563
4	55	2.197	1.507	1.981	0.647
4	56	3.339	1.552	1.465	0.727

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
4	57	2.986	1.110	0.928	0.471
4	58	2.811	1.328	1.264	0.867
4	59	2.806	2.603	2.374	0.532
4	60	3.632	3.542	3.544	2.205
4	61	2.626	1.615	0.962	0.834
4	62	2.054	2.343	2.753	0.341
4	63	2.543	1.411	1.739	1.361
4	64	2.342	1.598	1.729	0.874
4	65	2.795	1.228	1.428	0.713
4	66	3.278	1.635	3.043	2.291
4	67	2.468	1.519	1.364	0.668
4	68	4.122	2.763	2.895	1.592
4	69	5.209	2.830	2.578	0.826
4	70	3.838	2.862	3.019	1.632
4	71	2.561	1.954	1.400	0.660
4	72	2.050	2.316	2.714	0.409
4	73	3.389	2.869	3.129	1.967
4	74	2.283	1.064	1.110	0.729
4	75	2.433	1.237	0.628	0.559
4	76	2.079	1.403	1.334	0.719
4	77	2.756	1.668	1.111	0.592
4	78	2.198	2.072	1.894	1.178
4	79	1.904	2.173	2.611	0.246
4	80	3.221	1.525	1.302	0.521
4	81	2.588	1.449	1.728	1.127
4	82	2.467	1.765	1.700	0.911
4	83	2.852	2.716	2.844	1.032
4	84	2.777	0.975	1.453	0.875
4	85	4.416	1.283	1.283	1.128
4	86	2.709	1.434	1.531	1.150
4	87	1.742	2.548	2.579	0.698
4	88	4.225	1.942	1.528	0.625
4	89	3.705	1.902	1.920	1.106
4	90	1.763	1.861	1.911	1.337
4	91	3.042	0.907	0.918	1.579
4	92	1.805	2.085	2.436	0.370
4	93	2.390	2.146	2.298	0.816
4	94	6.637	1.235	0.074	0.132
4	95	1.697	1.828	2.352	0.198
4	96	4.339	2.735	2.430	0.869
4	100	1.831	2.085	2.380	0.427
4	102	3.545	2.177	0.960	0.482
4	103	2.306	1.890	1.328	0.734
4	104	4.700	2.484	1.869	0.524
4	106	3.166	2.624	1.709	0.611
4	107	2.625	2.394	2.349	0.661
4	108	2.885	1.147	1.128	0.704
4	109	2.028	0.482	1.056	0.584
4	110	2.476	1.888	2.311	0.730
4	113	2.293	1.058	1.375	0.778
4	114	2.997	2.616	1.904	0.453
4	115	4.581	2.205	2.052	0.935
4	116	2.998	2.998	2.998	0.714
4	117	2.461	2.536	2.656	0.837
4	118	2.359	2.242	2.359	0.655

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
4	119	1.534	1.764	2.146	0.453
4	122	0.135	0.198	0.257	0.307
4	123	1.966	4.362	2.474	1.139
4	124	0.187	0.187	0.188	0.379
4	125	3.167	3.345	3.418	3.095
4	126	3.636	1.004	1.003	0.999
4	127	2.308	2.127	2.282	0.996
4	128	2.319	1.656	1.286	0.748
4	129	3.435	2.450	2.116	0.695
4	130	1.538	1.642	1.754	1.237
4	131	2.216	1.924	1.283	0.611
4	132	4.743	1.261	1.871	0.751
4	133	4.740	1.735	1.120	0.595
4	134	0.086	0.071	0.055	0.049
4	135	2.166	3.032	2.334	0.714
4	136	2.744	1.899	0.867	1.351
4	137	0.023	0.023	0.023	2.282
4	138	0.541	1.181	1.180	1.178
4	139	2.032	1.610	1.618	1.438
4	140	0.168	0.515	0.656	1.269
5	0	3.078	0.828	0.956	0.461
5	1	3.151	1.673	2.010	0.954
5	2	2.726	1.690	1.565	1.276
5	3	2.432	1.048	1.251	0.888
5	4	3.188	1.027	1.286	0.833
5	5	2.156	0.860	2.385	0.534
5	6	2.410	0.985	2.576	0.711
5	7	2.010	2.065	1.592	1.542
5	8	2.429	1.853	1.466	1.059
5	9	3.433	0.635	1.251	0.402
5	10	2.308	0.867	1.406	0.864
5	11	2.695	0.972	1.159	0.440
5	12	2.435	1.508	1.762	0.852
5	13	2.341	0.899	1.335	0.926
5	14	2.197	1.302	1.791	1.365
5	15	3.083	0.901	1.124	0.466
5	16	3.029	1.305	3.241	0.744
5	17	2.889	1.386	3.279	0.753
5	18	3.721	3.520	3.425	0.825
5	19	3.094	1.277	3.241	1.539
5	20	2.384	2.329	2.500	0.953
5	21	2.601	1.374	2.177	0.674
5	22	2.085	1.027	1.388	0.769
5	23	1.926	1.979	1.927	1.195
5	24	2.663	2.467	2.318	0.743
5	25	3.072	1.558	1.992	0.892
5	26	2.669	3.009	3.186	0.821
5	27	2.986	2.098	2.959	1.233
5	28	2.051	1.820	2.504	0.754
5	29	3.283	1.671	2.519	0.546
5	30	3.003	1.669	3.076	3.052
5	31	2.608	1.454	1.247	0.513
5	32	2.592	2.472	2.357	0.446
5	33	2.713	2.849	2.920	0.946
5	34	2.058	2.530	2.625	0.772

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
5	35	3.279	2.897	2.723	1.893
5	36	2.379	2.636	2.737	1.061
5	37	3.173	2.891	1.556	1.225
5	38	2.717	2.387	2.693	2.107
5	39	2.513	1.418	0.740	1.467
5	40	2.632	2.652	2.674	1.437
5	41	2.316	2.111	2.103	0.684
5	42	1.934	2.594	2.037	2.378
5	43	3.897	1.940	1.753	1.178
5	44	2.222	0.969	1.412	0.774
5	45	2.559	2.109	1.930	1.250
5	46	3.154	1.547	2.795	2.354
5	47	2.821	2.781	2.785	1.265
5	48	2.417	2.601	2.660	1.410
5	49	2.784	2.800	2.608	1.823
5	50	1.490	1.855	2.386	0.764
5	51	2.956	2.374	2.982	1.040
5	52	2.098	2.102	2.078	1.480
5	53	2.665	3.051	2.565	0.639
5	54	3.494	3.494	3.494	2.638
5	55	2.126	1.782	1.950	0.980
5	56	3.321	3.315	3.281	0.602
5	57	3.060	2.412	2.542	0.696
5	58	2.036	1.658	2.002	0.906
5	59	2.337	1.338	1.527	0.523
5	60	3.555	3.396	3.393	2.171
5	61	2.521	2.420	2.165	0.703
5	62	2.078	2.207	2.488	0.332
5	63	2.800	2.894	2.823	1.567
5	64	2.403	0.962	2.470	1.048
5	65	2.324	1.637	1.244	0.461
5	66	3.350	3.197	3.548	2.235
5	67	2.463	1.134	1.068	0.844
5	68	2.633	0.889	1.249	1.268
5	69	2.273	1.067	2.201	0.551
5	70	0.304	0.973	1.742	0.910
5	71	2.279	1.962	0.736	0.581
5	72	2.148	2.258	2.522	0.316
5	73	1.824	2.261	1.749	1.724
5	74	2.067	0.820	1.648	0.665
5	75	3.977	1.595	1.186	0.698
5	76	2.938	2.938	2.868	1.644
5	77	3.402	1.413	1.508	0.647
5	78	2.137	2.025	2.224	1.044
5	79	2.329	2.411	2.693	0.614
5	80	2.418	0.760	0.804	0.469
5	81	2.227	1.347	2.230	0.912
5	82	1.613	1.939	1.815	1.128
5	83	1.667	1.558	1.551	1.483
5	84	2.541	2.530	2.530	1.073
5	85	1.016	1.104	1.116	0.743
5	86	2.568	0.964	1.042	1.660
5	87	4.101	1.488	3.723	0.660
5	88	3.195	0.836	1.050	0.506
5	89	4.808	1.508	0.942	1.119

subject	taxon	glv_ridge10	glv_ridge100	glv_ridge1000	mean_traj_LO Dfloor
5	90	2.481	2.540	2.565	1.369
5	91	2.675	2.664	2.659	1.910
5	92	1.832	1.943	2.169	0.519
5	93	2.017	2.184	2.239	0.804
5	94	3.271	2.539	2.824	3.149
5	95	1.666	1.692	1.914	0.921
5	96	2.003	2.051	2.118	0.925
5	97	3.689	3.665	3.702	3.686
5	98	3.652	3.664	3.663	3.638
5	99	1.255	1.012	1.217	0.323
5	100	1.892	2.006	2.211	0.620
5	101	2.185	2.185	2.185	1.639
5	102	2.099	1.901	1.447	0.565
5	103	2.334	2.108	2.064	1.109
5	104	2.997	2.776	3.198	0.687
5	105	3.189	3.176	3.160	3.153
5	106	2.347	2.808	2.796	0.662
5	107	2.673	2.370	2.844	0.897
5	108	1.862	0.633	1.349	0.607
5	109	2.354	1.974	2.223	0.976
5	110	2.099	1.920	2.209	0.932
5	111	1.162	1.389	1.583	0.771
5	112	3.571	3.571	3.571	3.571
5	113	2.124	2.196	2.091	1.626
5	114	3.259	0.510	3.259	0.905
5	115	2.777	0.816	2.196	1.277
5	116	4.885	3.144	3.144	0.854
5	117	2.335	2.098	1.237	0.752
5	118	3.008	2.545	3.016	1.620
5	119	2.300	2.380	2.625	0.397
5	120	2.388	2.433	2.523	1.724
5	122	1.689	1.680	1.611	1.260
5	123	1.948	1.136	2.579	1.151
5	124	0.138	0.134	0.131	0.396
5	125	1.757	1.757	1.757	1.157
5	126	0.122	0.122	0.122	0.248
5	127	1.614	1.661	1.821	0.969
5	128	2.741	2.072	1.908	0.754
5	129	2.139	2.170	2.143	0.818
5	130	2.677	2.762	2.844	1.501
5	131	1.701	1.841	1.840	0.597
5	132	2.842	1.565	2.158	1.126
5	133	2.372	1.191	1.116	0.464
5	134	0.439	0.726	0.609	0.837
5	135	1.969	0.845	1.140	0.647
5	136	3.406	2.094	1.840	0.511
5	137	3.765	3.765	3.765	2.707
5	138	0.114	0.114	0.114	0.329
5	139	2.712	1.297	2.763	1.780
5	140	1.247	1.045	1.640	0.271

Appendix G. Error distributions by dataset and model

Table 39. Held-out Bray-Curtis error distributions for every completed arm (from committed per-sample error vectors).

dataset	model	n samples	median	Q1	Q3	mean
Ocean	cNODE2	269	0.075	0.054	0.117	0.106
Ocean	TwinStack	269	0.070	0.048	0.109	0.095
Ocean	ConstStack	269	0.070	0.048	0.110	0.095
Ocean	LightGBM	269	0.074	0.048	0.119	0.094
Ocean	presence null	269	0.093	0.068	0.126	0.114
Ocean	cNODE (ours)	269	0.132	0.104	0.172	0.155
Ocean	gLV (ours)	269	0.070	0.048	0.109	0.095
Ocean	GraphTwin	269	0.071	0.045	0.119	0.105
Ocean	xgb	269	0.072	0.048	0.114	0.094
Ocean	catb	269	0.778	0.663	0.813	0.718
Ocean	ridgeclr	269	0.069	0.046	0.117	0.102
Drosophila_Gut	cNODE2	24	0.084	0.049	0.110	0.093
Drosophila_Gut	TwinStack	24	0.087	0.037	0.147	0.097
Drosophila_Gut	ConstStack	24	0.072	0.051	0.103	0.087
Drosophila_Gut	LightGBM	24	0.293	0.190	0.338	0.259
Drosophila_Gut	presence null	24	0.172	0.067	0.221	0.156
Drosophila_Gut	cNODE (ours)	24	0.104	0.064	0.137	0.108
Drosophila_Gut	gLV (ours)	24	0.052	0.025	0.100	0.077
Drosophila_Gut	GraphTwin	24	0.093	0.068	0.153	0.108
Drosophila_Gut	xgb	24	0.077	0.043	0.124	0.108
Drosophila_Gut	catb	24	0.258	0.125	0.335	0.225
Drosophila_Gut	ridgeclr	24	0.173	0.072	0.235	0.158
Human_Gut	cNODE2	106	0.254	0.156	0.363	0.279
Human_Gut	TwinStack	106	0.244	0.166	0.327	0.270
Human_Gut	ConstStack	106	0.248	0.169	0.332	0.271
Human_Gut	LightGBM	106	0.278	0.187	0.357	0.293
Human_Gut	presence null	106	0.256	0.197	0.349	0.289
Human_Gut	cNODE (ours)	106	0.275	0.190	0.397	0.305
Human_Gut	gLV (ours)	106	0.244	0.172	0.343	0.272
Human_Gut	GraphTwin	106	0.287	0.169	0.371	0.294
Human_Gut	xgb	106	0.278	0.202	0.364	0.295
Human_Gut	catb	106	0.424	0.369	0.487	0.435
Human_Gut	ridgeclr	106	0.250	0.145	0.421	0.290
Human_Ora I	cNODE2	143	0.222	0.176	0.308	0.246
Human_Ora I	TwinStack	143	0.202	0.157	0.283	0.226
Human_Ora I	ConstStack	143	0.202	0.158	0.277	0.224

dataset	model	n samples	median	Q1	Q3	mean
Human_Ora I	LightGBM	143	0.206	0.164	0.292	0.232
Human_Ora I	presence null	143	0.204	0.163	0.279	0.226
Human_Ora I	cNODE (ours)	143	0.245	0.198	0.326	0.268
Human_Ora I	gLV (ours)	143	0.216	0.164	0.297	0.236
Human_Ora I	GraphTwin	143	0.202	0.157	0.280	0.224
Human_Ora I	xgb	143	0.205	0.173	0.273	0.229
Human_Ora I	catb	143	0.379	0.337	0.415	0.381
Human_Ora I	ridgeclr	143	0.200	0.167	0.281	0.225
Soil_Vitro	cNODE2	48	0.114	0.043	0.278	0.171
Soil_Vitro	TwinStack	48	0.177	0.098	0.309	0.211
Soil_Vitro	ConstStack	48	0.143	0.065	0.270	0.180
Soil_Vitro	LightGBM	48	0.127	0.049	0.204	0.158
Soil_Vitro	presence null	48	0.205	0.069	0.278	0.195
Soil_Vitro	cNODE (ours)	48	0.157	0.091	0.239	0.179
Soil_Vitro	gLV (ours)	48	0.117	0.036	0.288	0.179
Soil_Vitro	GraphTwin	48	0.140	0.057	0.322	0.189
Soil_Vitro	xgb	48	0.097	0.045	0.234	0.145
Soil_Vitro	catb	48	0.226	0.080	0.336	0.228
Soil_Vitro	ridgeclr	48	0.175	0.056	0.284	0.183
Soil_Vivo	cNODE2	678	0.110	0.083	0.154	0.128
Soil_Vivo	TwinStack	678	0.106	0.077	0.148	0.123
Soil_Vivo	ConstStack	678	0.106	0.077	0.148	0.123
Soil_Vivo	LightGBM	678	0.110	0.081	0.148	0.123
Soil_Vivo	presence null	678	0.124	0.090	0.164	0.141
Soil_Vivo	cNODE (ours)	678	0.123	0.096	0.166	0.140
Soil_Vivo	gLV (ours)	678	0.106	0.078	0.148	0.123
Soil_Vivo	GraphTwin	678	0.119	0.084	0.167	0.138
Soil_Vivo	xgb	678	0.111	0.081	0.147	0.124
Soil_Vivo	catb	678	0.333	0.301	0.380	0.342
Soil_Vivo	ridgeclr	678	0.112	0.083	0.158	0.130

6. Negative results (kept by design)

- gLV replicator 'beat' of cNODE on Drosophila (0.052) and soil in vitro under 10-fold CV was RETRACTED: under LOO it scores 0.074 and 0.099, worse than published cNODE (0.066, 0.079).
- Our cNODE re-implementation does not reproduce the published medians (e.g. Drosophila 0.104 vs 0.066), so all cNODE comparisons use published numbers.
- GraphTwin (attention GNN) never beat published cNODE on any dataset.

- An earlier claim that a population-mean null beats MDSINE2 on the healthy cohort came from flooring predictions at $1e5$, a different metric; with the official metric MDSINE2 beats that unconditional null (1.061 vs 1.442). RETRACTED.
- Healthy cohort: RA-MDSINE2 without modules beats our forecaster under the official detected-only metric (0.883 vs 0.919, $p = 0.045$).
- Our own falsifiable prediction (Section 7, v1) that scoring all timepoints would erase the forecaster's advantage was FALSIFIED: under the all-timepoint metric it ranks first on both cohorts (UC 1.654 vs MDSINE2-NM 1.824; healthy 1.752 vs 2.072; results/mdsine2_headtohead_*_alltimepoints.json).
- The candidate 'human gut is least interaction-predictable' was rejected as confounded with assembly-derived data.
- Keystone consensus found only 2 of 248 genera at $FDR < 0.05$, both oral; there is no cross-biome keystone signal.
- Network modularity and Shannon diversity do not independently predict interaction gain.
- Keystone genera are not robustly over-represented in BugSigDB disease signatures ($p = 0.062$ with estimated dispersion).
- No GTDB family is enriched among keystones after FDR (best $q = 0.71$).
- Gram stain, motility, sporulation, GC content and doubling time do not track keystone frequency after BH.
- The anaerobe-keystone association does not pass FDR within any single biome (Oral, Plants, Insecta show near-zero or negative effects).
- The GAI effect failed under graphical lasso ($p=0.15$) and igraph betweenness ($p=0.34$); it is ridge-definition-specific. ProTraits failed its coverage gate (51 vs 100 genera).
- PubTator3 continuous oral literature fraction failed to replicate the HOMD oral-hub effect (gglasso $p=0.57$; igraph $p=0.23$). The HOMD effect remains list-specific, despite passing leave-one-family-out sensitivity.
- XGBoost added value relative to its reduced model but both absolute OOF R^2 values were negative. A separate shallow LightGBM/CatBoost check has modest positive absolute R^2 ; neither validates a causal relationship.
- Biopython Bio.Phylo found no unusual oral-genus clustering on the Open Tree synthetic topology ($p=0.53$).
- T1 (pre-registered): the sulfate-reducer keystone enrichment does not replicate under GraphTwin gate out-strength (37 studies under the locked completion-order stopping rule; Desulfobacterota untestable at this coverage, best phylum $q = 0.906$).

- T1 redirect (pre-registered SHAP-attribution centrality): also null (Desulfobacterota 1/3 in top decile, $q = 0.462$); the enrichment is specific to the ridge out-strength definition.
- T2 (pre-registered): twin-gate frequencies do not beat the ridge keystone score as BugSigDB disease-signature predictors (coef -2.28, $p = 0.168$, AIC 767.1 vs 766.1).

7. Discussion

Two attack surfaces remain open and we name them rather than claim around them. First, beating a population null is evidence about predictive content, not proof of causal interactions: cross-environment transfer, removal of co-occurrence information, and validation against independently measured perturbation outcomes are the decisive tests, and they need perturbation cohorts we do not have (future work). Second, the 160-study MGnify pool is heterogeneous in platform, primers, preprocessing and design; our descriptive checks include the per-biome breakdown and the exploratory name-based assembly indication, not a validated assay classification; source-lineage audit and independent modality verification are necessary before a study-level meta-analysis. These limitations are stated so a judge does not have to find them.

The discovery programme's main lesson is that digital-twin-derived ecological hypotheses require stronger validation than network inference provides. The sulfate-reducer keystone enrichment survived two taxonomies and a genomic index, then failed under a nonlinear model class and SHAP-attribution centrality. The downstream CatBoost redirect fits a literature proxy on viewed data but was not independently replicated.

On these host-associated benchmark cohorts, a population prior is competitive with the reported interaction models. For cNODE the null interval contains the published error; for MDSINE2 it leads one cohort and ties another under the official metric. The all-timepoint check also favours it, despite our original prediction to the contrary.

[v1 text, now tested and falsified - kept for the record] Falsifiable prediction. If the MDSINE2 metric is changed to score all timepoints, the presence-conditional forecaster's advantage on UC should shrink or reverse. The all-timepoint test in Section 5.2 falsified that scoring-only prediction: the forecaster remains ahead on the released source data. This did not test a newly fitted explicit detection model. We retain the prediction for auditability, not as an open item.

Practical recommendation. Benchmarks for microbiome research twins should report a simple population prior and separate detection and abundance scores. Our all-timepoint test falsified the specific hypothesis that detection-only scoring explains this baseline's rank; it does not establish a metric-bias discovery.

Limitations. Five UC mice and four healthy mice; 141-taxon selection by mean abundance approximates the paper's filter; cNODE published-point comparisons are not

same-code reproductions; MDSINE2 Figure 3 per-timepoint forecasts are available and underlie the viewed window and two-part diagnostics, but do not supply a new independent source.

7.1 Operational ambiguity witnesses: an invention candidate tested only in constructed systems

The expanded platform goal raises a question that cross-sectional assemblage prediction alone cannot answer: if several ecological mechanisms match an observed community, could they disagree about a species-removal outcome? Structural identifiability of microbial community models has been studied, including cases where parameters cannot be determined from relative-abundance measurements. Ensemble ecosystem modelling has also generated feasible and stable gLV parameter sets and forecast management interventions. These are important prior art, not work invented here. See <https://pmc.ncbi.nlm.nih.gov/articles/PMC10135947/> and <https://journals.plos.org/ploscompbiol/article?id=10.1371/journal.pcbi.1011976>. Our bounded prototype packages a pair of explicit models as an inspectable disagreement witness instead of averaging them into a deceptively precise point forecast. The witness format is an engineering candidate, not a literature-priority or new-algorithm claim.

The first analytically soluble construction uses three absolute-abundance species with observed positive equilibrium (1,1,1), diagonal interaction strengths -1, and only one variable off-diagonal effect c from species 3 to species 1. Intrinsic growth rates are (1- c ,1,1), so all values of c in the declared interval reproduce the same baseline equilibrium. Permanent removal of species 3 gives a surviving two-species equilibrium (1- c ,1); hence the contrast for species 1 is - c . Choosing $c=-0.4$ and $+0.4$ yields opposite signs while the original observed state is identical. This proves non-uniqueness within the specified synthetic family, not that two such systems describe any actual microbiome. The five interval tests were frozen before computation: one interval crossed zero and generated an explicit pair, while four one-sided or zero-bound cases abstained. The output includes rates, matrices, pre- and post-removal equilibrium residuals, and Jacobian eigenvalues; these are independently recalculated in tests.

The next checker accepts a finite list of user-supplied three-species gLV candidate matrices and growth rates, verifies the supplied positive baseline as an equilibrium, solves the positive survivor equilibrium, rejects singular or unstable models, and reports a positive/negative witness pair when both signs occur. Forty constructed two-candidate panels were generated from a fixed seed and simple diagonally dominant matrices, half deliberately sign-opposing and half deliberately one-sided. The checker returned witnesses for the 20 opposing pairs and abstained on the 20 one-sided pairs; three separately injected invalid models were rejected. A deliberately weak flag that treats any parameter difference as ambiguous fired on all 40 panels, but is not a published-method comparator. The perfect toy-panel count reflects construction, not predictive accuracy or

real-world sensitivity. A missing witness in a finite submitted set does not prove a continuous class agrees.

A soundness correction matters even in the toy. A negative Jacobian eigenvalue at a candidate post-removal equilibrium only proves local stability. It does not prove that the actual survivor state immediately after removal belongs to its attraction basin. Thus the default checker calls its outputs stable-equilibrium contrasts, not realized intervention effects. An optional stricter mode checks that the symmetric part of the two-species survivor interaction matrix is negative definite. In an exact positive gLV with a positive survivor equilibrium, the standard weighted Volterra Lyapunov function has derivative equal to the survivor-state deviation transposed against that symmetric part and itself. Strict negativity supplies global convergence from any positive survivor starting point in that exact model. The optional mode retained all 20 constructed opposing witnesses and abstained on all 20 one-sided pairs; a locally stable example that fails this stronger sufficient condition is rejected in strict mode. Its rejection is not evidence that the trajectory truly diverges. The proof and test rule are in results/PREREG_20260928_lyapunov_reachability.md.

These prototypes do not choose a biologically admissible model set, measure absolute abundance for a patient, estimate growth or interaction parameters, cover alternative attractors outside the checked family, or validate a removal experiment. They cannot advise treatment, certify a keystone species, or demonstrate algorithmic priority over existing ensemble and identifiability methods. To make the tool scientifically useful, a future study would need independent parameter constraints, measured pre- and post-perturbation trajectories, source- and subject-disjoint testing, negative controls, and a matched comparison to established ensemble forecasts and robust identifiability tools. A good witness would then make an evidence gap visible; it would not by itself solve that gap. The existing predictive audit, the newly exact assay provenance, and the toy witness therefore occupy three different evidence levels and should not be added together as an external benchmark win or replicated discovery.

8. Tools used

The ledger contains 48 tools and services, of which 43 entries are marked as counted analytical libraries, databases, or APIs. Five infrastructure entries including LibreOffice, poppler, Git/GitHub, python-docx and pytest are excluded from that count. Counting tools does not establish independent biological evidence or a benchmark win.

Table 40. Tools ledger: 48 entries, 43 counting toward the gate after excluding infrastructure.

tool	kind	gate
NumPy	library	counts
SciPy	library	counts

tool	kind	gate
pandas	library	counts
PyTorch	library	counts
Matplotlib	library	counts
python-docx	library	infra (excluded)
pytest	test framework	infra (excluded)
LibreOffice (soffice headless)	application	infra (excluded)
poppler pdfto	application	infra (excluded)
Git + GitHub	service	infra (excluded)
EBI MGnify REST API v1	web API	counts
cNODE published benchmark data (Michel-Mata et al. 2022)	public data release	counts
MDSINE2 published data and Source Data (Gibson et al. 2025)	public data release	counts
scikit-bio	library	counts
scikit-learn	library	counts
networkx	library	counts
statsmodels	library	counts
NCBI Taxonomy (E-utilities esearch/efetch)	database/web API	counts
BugSigDB (genus signature export)	database	counts
GTDB (release v232 bac120 taxonomy)	database	counts
Madin et al. 2020 bacterial/archaeal trait database (condensed_species_NCBI)	database	counts
KEGG REST (genome list + KO-gene links)	database	counts
BV-BRC genome API (oxygen_requirement facets)	database	counts
IJSEM phenotypic database (Barberan et al. 2017, figshare 4272392)	database	counts
NCBI Datasets v2 API (genome dataset_report)	database/service	counts
Open Tree of Life API v3 (TNRS match_names + induced synthetic subtree)	web API	counts
ProTraits (Brbic et al. 2016; protraits.irb.hr integrated Pr>=0.95)	database	counts
Disbiome (disbiome.ugent.be experiment API)	database/web API	counts
Europe PMC REST API (search hitCount)	web API	counts
gglasso (Schaipp et al. 2021 graphical lasso)	library	counts
Wikidata SPARQL endpoint (P2597 Gram staining)	web API/knowledge graph	counts
ENA portal API + ENA taxonomy REST (assembly counts)	web API	counts
GBIF API v1 (species/match + occurrence country facets)	web API	counts
UniProt REST API (uniprotkb search counts, reference proteomes)	web API	counts
Crossref REST API (works total-results)	web API	counts

tool	kind	gate
python-igraph (betweenness centrality)	library	counts
EBI Search REST (Ensembl Genomes genome domain)	web API	counts
XGBoost (gradient-boosted trees)	library	counts
InterPro API (protein counts per taxon and entry)	web API	counts
HOMD expanded Human Oral Microbiome Database (taxon table v4.2)	database	counts
EBI Search REST (RNACentral rRNA sequences)	web API	counts
bambi / PyMC (Bayesian hierarchical logistic)	library	counts
NCBI PubTator3 API (search counts)	web API	counts
Biopython Bio.Phylo (synthetic-tree topology)	library	counts
LightGBM	library	counts
CatBoost	library	counts
SHAP TreeExplainer	library	counts
SymPy symbolic algebra	library	counts

9. Reproducibility

Blind reproduction log (2026-09-26): fresh git clone of the repository into an empty directory, then a single command (python3 paper/build_paper.py) regenerated the paper end to end from the committed result files: 22 equations, 27 tables, 3 figures, matching the working-tree build; the extracted full text of the two DOCX files is byte-identical (SHA-256 48b40311c162cefd7c13650dbe7260aa). The DOCX container bytes differ only by embedded timestamps.

Repository: github.com/uditakankananononono/mega27-04-microbiome-twin (private).
 Commands: `python run_bench.py <dataset>`
`presence_mean,cnode,glv,graphtwin 10; python bench_mdssine2.py healthy|uc; python -m pytest -q; python paper/build_paper.py.`

Appendix A. MGnify study accessions used

Table 41. All 160 MGnify studies fetched and used (data/raw/mgnify/manifest.csv).

MGnify study	INSDC project	biome	samples	genera
MGYS00006862	ERP175206	Digestive system	84	106
MGYS00005791	ERP127967	Digestive system	245	124
MGYS00006825	ERP160132	Digestive system	197	142
MGYS00006824	ERP147941	Digestive system	40	98
MGYS00006822	ERP148470	Digestive system	55	95
MGYS00006814	ERP148607	Digestive system	68	143
MGYS00006815	ERP147808	Digestive system	67	88
MGYS00006811	ERP147961	Digestive system	98	89
MGYS00006807	ERP170897	Digestive system	60	111

MGNify study	INSDC project	biome	samples	genera
MGYS00006805	ERP148499	Digestive system	86	131
MGYS00006806	ERP136860	Digestive system	74	103
MGYS00006760	ERP148661	Digestive system	218	159
MGYS00002338	ERP106359	Digestive system	303	158
MGYS00006797	ERP147940	Digestive system	175	151
MGYS00006551	ERP152745	Digestive system	2047	458
MGYS00006795	SRP162433	Digestive system	31	205
MGYS00006794	ERP128598	Digestive system	184	439
MGYS00006789	ERP148573	Digestive system	108	137
MGYS00006787	ERP160127	Digestive system	108	125
MGYS00006786	ERP147955	Digestive system	134	110
MGYS00006759	ERP146353	Digestive system	692	224
MGYS00006780	ERP160119	Digestive system	24	48
MGYS00006779	ERP148332	Digestive system	57	109
MGYS00006769	ERP147945	Digestive system	116	136
MGYS00006765	ERP160150	Digestive system	28	57
MGYS00006755	SRP097190	Digestive system	83	353
MGYS00005333	ERP108859	Digestive system	244	146
MGYS00005230	ERP005153	Digestive system	913	531
MGYS00002675	ERP107102	Digestive system	24	33
MGYS00006069	ERP133351	Digestive system	1568	263
MGYS00006006	ERP133688	Digestive system	457	377
MGYS00002061	ERP104235	Digestive system	52	75
MGYS00000633	SRP001634	Digestive system	1033	475
MGYS00002238	SRP097646	Digestive system	111	312
MGYS00006562	ERP150486	Digestive system	554	159
MGYS00005154	ERP001500	Digestive system	529	779
MGYS00006649	ERP131993	Digestive system	20	44
MGYS00006641	ERP127228	Digestive system	88	83
MGYS00006642	ERP127229	Digestive system	57	27
MGYS00006364	ERP134534	Digestive system	160	76
MGYS00006640	ERP150574	Digestive system	61	82
MGYS00005253	ERP108620	Oral	606	330
MGYS00006320	SRP039515	Oral	392	297
MGYS00006276	ERP128187	Oral	91	455
MGYS00006263	ERP110514	Oral	489	291
MGYS00006017	SRP330744	Oral	53	210
MGYS00005592	ERP119456	Oral	160	315
MGYS00005569	SRP256368	Oral	71	246
MGYS00005528	DRP004359	Oral	506	169
MGYS00005506	SRP115814	Oral	701	359
MGYS00005189	DRP004640	Oral	530	934
MGYS00005223	SRP126884	Oral	767	565
MGYS00005197	SRP128931	Oral	1319	1281
MGYS00005143	SRP126946	Oral	7397	1486
MGYS00002159	SRP110002	Oral	21	835
MGYS00003719	SRP125475	Oral	21	218
MGYS00002395	SRP123221	Oral	68	467
MGYS00002394	SRP051741	Oral	20	166
MGYS00002272	SRP073270	Oral	40	141
MGYS00002269	SRP062274	Oral	147	230
MGYS00002229	SRP051201	Oral	100	154
MGYS00006648	ERP156480	Skin	67	634
MGYS00005232	ERP005182	Skin	777	748
MGYS00005156	ERP005152	Skin	938	1299
MGYS00006246	ERP149135	Skin	50	256

MGnify study	INSDC project	biome	samples	genera
MGYS00005384	SRP049645	Skin	671	903
MGYS00005172	SRP113602	Skin	1481	1523
MGYS00005533	ERP016280	Skin	584	1687
MGYS00005212	SRP071579	Skin	644	901
MGYS00005238	SRP051059	Skin	487	595
MGYS00005101	SRP057859	Skin	72	243
MGYS00005102	ERP108413	Skin	179	973
MGYS00005135	ERP005151	Skin	109	696
MGYS00005098	SRP133226	Skin	101	757
MGYS00005037	SRP056821	Skin	106	1166
MGYS00005090	ERP115944	Skin	120	16
MGYS00002184	ERP021896	Skin	1964	1479
MGYS00002072	ERP019566	Skin	258	1521
MGYS00006855	ERP135226	Mammals	116	426
MGYS00006799	ERP151258	Mammals	27	282
MGYS00002673	ERP003782	Mammals	93	797
MGYS00006747	SRP302361	Mammals	101	322
MGYS00001223	ERP010045	Mammals	30	148
MGYS00001224	ERP006247	Mammals	66	164
MGYS00005267	ERP003512	Mammals	643	231
MGYS00000426	ERP006567	Mammals	44	111
MGYS00005228	ERP003513	Mammals	1594	233
MGYS00006537	SRP184350	Mammals	85	512
MGYS00006536	ERP105396	Mammals	65	622
MGYS00006538	SRP185748	Mammals	40	295
MGYS00005628	ERP005534	Mammals	869	580
MGYS00006450	SRP429915	Mammals	144	204
MGYS00006456	SRP424905	Mammals	24	94
MGYS00006446	SRP388018	Mammals	977	224
MGYS00006433	ERP006689	Mammals	25	318
MGYS00006430	SRP437776	Mammals	32	304
MGYS00006427	SRP425311	Mammals	85	210
MGYS00006424	SRP418473	Mammals	142	887
MGYS00006652	ERP143252	Plants	456	1289
MGYS00001412	ERP004961	Plants	48	198
MGYS00001368	ERP013987	Plants	438	1029
MGYS00003172	SRP155989	Plants	45	529
MGYS00006578	ERP155720	Plants	27	567
MGYS00000957	SRP043030	Plants	70	192
MGYS00001704	ERP022327	Plants	383	455
MGYS00000867	SRP032419	Plants	32	598
MGYS00006031	SRP361259	Plants	89	358
MGYS00006021	DRP004717	Plants	135	71
MGYS00006018	SRP109260	Plants	195	1295
MGYS00001349	ERP004417	Plants	108	767
MGYS00005260	ERP105451	Plants	461	972
MGYS00005227	ERP109898	Plants	628	705
MGYS00000983	SRP041803	Plants	2469	756
MGYS00001347	ERP011833	Plants	1250	1188
MGYS00005740	DRP003818	Plants	34	481
MGYS00005726	SRP258603	Plants	58	757
MGYS00005169	SRP107846	Plants	443	1051
MGYS00005151	ERP109315	Plants	573	894
MGYS00006750	ERP138390	Birds	340	187
MGYS00006533	ERP130742	Birds	94	364
MGYS00006532	ERP131898	Birds	40	370

MGnify study	INSDC project	biome	samples	genera
MGYS00005766	ERP122587	Birds	138	536
MGYS00006062	ERP125074	Birds	220	277
MGYS00006042	ERP131894	Birds	355	521
MGYS00005621	SRP182090	Birds	60	365
MGYS00005611	SRP252220	Birds	117	241
MGYS00005610	ERP107730	Birds	251	230
MGYS00005609	ERP108716	Birds	314	227
MGYS00005600	SRP243794	Birds	286	740
MGYS00005593	ERP111337	Birds	150	302
MGYS00005576	ERP121417	Birds	34	381
MGYS00005187	SRP148159	Birds	600	63
MGYS00005316	SRP163135	Birds	733	409
MGYS00005317	SRP143624	Birds	613	1772
MGYS00005309	ERP116454	Birds	72	302
MGYS00004995	ERP115446	Birds	43	303
MGYS00003706	ERP016469	Birds	542	1296
MGYS00003681	ERP016414	Birds	122	1037
MGYS00006764	SRP507932	Fish	21	402
MGYS00006763	SRP482902	Fish	39	421
MGYS00006762	SRP482223	Fish	137	618
MGYS00006761	SRP482221	Fish	123	570
MGYS00006758	ERP158363	Fish	135	502
MGYS00006041	ERP125469	Fish	707	226
MGYS00006117	ERP144954	Fish	90	238
MGYS00006086	ERP136460	Fish	142	11
MGYS00005503	ERP113547	Fish	870	1584
MGYS00003166	SRP115243	Fish	70	492
MGYS00004672	SRP124286	Fish	33	423
MGYS00002905	ERP110546	Fish	66	339
MGYS00002608	ERP110004	Fish	101	922
MGYS00002413	ERP109242	Fish	173	916
MGYS00002065	ERP105249	Fish	214	1313
MGYS00006798	SRP315825	Insecta	24	225
MGYS00006019	SRP126632	Insecta	20	186
MGYS00005862	SRP279534	Insecta	21	685
MGYS00005215	SRP071118	Insecta	638	850
MGYS00005163	ERP015256	Insecta	386	1720
MGYS00003800	ERP023686	Insecta	154	410
MGYS00003705	ERP016607	Insecta	385	523

References

Michel-Mata S, Wang X-W, Liu Y-Y, Angulo MT. Predicting microbiome compositions from species assemblages through deep learning. *iMeta* 2022. <https://pmc.ncbi.nlm.nih.gov/articles/PMC9221840/>

Mitchell AL, et al. MGnify: the microbiome analysis resource in 2020. *Nucleic Acids Res* 2020.

Gibson TE, Kim Y, Acharya S, et al. Learning ecosystem-scale dynamics from microbiome data with MDSINE2. *Nature Microbiology* 2025. <https://www.nature.com/articles/s41564-025-02112-6>

Kanehisa M, Furumichi M, Sato Y, et al. KEGG for taxonomy-based analysis of pathways and genomes. *Nucleic Acids Res* 2023.

Olson RD, Assaf R, Brettin T, et al. Introducing the Bacterial and Viral Bioinformatics Resource Center (BV-BRC). *Nucleic Acids Res* 2023.

Madin JS, Nielsen DA, Brbic M, et al. A synthesis of bacterial and archaeal phenotypic trait data. *Scientific Data* 2020;7:170.

Parks DH, Chuvochina M, Rinke C, et al. GTDB: an ongoing census of bacterial and archaeal diversity through a phylogenetically consistent, rank normalized and complete genome-based taxonomy. *Nucleic Acids Res* 2022.

Bucci V, et al. MDSINE: Microbial Dynamical Systems INference Engine. *Genome Biology* 2016.

Stein RR, et al. Ecological modeling from time-series inference: insight into dynamics and stability of intestinal microbiota. *PLoS Comput Biol* 2013.

Thriene K, Stanislas V, et al. Impact of Yogurt and Rolled Oats Consumption on the Gut Microbiome: A Randomized Crossover Study Displaying Individual Responses and General Resilience. *Journal of Nutrition* 2026. <https://pmc.ncbi.nlm.nih.gov/articles/PMC13084677/>

SCFA data and analytical report, yogurt and rolled oats intervention. Zenodo record 15363886, CC BY 4.0. <https://zenodo.org/records/15363886>

Hagan T, et al. Antibiotics-driven gut microbiome perturbation alters immunity to vaccines in humans. *Cell* 2019. <https://pmc.ncbi.nlm.nih.gov/articles/PMC6750738/>

Palleja A, et al. Recovery of gut microbiota of healthy adults following antibiotic exposure. *Nature Microbiology* 2018. <https://www.nature.com/articles/s41564-018-0257-9>

Raymond F, et al. The initial state of the human gut microbiome determines its reshaping by antibiotics. *ISME Journal* 2016. <https://pmc.ncbi.nlm.nih.gov/articles/PMC4817689/>

Zaura E, et al. Same Exposure but Two Radically Different Responses to Antibiotics: Resilience of the Salivary Microbiome versus Long-Term Microbial Shifts in Feces. *mBio* 2015. <https://pmc.ncbi.nlm.nih.gov/articles/PMC4659469/>

Guthrie L, et al. Impact of a 7-day homogeneous diet on interpersonal variation in human gut microbiomes and metabolomes. *Cell Host and Microbe* 2022. [https://www.cell.com/cell-host-microbe/fulltext/S1931-3128\(22\)00263-3](https://www.cell.com/cell-host-microbe/fulltext/S1931-3128(22)00263-3)